

ECE 661 Homework 7

Michael Goldberg

October 30 2024

1 Theoretical Question

1.1 Custom Texture Detector

I think a cool texture detector could extend on the idea behind LBP, but instead of using a single channel, use multiple channels. Texture is encoded and interpreted very richly with the human experience, and doesn't solely lie within hue or brightness alone. Therefore, I propose a 3-channel concatenated LBP algorithm. The way this works is one first converts their image into the HSV colorspace, as the three channels are more independent of image features from one another than something like RGB. For example, if we were to image a yellow ball, the red channel and green channel would encode relatively the same information - however the hue channel and saturation channel would encode vastly different information. There could probably be some cases pointed out about where HSV has this same phenomenon happen, but with my experience with image processing, the HSV map tends to have more independent channels, so I will explore using HSV for now. The next step is to apply LBP to each of these channels individually. We then finally concatenate all three LBP histograms together to form a vector of size $(num_{neighbors} + 2) * 3$. The larger feature vector allows us to encode a richer representation of the image's texture, and should fare better than traditional LBP in a classifier. I will show a comparison of LBP versus complex LBP when subjected to an SVM classifier in the later programming section when I go over my SVM.

2 Programming Tasks

The following task was to take a dataset of images taken in four classifications of weather: cloudy, rain, shine, and sunrise and to train an SVM classifier with varying texture descriptors. The four descriptors used were Gram Matrices from VGG19 encodings, Gram Matrices from ResNet50-Coarse encodings, Gram Matrices from ResNet50-Fine encodings, and LBP histograms. The following are examples of an image from each class:



(a) Cloudy Image



(b) Rain Image



(c) Shine Image



(d) Sunrise Image

2.1 LBP Implementation

The first step was to create the LBP feature extraction algorithm from scratch. LBP is performed on a single-channel image, and we were instructed to use the hue channel of the images converted from RGB to HSV. The LBP algorithm assigns a texture value to a pixel by counting how many neighboring pixel values are larger than it and forming a histogram of these counts, with some exceptions which I will go into later. LBP takes a few parameters, mainly the radius of the search r as well as how many neighbor points to consider N . The algorithm starts by calculating the angle that each of these neighbor points form, and calculating their subpixel locations using the following equation.

$$x_n = j + r \cos \theta_n, y_n = i - r \sin \theta_n \quad (1)$$

Here, i and j represent the integer coordinates of the central pixel, $n \in N$ represents the index of the n th neighbor value, and θ_n represents the angle of the n th neighbor point. The subpixel value of this location can then be calculated using bilinear interpolation.

Once this subpixel value is calculated, we can create a binary word of size N bits, where each bit corresponds to a certain neighbor. If that particular neighbor's value is larger than the central pixel's value, its corresponding bit is set to 1.


This word is then circularly rotated such that the longest string of 0s are at the left. We then can assign the value of the central pixel using the following rules:

- If the word is uniform, meaning there is a single transition from 0s to 1s when reading left to right, then the central pixel is assigned an integer equal to the number of 1s in the word


- If the word is not uniform, meaning there is a transition from 1s to 0s, then the central pixel is assigned an integer equal to $N + 1$

This process is repeated for all pixels across the image, and a histogram is created counting up the values assigned to each pixel and placing it in each bin. Since the binary word is N elements long, there are $N + 2$ bins (considering the cases of non-uniform words and words equal to 0).


Here are what the LBP images and histograms look like for each of the original images shown above with $N = 8$ and $r = 1$:




(a) Cloudy LBP Image




(b) Rain LBP Image




(c) Shine LBP Image




(d) Sunrise LBP Image




(a) Cloudy Histogram



(b) Rain Histogram



(c) Shine Histogram




(d) Sunrise Histogram

2.2 Gram Matrices


Gram matrices are formed by simply taking a feature vector F_l and multiplying it by its transpose:

$$G = F_l * F_l^T \tag{2}$$


Both VGG19 and ResNet encode images into feature vectors, so I just had to perform the above calculation on the results of each of the encodings to find their gram matrices. Below are the gram matrices formed from VGG19, ResNet50-Coarse, and ResNet50-Fine on the above original images. Additionally, from here on out, I will refer to these encodings as vgg, resnet-coarse, and resnet-fine. Please zoom in to see the slight differences between them.




(a) Cloudy Gram Matrices



(b) Rain Gram Matrices



(c) Shine Gram Matrices



(d) Sunrise Gram Matrices

However, because G is a square symmetric matrix this is both redundant and incompatible with an SVM. We therefore need to transform G into a 1D vector, and remove the duplicate entries due to symmetry. Therefore, all that was done was to flatten the upper-triangular representation of G , and that was the feature vector used for training the SVM with vgg, resnet-coarse, and resnet-fine.

2.3 Results

An SVM was trained using scikit-learn with the following best-performing hyperparameters for vgg, resnet-coarse, and resnet-fine:

- batch size = 1024
- kernel = "linear"
- probability = True
- C = 1

However, these hyperparameters performed best for LBP:

- batch size = 1024
- kernel = "rbf"
- probability = True
- C = 10000
- N (number of neighbors for LBP) = 16
- r (radius for LBP) = 2

Here are the resultant confusion matrices and accuracies for each texture descriptor:




Figure 5: VGG Confusion Matrix, Accuracy = 91%




Figure 6: Resnet-Coarse Confusion Matrix, Accuracy = 88.5%




Figure 7: Resnet-Fine Confusion Matrix, Accuracy = 96%




Figure 8: LBP Confusion Matrix, Accuracy = 50.5%

Let's take at some images that were correctly classified:

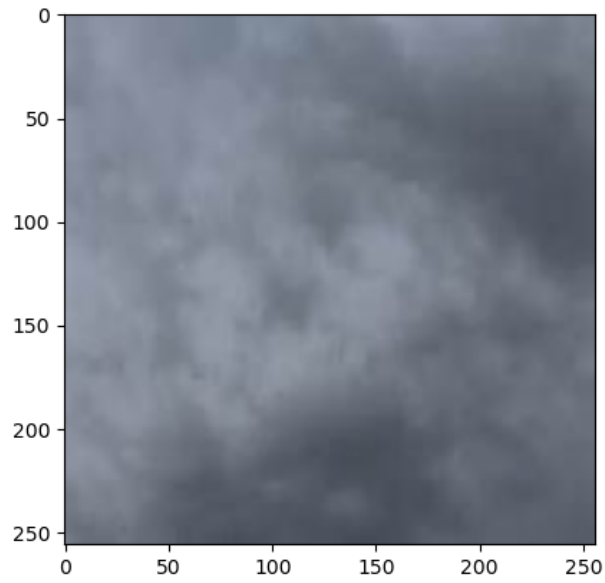


Figure 9: Correctly Classified Image with VGG Descriptors: Predicted=Actual=Cloudy




Figure 10: Correctly Classified Image with Resnet-Coarse Descriptors: Predicted=Actual=Cloudy




Figure 11: Correctly Classified Image with Resnet-Fine Descriptors: Predicted=Actual=Shine

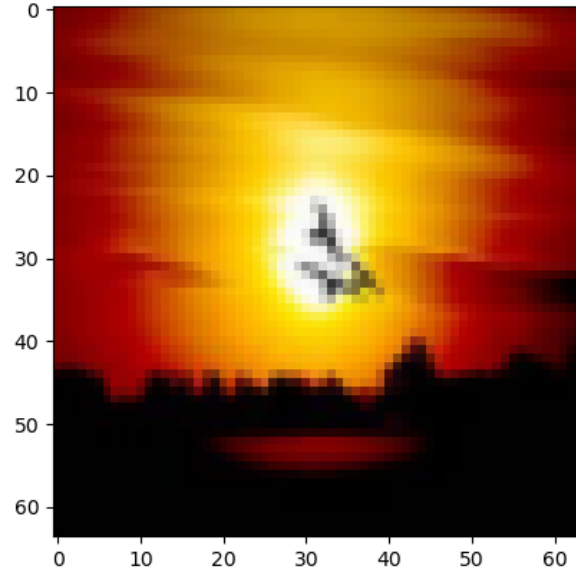


Figure 12: Correctly Classified Image with LBP Descriptors: Predicted=Actual=Sunrise

For vgg and both resnet feature descriptors, I was able to get classification accuracy at close to or above 90%, with the best performer being resnet-fine at 96% accuracy. It is clear, however, why LBP works significantly worse than the other texture descriptors, and that is because a histogram of size N is too small to depict the richer representation of texture that the other strategies can convey. This is why I settled on my custom texture descriptor being the 3-channel LBP algorithm, as this triples the size of the feature vector, while also taking advantage of the other channels to extract unique texture features from the image. Below is a comparison of the complex LBP trained on the same SVM with the same hyperparameters as the standard LBP implementation:

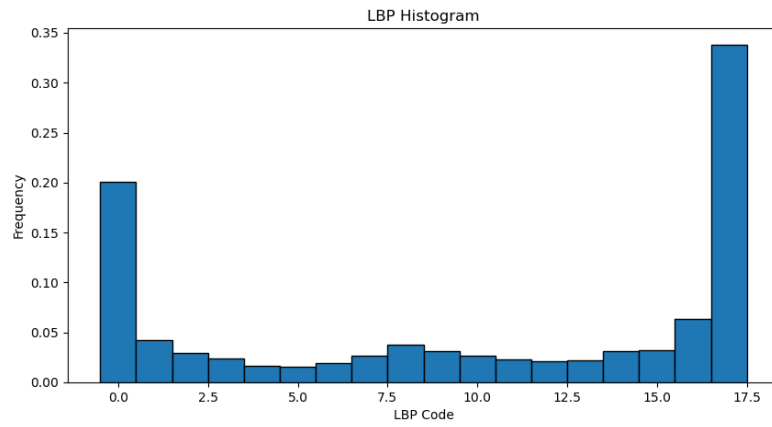


Figure 13: LBP Histogram

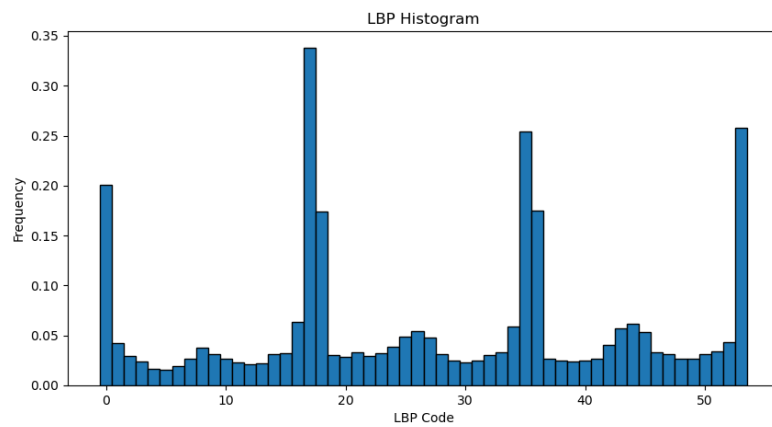


Figure 14: Complex LBP Histogram




Figure 15: LBP Confusion Matrix, Accuracy = 50.5%




Figure 16: Complex LBP Confusion Matrix, Accuracy = 87%

This definitely increased the performance of the LBP methodology significantly, and it's pretty cool to see a custom implementation come to fruition like this.

However none of these descriptors were without error. Below are some example images that were misclassified by the networks:




Figure 17: Incorrectly Classified Image with VGG Descriptors: Predicted=Shine, Actual=Rain




Figure 18: Incorrectly Classified Image with Resnet-Coarse Descriptors:
Predicted-Shine, Actual=Rain




Figure 19: Incorrectly Classified Image with Resnet-Fine Descriptors: Predicted=Shine, Actual=Cloudy

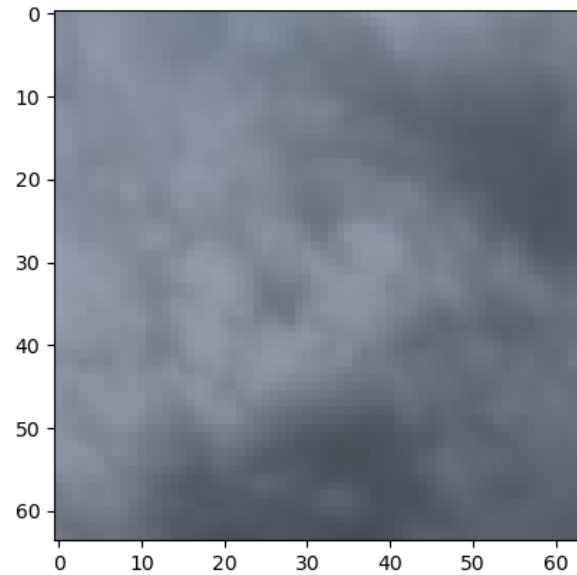


Figure 20: Incorrectly Classified Image with LBP Descriptors: Predicted=Rain, Actual=Cloudy

3 Code

```

1  import os
2  import numpy as np
3  import torch
4  import torch.nn as nn
5  import matplotlib.pyplot as plt
6  import importlib
7  import seaborn as sns
8
9  from skimage import io, transform
10 from skimage.measure import block_reduce
11 from torchvision.models import ResNet50_Weights
12 from sklearn import svm
13 from sklearn.metrics import classification_report, accuracy_score,
    confusion_matrix, ConfusionMatrixDisplay
14 from sklearn.preprocessing import StandardScaler
15 from sklearn.model_selection import GridSearchCV
16 from tqdm import tqdm
17
18 num_neighbors = 16
19 lbp_radius = 2
20
21 downsample_size = 32

```

```

22 device = torch.device("mps" if torch.backends.mps.is_available()
23                        else "cpu")
24
25 label_map = {"cloudy": 0, "rain": 1, "shine": 2, "sunrise": 3}
26 reverse_label_map = ["cloudy", "rain", "shine", "sunrise"]
27
28 class VGG19(nn.Module):
29     def __init__(self):
30         super().__init__()
31         self.model = nn.Sequential(
32             # encode 1-1
33             nn.Conv2d(3, 3, kernel_size=(1, 1), stride=(1, 1)),
34             nn.Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
35             nn.ReLU(inplace=True), # relu 1-1
36             # encode 2-1
37             nn.Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
38             nn.ReLU(inplace=True),
39             nn.MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False), #1/2
40
41             nn.Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
42             nn.ReLU(inplace=True), # relu 2-1
43             # encoder 3-1
44             nn.Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
45             nn.ReLU(inplace=True),
46
47             nn.MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False), #1/4
48             nn.Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
49             nn.ReLU(inplace=True), # relu 3-1
50             # encoder 4-1
51             nn.Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
52             nn.ReLU(inplace=True),
53             nn.Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
54             nn.ReLU(inplace=True),
55             nn.Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
56             nn.ReLU(inplace=True),
57             nn.MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False), #1/8
58
59             nn.Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
60             nn.ReLU(inplace=True), # relu 4-1
61             # rest of vgg not used
62             nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
63             nn.ReLU(inplace=True),

```

```

64         nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
65         nn.ReLU(inplace=True),
66         nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
67         nn.ReLU(inplace=True),
68         nn.MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False), #1/16
69
70         nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), padding_mode='reflect'),
71         nn.ReLU(inplace=True), # relu 5-1
72         # nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), padding_mode='reflect'),
73         # nn.ReLU(inplace=True),
74         # nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), padding_mode='reflect'),
75         # nn.ReLU(inplace=True),
76         # nn.Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), padding_mode='reflect'),
77         # nn.ReLU(inplace=True)
78     )
79
80     def load_weights(self, path_to_weights):
81         vgg_model = torch.load(path_to_weights, weights_only=True)
82         # Don't care about the extra weights
83         self.model.load_state_dict(vgg_model, strict=False)
84         for parameter in self.model.parameters():
85             parameter.requires_grad = False
86
87     def forward(self, x):
88         # Input is numpy array of shape (H, W, 3)
89         # Output is numpy array of shape (N_l, H_l, W_l)
90         x = torch.from_numpy(x).permute(0, 3, 1, 2).float().to(
device)
91         out = self.model(x)
92         out = out.cpu().numpy()
93         return out
94
95     def extract_features(self, images, batch_size=32):
96         vgg_features = []
97
98         for i in range(0, len(images), batch_size):
99             batch_images = images[i:i + batch_size]
100
101             features = self.forward(batch_images)
102             vgg_features.append(features)
103
104             torch.mps.empty_cache()
105
106         return np.vstack(vgg_features)
107
108     def class_for_name(module_name, class_name):
109         # load the module, will raise ImportError if module cannot be
loaded
110         m = importlib.import_module(module_name)
111         return getattr(m, class_name)

```

```

112
113 class CustomResNet(nn.Module):
114     def __init__(self,
115                 encoder='resnet50',
116                 weights=ResNet50_Weights.DEFAULT):
117
118         super(CustomResNet, self).__init__()
119         assert encoder in ['resnet18', 'resnet34', 'resnet50', '
resnet101', 'resnet152'], "Incorrect encoder type"
120         # if encoder in ['resnet18', 'resnet34']:
121         #     filters = [64, 128, 256, 512]
122         # else:
123         #     filters = [256, 512, 1024, 2048]
124         resnet = class_for_name("torchvision.models", encoder)(
weights=weights)
125
126         for parameter in resnet.parameters():
127             parameter.requires_grad = False
128
129         self.firstconv = resnet.conv1 # H/2
130         self.firstbn = resnet.bn1
131         self.firstrelu = resnet.relu
132         self.firstmaxpool = resnet.maxpool # H/4
133
134         # encoder
135         self.layer1 = resnet.layer1 # H/4
136         self.layer2 = resnet.layer2 # H/8
137         self.layer3 = resnet.layer3 # H/16
138
139     def forward(self, x):
140         """
141         Coarse and Fine Feature extraction using ResNet
142         Coarse Feature Map has smaller spatial sizes.
143         Arg:
144             x: (np.array) [H,W,C]
145         Return:
146             xc: (np.array) [C_coarse, H/16, W/16]
147             xf: (np.array) [C_fine, H/8, W/8]
148         """
149         x = torch.from_numpy(x).permute(0, 3, 1, 2).float().to(
device)
150
151         x = self.firstrelu(self.firstbn(self.firstconv(x))) #1/2
152         x = self.firstmaxpool(x) #1/4
153
154         x = self.layer1(x) #1/4
155         xf = self.layer2(x) #1/8
156         xc = self.layer3(xf) #1/16
157
158         # convert xc, xf to numpy
159         xc = xc.cpu().detach().numpy()
160         xf = xf.cpu().detach().numpy()
161         return xc, xf
162
163     def extract_features(self, images, batch_size=32):
164         coarse_features = []
165         fine_features = []

```



```

166         for i in range(0, len(images), batch_size):
167             batch_images = images[i:i + batch_size]
168
169             features_coarse, features_fine = self.forward(
170                 batch_images)
171
172             coarse_features.append(features_coarse)
173             fine_features.append(features_fine)
174
175             torch.mps.empty_cache()
176
177         return np.vstack(coarse_features), np.vstack(fine_features)
178
179 def bilinear_interpolate(image, x, y):
180     # get the four surrounding pixel values
181     x0, y0 = int(x), int(y)
182     x1, y1 = min(x0 + 1, image.shape[1] - 1), min(y0 + 1, image.
183         shape[0] - 1)
184
185     # calculate weights for each pixel based on distance from
186     # integer pixel locations
187     wa = (x1 - x) * (y1 - y)
188     wb = (x1 - x) * (y - y0)
189     wc = (x - x0) * (y1 - y)
190     wd = (x - x0) * (y - y0)
191
192     # weighted sum to find interpolated value
193     interpolated_value = wa * image[y0, x0] + wb * image[y1, x0] +
194         wc * image[y0, x1] + wd * image[y1, x1]
195     return interpolated_value
196
197 def rgb2hsv(image):
198     # cv2.imshow("original", image)
199
200     # normalize the image and convert from bgr to rgb
201     image = np.asarray(image, dtype=float) / 255.0
202
203     r, g, b = image[:, :, 0], image[:, :, 1], image[:, :, 2]
204
205     chroma_max = np.maximum(np.maximum(r, g), b) # maximum chroma
206     # is the maximum value across all three RGB channels
207     chroma_min = np.minimum(np.minimum(r, g), b) # minimum chroma
208     # is the minimum value across all three RGB channels
209     delta = chroma_max - chroma_min # delta is the difference in
210     # max vs. min chroma
211
212     hue = np.zeros_like(chroma_max)
213     mask = delta != 0
214
215     # Using well-established conversion between rgb and hue, using
216     # masking to prevent division by 0 and max_chroma checking to
217     # determine which equation to use
218     hue[mask & (chroma_max == r)] = ((g[mask & (chroma_max == r)] -
219         b[mask & (chroma_max == r)]) / delta[mask & (chroma_max == r)
220         ]) % 6

```

```

211     hue[mask & (chroma_max == g)] = ((b[mask & (chroma_max == g)] -
    r[mask & (chroma_max == g)]) / delta[mask & (chroma_max == g)
    ]) + 2
212     hue[mask & (chroma_max == b)] = ((r[mask & (chroma_max == b)] -
    g[mask & (chroma_max == b)]) / delta[mask & (chroma_max == b)
    ]) + 4
213
214     hue *= 60 # convert to degrees on the color wheel
215     hue[hue < 0] += 360 # prevent negative values
216
217     # Using well-established saturation calculation where S = 0 if
    chroma_max = 0 and S = delta / chroma_max otherwise
218     saturation = np.zeros_like(chroma_max)
219     saturation[chroma_max != 0] = delta[chroma_max != 0] /
    chroma_max[chroma_max != 0]
220
221     value = chroma_max
222
223     hsv_image = np.stack([hue, saturation, value], axis=-1)
224
225     return hsv_image
226
227 def check_if_code_uniform(code):
228
229     # if code is 0 or negative, nonuniform for purposes of lbp
230     if code <= 0:
231         return False
232
233     # collect the unlabelled binary representation of the code
234     binary = bin(code)[2:]
235
236     # set previous bit to first bit
237     prev_bit = binary[0]
238
239     # if there is ever a situation where the previous bit is 1 and
    the next bit is 0, nonuniform
240     for bit in binary[1:]:
241         if (prev_bit == "1" and bit == "0"):
242             return False
243         prev_bit = bit
244
245     return True
246
247 def lbp_histogram(lbp_image, num_bins=num_neighbors + 2):
248     histogram, _ = np.histogram(lbp_image.ravel(), bins=num_bins,
    range=(0, num_bins))
249
250     # normalization
251     histogram = histogram.astype("float")
252     histogram /= (histogram.sum() + 1e-6)
253
254     return histogram
255
256 def lbp_descriptor(image, radius=lbp_radius, num_neighbors=
    num_neighbors):
257
258     # collect hue channel of image and normalize to 0-255

```

```

259 image = (rgb2hsv(image)[: , : , 0] / 360.0) * 255.0
260
261
262 height, width = image.shape
263
264 lbp_image = np.zeros((height, width), dtype=np.uint8)
265
266 # creation of array of angles that each neighbor point makes
    with the central point
267 angles = [2 * np.pi * i / num_neighbors for i in range(
    num_neighbors)]
268
269 for i in range(radius, height - radius):
270     for j in range(radius, width - radius):
271         center = image[i, j]
272         lbp_code = 0
273
274         for idx, angle in enumerate(angles):
275             # collecting subpixel value of neighbor point
276             x = j + radius * np.cos(angle)
277             y = i - radius * np.sin(angle)
278
279             # bilinear interpolation to determine neighbor's
    subpixel value
280             neighbor = bilinear_interpolate(image, x, y)
281
282             # setting bits to 1 if their corresponding neighbor
    is greater than the center pixel
283             lbp_code |= (neighbor > center) << idx
284
285             min_val = lbp_code
286
287             # circularly shifting to produce the largest number of
    zeros on the left of the lbp code
288             for _ in range(num_neighbors):
289                 lbp_code = (lbp_code >> 1) | ((lbp_code & 1) << (
    num_neighbors - 1))
290                 min_val = min(min_val, lbp_code)
291
292             # if code is 0, set lbp label to 0. If code uniform,
    set to number of bits that are 1. If nonuniform, set label to
    num_neighbors + 1. This is in accordance with the handout
293             if(min_val == 0):
294                 lbp_image[i, j] = 0
295             elif(check_if_code_uniform(min_val)):
296                 lbp_image[i, j] = bin(min_val).count('1')
297             else:
298                 lbp_image[i, j] = num_neighbors + 1
299
300     return lbp_histogram(lbp_image)
301
302 def complex_lbp_descriptor(image, radius=lbp_radius, num_neighbors=
    num_neighbors):
303     height, width = image.shape
304
305     lbp_image = np.zeros((height, width), dtype=np.uint8)
306

```

```

307 # creation of array of angles that each neighbor point makes
    with the central point
308 angles = [2 * np.pi * i / num_neighbors for i in range(
    num_neighbors)]
309
310 for i in range(radius, height - radius):
311     for j in range(radius, width - radius):
312         center = image[i, j]
313         lbp_code = 0
314
315         for idx, angle in enumerate(angles):
316             # collecting subpixel value of neighbor point
317             x = j + radius * np.cos(angle)
318             y = i - radius * np.sin(angle)
319
320             # bilinear interpolation to determine neighbor's
    subpixel value
321             neighbor = bilinear_interpolate(image, x, y)
322
323             # setting bits to 1 if their corresponding neighbor
    is greater
    than the center pixel
324             lbp_code |= (neighbor > center) << idx
325
326             min_val = lbp_code
327
328             # circularly shifting to produce the largest number of
    zeros on the left of the lbp code
329             for _ in range(num_neighbors):
330                 lbp_code = (lbp_code >> 1) | ((lbp_code & 1) << (
    num_neighbors - 1))
331                 min_val = min(min_val, lbp_code)
332
333             # if code is 0, set lbp label to 0. If code uniform,
    set to number of bits that are 1. If nonuniform, set label to
    num_neighbors + 1. This is in accordance with the handout
334             if(min_val == 0):
335                 lbp_image[i, j] = 0
336             elif(check_if_code_uniform(min_val)):
337                 lbp_image[i, j] = bin(min_val).count('1')
338             else:
339                 lbp_image[i, j] = num_neighbors + 1
340
341         return lbp_histogram(lbp_image)
342
343 def form_gram_matrix_vector(feature_tensor):
344     # Collect sizes of input tensor
345     batch_size = feature_tensor.shape[0]
346
347     C = int(feature_tensor.shape[1] / downsample_size)
348
349     # Setting to size int(C * (C + 1) / 2) due to only collecting
    upper triangular portion of gram matrix
350     gram_tensor = np.zeros((batch_size, int(downsample_size * (
    downsample_size + 1) / 2)))
351
352     # Go through each element in tensor and calculate gram matrix
    and turn to flattened upper triangular vector

```

```

353     for i in range(batch_size):
354         feature_vector = feature_tensor[i]
355
356         F_1 = feature_vector.reshape(feature_vector.shape[0],
feature_vector.shape[1] * feature_vector.shape[2])
357         G = F_1 @ F_1.T
358
359         G = block_reduce(G, block_size=(C, C), func=np.mean)
360
361         upper_triangular = G[np.triu_indices_from(G)]
362
363         gram_tensor[i] = upper_triangular.flatten()
364
365     return gram_tensor
366
367 def plot_gram_matrices(vgg_gram, resnet_coarse_gram,
resnet_fine_gram):
368     # This function only works if you input the 2D (nontensor)
versions of the gram matrices and is for visualization purposes
    only
369     fig, axis = plt.subplots(1, 3, figsize=(15, 5))
370
371     axis[0].imshow(vgg_gram, cmap='viridis')
372     axis[0].set_title("VGG Gram Matrix")
373
374     axis[1].imshow(resnet_coarse_gram, cmap='viridis')
375     axis[1].set_title("ResNet Coarse Gram Matrix")
376
377     axis[2].imshow(resnet_fine_gram, cmap='viridis')
378     axis[2].set_title("ResNet Fine Gram Matrix")
379
380     plt.show()
381
382 def load_images(directory):
383     images, labels, lbp_images = [], [], []
384
385     # Loop through all jpg files in given directory, extract the
label and append image and label to output lists
386     for filename in os.listdir(directory):
387         if filename.endswith(".jpg"):
388             img_path = os.path.join(directory, filename)
389             image = io.imread(img_path)
390
391             if len(image.shape) == 2: # Grayscale image
392                 image = np.stack((image,)*3, axis=-1) # Convert to
RGB by repeating the channel
393                 elif image.shape[2] == 4: # RGBA image
394                     image = image[:, :, :3] # Convert to RGB by
discarding the alpha channel
395
396                 image = transform.resize(image, (256, 256),
anti_aliasing=True, mode='reflect')
397                 lbp_image = transform.resize(image, (64, 64),
anti_aliasing=True, mode='reflect')
398
399                 images.append(image)
400                 lbp_images.append(lbp_image)

```

```

401         if filename.startswith("cloudy"):
402             labels.append(label_map["cloudy"])
403         elif filename.startswith("rain"):
404             labels.append(label_map["rain"])
405         elif filename.startswith("shine"):
406             labels.append(label_map["shine"])
407         else:
408             labels.append(label_map["sunrise"])
409
410     # Returning numpy arrays of output lists
411     return np.array(images), np.array(labels), np.array(lbp_images)
412
413 def create_lbp_tensor(images):
414     lbp_histograms = []
415
416     # Run lbp on each image and form a tensor for all images in set
417     for i in range(int(images.shape[0])):
418         lbp_histograms.append(lbp_descriptor(images[i]))
419         print("processed LBP for {} images out of {}".format(i + 1,
420             images.shape[0]))
421
422     return np.vstack(lbp_histograms)
423
424 def create_complex_lbp_tensor(images):
425     complex_lbp_histograms = []
426
427     for i in range(int(images.shape[0])):
428         image = rgb2hsv(images[i])
429         image_hue = (image[:, :, 0] / 360.0) * 255.0
430         image_sat = (image[:, :, 1]) * 255.0
431         image_val = (image[:, :, 2]) * 255.0
432
433         hue_hist = complex_lbp_descriptor(image_hue)
434         sat_hist = complex_lbp_descriptor(image_sat)
435         val_hist = complex_lbp_descriptor(image_val)
436
437         histogram = np.hstack([hue_hist, sat_hist, val_hist])
438         complex_lbp_histograms.append(histogram)
439         print("processed Complex LBP for {} images out of {}".
440             format(i + 1, images.shape[0]))
441
442     return np.vstack(complex_lbp_histograms)
443
444 def train_svm(train_features, train_labels, test_features,
445     test_labels, batch_size, namestring):
446
447     # Create the SVM classifier
448     if(namestring == "lbp"):
449         # Best performing classifier for LBP
450         classifier = svm.SVC(kernel='rbf', probability=True, C
451             =10000)
452     else:
453         # Best performing classifier for non-LBP
454         classifier = svm.SVC(kernel='linear', probability=True)

```

```

454 # Calculate the number of batches
455 num_batches = int(np.ceil(len(train_features) / batch_size))
456
457 # Train the SVM in batches with a progress bar
458 for i in tqdm(range(num_batches), desc='Training SVM'):
459     start_idx = i * batch_size
460     end_idx = min((i + 1) * batch_size, len(train_features))
461     X_batch = train_features[start_idx:end_idx]
462     y_batch = train_labels[start_idx:end_idx]
463
464     # Fit the model on the current batch
465     classifier.fit(X_batch, y_batch)
466
467 # Make predictions on the test set
468 y_pred = classifier.predict(test_features)
469
470 accuracy = accuracy_score(test_labels, y_pred)
471 print("Accuracy:", accuracy)
472
473 # Generate the confusion matrix
474 conf_matrix = confusion_matrix(test_labels, y_pred)
475 print("Confusion Matrix:")
476 print(conf_matrix)
477
478 # Visualization of confusion matrix
479 sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
480             cbar=False, xticklabels=reverse_label_map, yticklabels=
481             reverse_label_map)
482 plt.title("Confusion Matrix for {} With Accuracy: {}".format(
483     namestring, accuracy * 100))
484 plt.show()
485
486 return classifier
487
488 def plot_histogram(histogram):
489     plt.figure(figsize=(10, 5))
490     plt.bar(np.arange(len(histogram)), histogram, width=1,
491            edgecolor='black')
492     plt.xlabel('LBP Code')
493     plt.ylabel('Frequency')
494     plt.title("LBP Histogram")
495     plt.show()
496
497 def find_correct_and_incorrect_predictions(classifier,
498     test_feature_vector, test_images, test_labels):
499     correct_found = 0 # Setting to counter because all classifiers
500     were identifying same correct image. By collecting a large
501     amount of correct images (in this case 15) I can get some
502     variety
503     incorrect_found = False
504
505     for i in range(test_feature_vector.shape[0]):
506         y_pred = classifier.predict(test_feature_vector[i].reshape
507         (1, -1))[0]
508
509         if((y_pred != test_labels[i]) and not incorrect_found):

```

```

502         incorrect_found = True
503         print("Incorrect Image Found")
504         print("Predicted class: ", reverse_label_map[y_pred])
505         print("Actual class: ", reverse_label_map[test_labels[i
506     ])
507
508     plt.imshow(test_images[i])
509     plt.show()
510     if((y_pred == test_labels[i]) and correct_found < 15):
511         correct_found += 1
512         print("Correct Image Found")
513         print("Predicted class: ", reverse_label_map[y_pred])
514         print("Actual class: ", reverse_label_map[test_labels[i
515     ])
516
517     plt.imshow(test_images[i])
518     plt.show()
519     if(correct_found >= 15 and incorrect_found):
520         break
521
522 if __name__ == '__main__':
523     train_dir = "data/training"
524     test_dir = "data/testing"
525     encoder_name = "resnet50"
526
527     preprocessing = False
528
529     # Loading images
530     print("loading training images")
531     train_images, train_labels, train_lbp_images = load_images(
532         train_dir)
533     print("done loading training images, now loading testing")
534     test_images, test_labels, test_lbp_images = load_images(
535         test_dir)
536     print("done loading testing images")
537
538     if(preprocessing):
539         # ----- Pre processing -----#
540         # Saving labels so preprocessing only needs to be run once
541         torch.save(train_labels, "train_labels.pt")
542         torch.save(test_labels, "test_labels.pt")
543
544         # Encoding images with VGG19
545         vgg = VGG19()
546         vgg.load_weights("vgg_normalized.pth")
547         vgg.to(device)
548         vgg_feature_train = vgg.extract_features(train_images)
549         vgg_feature_test = vgg.extract_features(test_images)
550
551         # Encoding images with ResNet Coarse and Fine
552         resnet = CustomResNet(encoder=encoder_name)
553         resnet.to(device)
554         resnet_coarse_feature_train, resnet_fine_feature_train=
555         resnet(train_images)
556         resnet_coarse_feature_test, resnet_fine_feature_test=
557         resnet(test_images)

```



```

553     # Calculating gram matrices for vgg and resnet feature
        vectors
554         vgg_gram_train = form_gram_matrix_vector(vgg_feature_train)
555         resnet_coarse_gram_train = form_gram_matrix_vector(
resnet_coarse_feature_train)
556         resnet_fine_gram_train = form_gram_matrix_vector(
resnet_fine_feature_train)
557         vgg_gram_test = form_gram_matrix_vector(vgg_feature_test)
558         resnet_coarse_gram_test = form_gram_matrix_vector(
resnet_coarse_feature_test)
559         resnet_fine_gram_test = form_gram_matrix_vector(
resnet_fine_feature_test)
560
561
562     # Encoding images with LBP
563     lbp_feature_train = create_lbp_tensor(train_lbp_images)
564     lbp_feature_test = create_lbp_tensor(test_lbp_images)
565
566     # Custom texture extractor Complex LBP
567     complex_lbp_feature_train = create_complex_lbp_tensor(
train_lbp_images)
568     complex_lbp_feature_test = create_complex_lbp_tensor(
test_lbp_images)
569
570     print("VGG Gram Train Size: ", vgg_gram_train.shape)
571     print("VGG Gram Test Size: ", vgg_gram_test.shape)
572
573     print("ResNet Coarse Gram Train Size: ",
resnet_coarse_gram_train.shape)
574     print("ResNet Coarse Gram Test Size: ",
resnet_coarse_gram_test.shape)
575
576     print("ResNet Fine Gram Train Size: ",
resnet_fine_gram_train.shape)
577     print("ResNet Fine Gram Test Size: ", resnet_fine_gram_test
.shape)
578
579     print("LBP Feature Train Size: ", lbp_feature_train.shape)
580     print("LBP Feature Test Size: ", lbp_feature_test.shape)
581
582     print("Complex LBP Feature Train Size: ",
complex_lbp_feature_train.shape)
583     print("Complex LBP Feature Test Size: ",
complex_lbp_feature_test.shape)
584
585     # Saving all tensors so preprocessing only needs to be run
        once
586     torch.save(vgg_gram_train, 'vgg_train.pt', pickle_protocol
=4)
587     torch.save(resnet_coarse_gram_train, 'resnet_coarse_train.
pt', pickle_protocol=4)
588     torch.save(resnet_fine_gram_train, 'resnet_fine_train.pt',
pickle_protocol=4)
589     torch.save(lbp_feature_train, 'lbp_train.pt',
pickle_protocol=4)
590     torch.save(complex_lbp_feature_train, 'complex_lbp_train.pt
', pickle_protocol=4)

```

```

591         torch.save(vgg_gram_test, 'vgg_test.pt', pickle_protocol=4)
592         torch.save(resnet_coarse_gram_test, 'resnet_coarse_test.pt',
593             pickle_protocol=4)
594         torch.save(resnet_fine_gram_test, 'resnet_fine_test.pt',
595             pickle_protocol=4)
596         torch.save(lbp_feature_test, 'lbp_test.pt', pickle_protocol
597             =4)
598         torch.save(complex_lbp_feature_test, 'complex_lbp_test.pt',
599             pickle_protocol=4)
600
601         print("TENSORS SAVED")
602
603     else:
604         # Loading tensors from paths
605         train_labels = torch.load('train_labels.pt')
606         test_labels = torch.load('test_labels.pt')
607
608         vgg_gram_train = torch.load('vgg_train.pt')
609         resnet_coarse_gram_train = torch.load('resnet_coarse_train.
610 pt')
611         resnet_fine_gram_train = torch.load('resnet_fine_train.pt')
612         lbp_feature_train = torch.load('lbp_train.pt')
613         complex_lbp_feature_train = torch.load('complex_lbp_train.
614 pt')
615
616         vgg_gram_test = torch.load('vgg_test.pt')
617         resnet_coarse_gram_test = torch.load('resnet_coarse_test.pt
618 ')
619         resnet_fine_gram_test = torch.load('resnet_fine_test.pt')
620         lbp_feature_test = torch.load('lbp_test.pt')
621         complex_lbp_feature_test = torch.load('complex_lbp_test.pt'
622 )
623
624         # Printing size of everything to make sure tensors loaded
625         correctly
626         print("Train Labels Size: ", train_labels.shape)
627         print("Test Labels Size: ", test_labels.shape)
628
629         print("VGG Gram Train Size: ", vgg_gram_train.shape)
630         print("VGG Gram Test Size: ", vgg_gram_test.shape)
631
632         print("ResNet Coarse Gram Train Size: ",
633             resnet_coarse_gram_train.shape)
634         print("ResNet Coarse Gram Test Size: ",
635             resnet_coarse_gram_test.shape)
636
637         print("ResNet Fine Gram Train Size: ",
638             resnet_fine_gram_train.shape)
639         print("ResNet Fine Gram Test Size: ", resnet_fine_gram_test
640             .shape)
641
642         print("LBP Feature Train Size: ", lbp_feature_train.shape)
643         print("LBP Feature Test Size: ", lbp_feature_test.shape)
644
645         print("Complex LBP Feature Train Size: ",
646             complex_lbp_feature_train.shape)

```

```

634     print("Complex LBP Feature Test Size: ",
        complex_lbp_feature_test.shape)
635
636
637     # Normalizing input feature tensors
638     scaler = StandardScaler()
639     vgg_gram_train = scaler.fit_transform(vgg_gram_train)
640     vgg_gram_test = scaler.transform(vgg_gram_test)
641     resnet_coarse_gram_train = scaler.fit_transform(
        resnet_coarse_gram_train)
642     resnet_coarse_gram_test = scaler.transform(
        resnet_coarse_gram_test)
643     resnet_fine_gram_train = scaler.fit_transform(
        resnet_fine_gram_train)
644     resnet_fine_gram_test = scaler.transform(
        resnet_fine_gram_test)
645
646     # Setting batch size for training and feature descriptor
        type
647     batch_size = 1024
648     feature_type = "vgg" # "vgg", "resnet_coarse", "resnet_fine
        ", "lbp"
649
650
651     if(feature_type == "vgg"):
652         classifier = train_svm(vgg_gram_train, train_labels,
        vgg_gram_test, test_labels, batch_size, feature_type)
653         find_correct_and_incorrect_predictions(classifier,
        vgg_gram_test, test_images, test_labels)
654     elif(feature_type == "resnet_coarse"):
655         classifier = train_svm(resnet_coarse_gram_train,
        train_labels, resnet_coarse_gram_test, test_labels, batch_size,
        feature_type)
656         find_correct_and_incorrect_predictions(classifier,
        resnet_coarse_gram_test, test_images, test_labels)
657     elif(feature_type == "resnet_fine"):
658         classifier = train_svm(resnet_fine_gram_train,
        train_labels, resnet_fine_gram_test, test_labels, batch_size,
        feature_type)
659         find_correct_and_incorrect_predictions(classifier,
        resnet_fine_gram_test, test_images, test_labels)
660     elif(feature_type == "lbp"):
661         classifier_lbp = train_svm(lbp_feature_train,
        train_labels, lbp_feature_test, test_labels, batch_size,
        feature_type)
662         classifier_complex_lbp = train_svm(
        complex_lbp_feature_train, train_labels,
        complex_lbp_feature_test, test_labels, batch_size, feature_type
        )
663
664         find_correct_and_incorrect_predictions(classifier_lbp,
        lbp_feature_test, test_lbp_images, test_labels)
665
666     else:
667         print("Invalid feature type, please enter 'vgg', '
        resnet_coarse', 'resnet_fine', or 'lbp'")
668

```

669
670
671



Listing 1: Python Code