

ECE 661 Homework 7

Michael Goldberg

October 30 2024

1 Theoretical Question

1.1 Question 1

The reason this is a fundamental observation is because if we can sample the absolute conic Ω_∞ , we can construct ω , which is the image of the absolute conic. Because $\omega = K^{-T}K^{-1}$, where K is the intrinsic calibration matrix of the camera. Therefore, if ω can be found through sufficient number of sampling, then by proxy K can be determined and our camera's intrinsic parameters can be estimated.

1.2 Question 2

The algebraic relationship between Ω and ω can be realized through understanding the relationship between world3D points and their corresponding image within the camera. The projection of a 3D point (X, Y, Z) onto the image plane is:

$$x = K \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} \quad (1)$$

This can be compared with the general transformation of a conic C through a homography H with:

$$C' = H^{-T}CH^{-1} \quad (2)$$

Here, K is our homography, as it projects a world point to a camera point. Therefore,

$$\omega = K^{-T}\Omega_\infty K^{-1} \quad (3)$$

2 Programming Tasks

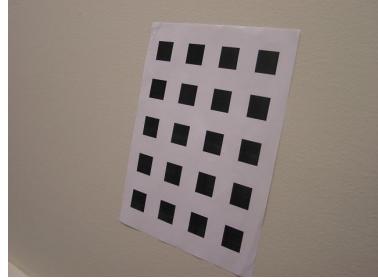
The following task was to use two datasets of a calibration board and calibrate the intrinsic and extrinsic parameters of the cameras that took the images. This was split into a series of steps:

1. Corner Detection

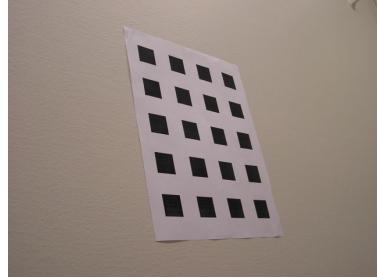
2. Zhang's Algorithm

3. Levenberg-Marquardt Nonlinear Optimization

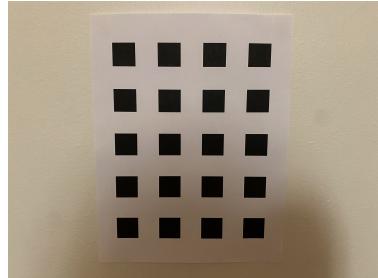
The following shows some example images of the calibration board from each dataset:



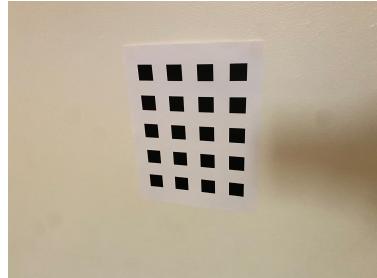
(a) Dataset 1 Image 1



(b) Dataset 1 Image 2



(c) Dataset 2 Image 1



(d) Dataset 2 Image 2

I will now go over how I addressed each task.

2.1 Corner Detection

This process was further broken down into a series of steps. The image was first converted to grayscale and Canny edge detection was performed. Then, Hough transform was performed on the Canny results to find lines that created the corners of the black squares. This process resulted in a lot of false positives, so a filtering process needed to be employed to ensure no duplicate lines or false positives were produced. Then, the intersections of each of these lines could be determined to find the corners. The lines also were sorted to ensure that corners were identified in the same order for any given image.

The most in-depth process was the line filtering and sorting. This was achieved in the following manner: First, the lines

The following shows the results of the Canny edge detection followed by the Hough Transform.

My best parameters for Canny edge detection were:

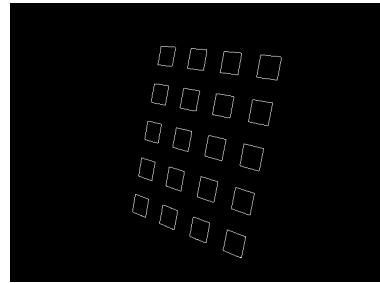
- Threshold1 = 200

- Threshold2 = 400

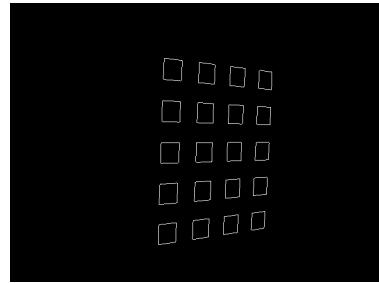
My best parameters for Hough transform were:

- Rho Threshold = 1
- Theta Threshold = $\pi/180$
- Hough Threshold = 50

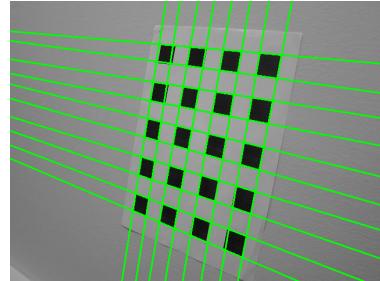
Dataset 1:



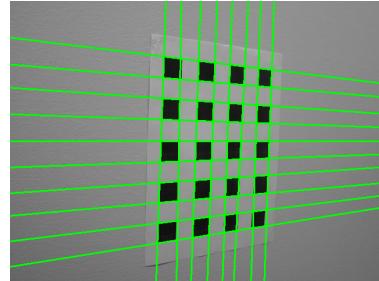
(a) Dataset 1 Edges 1



(b) Dataset 1 Edges 2

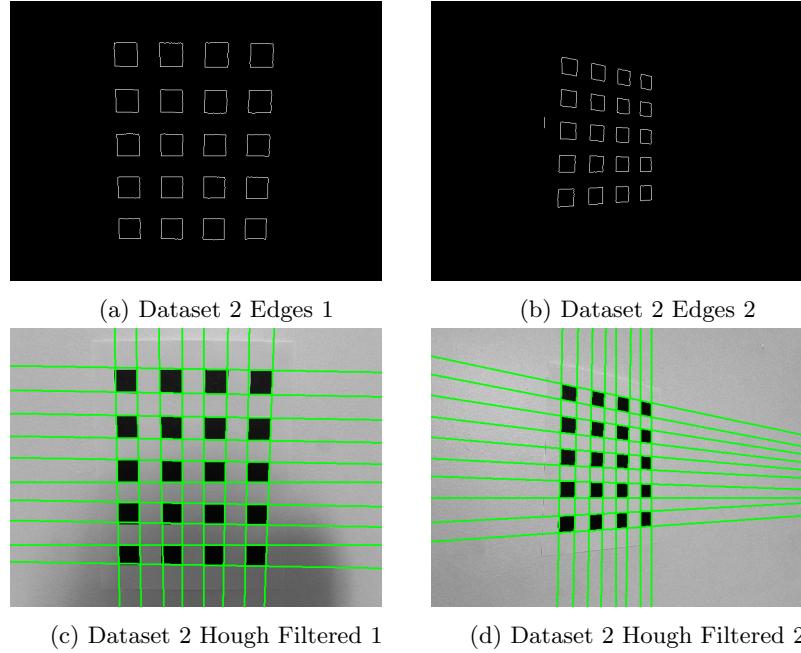


(c) Dataset 1 Hough Filtered 1



(d) Dataset 1 Hough Filtered 2

Dataset 2:



The Hough transform returns a list of lines given in the form (ρ, θ) , where ρ is the length of the perpendicular distance between the origin and the line that passes through a point of interest, and θ is the angle in radians of this perpendicular segment from the horizontal. The origin that this is measured from was observed to be from the image's origin, which is the top-left corner of the image. Therefore, to filter the lines, the lines were split into two groups - those that fell within a certain θ buffer = 50° around pure vertical (e.g. $abs(\theta) < deg2rad(50)$ or $abs(\theta - \pi) < deg2rad(50)$) and those that fell within the same buffer around pure horizontal (e.g. $abs(\theta - \pi/2) < deg2rad(50)$ or $abs(\theta - 3\pi/2) < deg2rad(50)$). Across these two groups, I could then check if any other given line within the group fell within a small buffer around the same value of ρ and θ . If these two values are too similar, then that means a slightly different line describing the same edge was picked up in the Hough transform and needs to be discarded. The buffer I used for this checking was 15 for ρ and 5° for θ . Now that we had discarded the duplicates, we were left with a single unique line for each edge.

For any pair of lines l_1 and l_2 in homogeneous coordinates, their intersection is given by

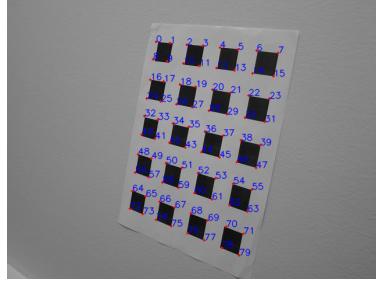
$$x_{intersection} = l_1 \times l_2 \quad (4)$$

This was done for all filtered lines, and any ideal points or points outside of the image bounds were discarded. This allowed me to find the corner points of the black squares on the calibration board. Now, all that was left was labelling the lines in a consistent manner, which could be achieved by sorting the lines in each group by increasing ρ . This resulted in line 0 of the horizontal group

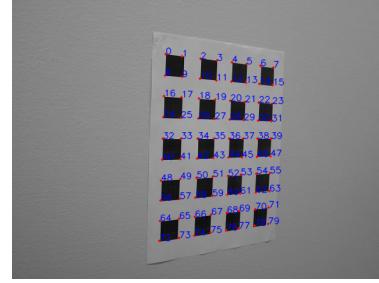
being the topmost line, and line 0 of the vertical group being the leftmost line. With a consistent sorting, the corners could be trivially labeled by which lines made up the intersection.

Below is the results of the corner detection for both datasets:

Dataset 1:

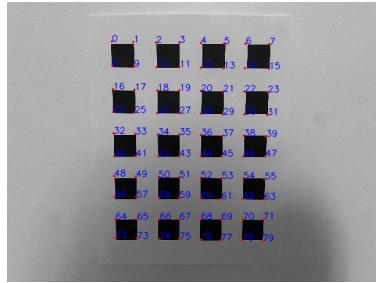


(a) Dataset 1 Detected Corners 1

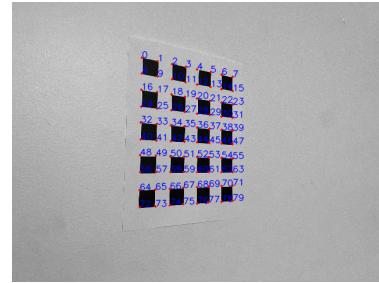


(b) Dataset 1 Detected Corners 2

Dataset 2:



(a) Dataset 2 Detected Corners 1



(b) Dataset 2 Detected Corners 2

Now, Zhang's calibration algorithm can be implemented.

2.2 Zhang's Algorithm Implementation

There are two main steps to Zhang's Camera Calibration Algorithm:

1. Intrinsic Calibration
2. Extrinsic Calibration

We begin with intrinsic calibration. We assume that the calibration pattern is on the world XY plane, which allows us to construct the following homography relationship:

$$x_{cam} = K[R|t] \begin{bmatrix} X \\ Y \\ 0 \\ 1 \end{bmatrix} \quad (5)$$

where we can express $K[R|t]$ as H .

The homography can be estimated by understanding where the detected corners lie in the world coordinate frame and using SVD with point-to-point correspondencies. I estimated the size of the black squares to be 25mm along each axis, as well as being separated by the same distance. I also assumed the top left corner (corner 0 labeled above) to be at world coordinates $X = 0$ and $Y = 0$. With these assumptions, I had the coordinates of all of the corners both in the world coordinate frame as well as the camera image. The following relationship could then be established:

$$Ah = x \quad (6)$$

where

$$A = \begin{bmatrix} -X & -Y & -1 & 0 & 0 & 0 & xX & xY & x \\ 0 & 0 & 0 & -X & -Y & -1 & yX & yY & y \end{bmatrix} \quad (7)$$

, where X and Y are the world coordinates of a point and x and y are the image points. This matrix A can be populated with all 80 corner correspondencies, leading to 160 rows and 9 columns. This can be solved using SVD to find H .

Once H is established, we want to estimate the intrinsic parameters using the relationship that $\omega = K^{-T}K^{-1}$. This can be done using the homography we found for the image as well as the sampled circular points where we know the following to be true:

$$V_{01}^T b = 0 \quad (8)$$

and

$$(V_{00} - V_{11})^T b = 0 \quad (9)$$

, where

$$V_{i,j} = \begin{bmatrix} h_{i,0}h_{j,0} \\ h_{i,0}h_{j,1} + h_{i,1}h_{j,0} \\ h_{i,1}h_{j,1} \\ h_{i,2}h_{j,0} + h_{i,0}h_{j,2} \\ h_{i,2}h_{j,1} + h_{i,1}h_{j,2} \\ h_{i,2}h_{j,2} \end{bmatrix} \quad (10)$$

and

$$b = \begin{bmatrix} \omega_{00} \\ \omega_{01} \\ \omega_{11} \\ \omega_{02} \\ \omega_{12} \\ \omega_{22} \end{bmatrix} \quad (11)$$

By calculating $V_{01}^T b$ and $(V_{00} - V_{11})^T b$ for each calibration image, we can

construct the overall relationship:

$$\begin{bmatrix} V_1 \\ V_2 \\ \vdots \\ \vdots \\ V_{n-1} \\ V_n \end{bmatrix} b = 0 \quad (12)$$

We can then use SVD to solve for b , which gives us ω . We now need to solve for K . This can be done through the following process:

$$x_0 = \frac{\omega_{01}\omega_{02} - \omega_{00}\omega_{12}}{\omega_{00}\omega_{11} - \omega_{01}^2} \quad (13)$$

$$\lambda = \omega_{22} - \frac{\omega_{02}^2 + x_0(\omega_{01}\omega_{02} - \omega_{00}\omega_{12})}{\omega_{00}} \quad (14)$$

$$\alpha_x = \sqrt{\frac{\lambda}{\omega_{00}}} \quad (15)$$

$$\alpha_x = \sqrt{\frac{\lambda\omega_{00}}{\omega_{00}\omega_{11} - \omega_{01}^2}} \quad (16)$$

$$s = -\frac{\omega_{01}\alpha_x^2\alpha_y}{\lambda} \quad (17)$$

$$y_0 = \frac{s x_0}{\alpha_y} - \frac{\omega_{02}\alpha_x^2}{\lambda} \quad (18)$$

Through these calculations, we can construct the intrinsic matrix, as its form is:

$$K = \begin{bmatrix} \alpha_x & s & x_0 \\ 0 & \alpha_y & y_0 \\ 0 & 0 & 1 \end{bmatrix} \quad (19)$$

Through this process, I achieved the following intrinsic matrices for the two cameras:

Dataset 1:

$$K_1 = \begin{bmatrix} 725.968 & 5.492 & 240.793 \\ 0 & 721.778 & 319.591 \\ 0 & 0 & 1 \end{bmatrix} \quad (20)$$

Dataset 2:

$$K_2 = \begin{bmatrix} 484.402 & 3.856 & 229.551 \\ 0 & 487.079 & 309.5772 \\ 0 & 0 & 1 \end{bmatrix} \quad (21)$$

We can now calculate the extrinsic parameters. This is done through the following process: It can be seen from Equation 5 that

$$[r1, r2, t] = K^{-1}H \quad (22)$$

However, we need to implement a normalization term in order to preserve homogeneity. Therefore, solving for $r1$, $r2$ and t can be done through:

$$r1 = \zeta K^{-1}h1 \quad (23)$$

$$r2 = \zeta K^{-1}h2 \quad (24)$$

$$t = \zeta K^{-1}h3 \quad (25)$$

, where $\zeta = \frac{1}{\|K^{-1}h1\|}$

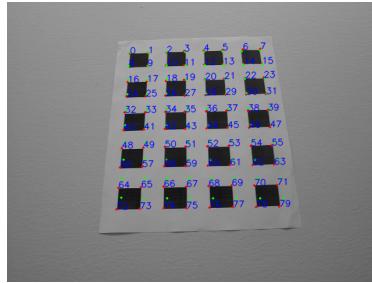
The final parameter $r3$ can be found through the cross product of $r1$ and $r2$, yielding $r3 = r1 \times r2$.

This then provides us with our three rotation vectors and our single translation vector that estimates the extrinsic parameters of the image. The final extrinsic matrix can be constructed through:

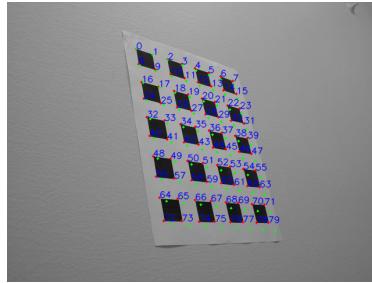
$$[R|t] = [r1 \ r2 \ r3 \ t] \quad (26)$$

The initial results of this can be observed through reprojection of the world coordinate points back onto the image and comparing the locations of the originally found corners vs. the reprojected corners. The results of this can be seen across both datasets below:

Dataset 1:

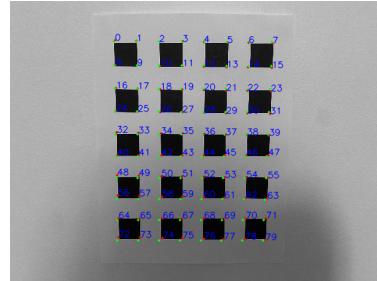


(a) Dataset 1 Reprojected Corners 1

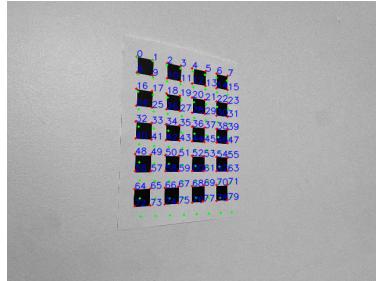


(b) Dataset 1 Reprojected Corners 2

Dataset 2:



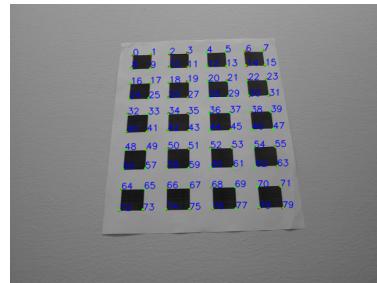
(a) Dataset 2 Reprojected Corners 1



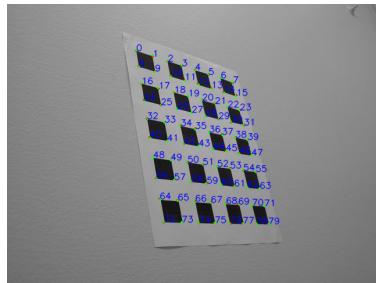
(b) Dataset 2 Reprojected Corners 2

It can be seen at large angles, the extrinsic parameters break down a bit and are not that accurate. I therefore applied Levenberg-Marquadt (LM) nonlinear least squares to minimize the reprojection error and therefore optimize the extrinsic calibration. The results of this can be seen below:

Dataset 1:



(a) Dataset 1 Reprojected Corners after LM 1



(b) Dataset 1 Reprojected Corners after LM 2

The final estimated extrinsic matrices for these two images from Dataset 1 were:

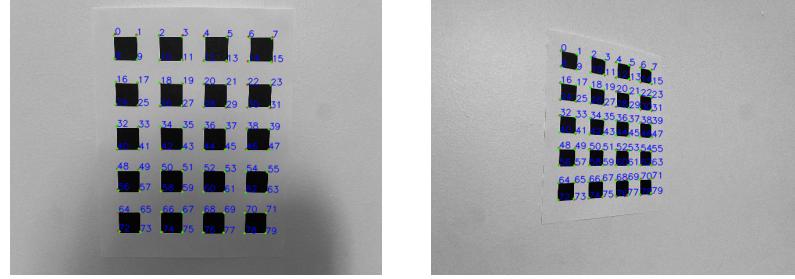
$$[R|t]_{dataset1,image1} = \begin{bmatrix} 1.061 & -0.247 & 0.608 & 20.018 \\ 0.389 & 1.429 & 0.039 & -275.63 \\ -0.829 & 0.119 & 0.793 & 803.36 \end{bmatrix} \quad (27)$$

and

$$[R|t]_{dataset1,image2} = \begin{bmatrix} 807.17 & -27.38 & -0.526 & 1.482e4 \\ -56.97 & 894.133 & -0.079 & -1.39e5 \\ 499.617 & 69.128 & 0.847 & 4.486e5 \end{bmatrix} \quad (28)$$

For image 1 the mean reprojection error (Euclidean distance) was 0.91 and the variance was 0.2. For image 2 the mean reprojection error was 0.89 and the variance was 0.45.

Dataset 2:



(a) Dataset 2 Reprojected Corners after LM 1 (b) Dataset 2 Reprojected Corners after LM 2

The final extrinsic matrices for these two images from Dataset 2 were:

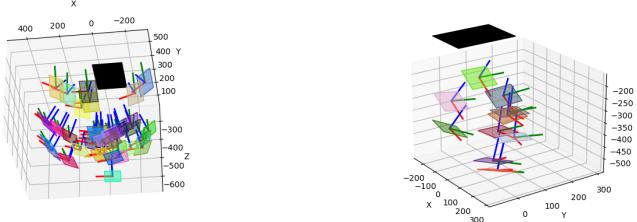
$$[R|t]_{dataset2,image1} = \begin{bmatrix} 1.04e2 & 1.18 & -5.82e-3 & -3.29e3 \\ 1.36 & 1.02e2 & -1.2e-1 & -1.6e4 \\ 8.97e-1 & 1.27e1 & 9.93e-1 & 3.25e4 \end{bmatrix} \quad (29)$$

and

$$[R|t]_{dataset2,image2} = \begin{bmatrix} 2.9 & -9.23e-2 & -4.8e-1 & -8.18 \\ -7.1e-2 & 3.1 & 5.18e-2 & -5.78e2 \\ 1.58 & -2.26e-1 & 8.77e-1 & 1.32e3 \end{bmatrix} \quad (30)$$

For image 1 the mean reprojection error (Euclidean distance) was 1.17 and the variance was 0.37. For image 2 the mean reprojection error was 0.78 and the variance was 0.14.

The camera poses could then be plotted in 3D to see all of the different positions and orientations of the camera in relation to the calibration board - this serves as yet another visual check that the extrinsic calibration is correct. This can be seen below:



(a) Dataset 1 Camera Poses

(b) Dataset 2 Camera Poses

3 Code

```

1 import cv2
2 import numpy as np
3 import time
4 import random
5 import matplotlib.pyplot as plt
6 from mpl_toolkits.mplot3d import Axes3D
7 from mpl_toolkits.mplot3d.art3d import Poly3DCollection
8 from scipy.optimize import least_squares
9
10
11 # Canny edge detection parameters
12 canny_threshold1 = 200
13 canny_threshold2 = 400
14
15 # Hough transform parameters
16 rho_threshold = 1
17 theta_threshold = np.pi / 180
18 hough_threshold = 50
19
20 # Calibration pattern parameters
21 square_distance = 25
22 pattern = [5, 4]
23
24 visualize_reprojection = True
25 visualize_poses = True
26
27 perform_LM = True
28
29
30 def HC(point):
31     # Helper function to convert a point to homogeneous coordinates
32     return np.array([point[0], point[1], 1])
33
34 def polar_to_cartesian(rho, theta, length=1500):
35     # Converting from polar coordinates to cartesian coordinates to
36     # get arbitrary endpoints of the line
37     x0 = rho * np.cos(theta)
38     y0 = rho * np.sin(theta)
39     x1 = int(x0 + length * (-np.sin(theta)))
40     y1 = int(y0 + length * (np.cos(theta)))
41     x2 = int(x0 - length * (-np.sin(theta)))
42     y2 = int(y0 - length * (np.cos(theta)))
43
44     return (x1, y1, x2, y2)
45
46 def detect_corners(image, visualize_corner_detection=False):
47
48     # Performing canny edge detection
49     edges = cv2.Canny(image, canny_threshold1, canny_threshold2)
50
51     detected_edge_image = cv2.cvtColor(image.copy(), cv2.COLOR_GRAY2BGR)
52     detected_corners_image = detected_edge_image.copy()
53
54     image_height, image_width = image.shape
55
56     # Performing hough transform to extract lines

```

```

56     lines = cv2.HoughLines(edges, rho_threshold, theta_threshold,
57                             hough_threshold)
58
59     # Filtering out duplicate lines for the same edge that deviate
60     # by small angles
61     rho_filter = 15
62     theta_filter = np.deg2rad(5)
63
64     filtered_lines = []
65
66     # Sorting lines by rho
67     if lines is not None:
68         lines = sorted(lines, key=lambda x: x[0][0])
69
70     # Checking if both rho and theta differences are large to
71     # qualify as a unique line
72     for line in lines:
73         # collect individual line's rho and theta values
74         rho, theta = line[0]
75         is_unique = True
76
77         for unique_rho, unique_theta, _ in filtered_lines:
78             # Checking if rho falls within a threshold of any
79             # previous rho as well as theta falling within a theta threshold
80             # of any previous theta to check if lines are duplicates. Check
81             # is complicated due to bounds of rho and theta for Hough
82             # First check is for positive rho, and therefore angle
83             # differences close to zero. If rho and unique_rho are opposite
84             # in sign, angle will be close to pi if lines are duplicates
85             if (abs(rho - unique_rho) < rho_filter and abs(theta -
86                 unique_theta) < theta_filter) or (abs(abs(rho) - abs(unique_rho))
87                 ) < rho_filter and abs(theta - unique_theta) > (np.pi -
88                 theta_filter):
89                 is_unique = False
90                 break
91
92         # If unique flag still true, line passed checking against
93         # all other lines and is in fact not a duplicate
94         if is_unique:
95             x1, y1, x2, y2 = polar_to_cartesian(rho, theta)
96             line_cartesian = (x1, y1, x2, y2)
97
98             filtered_lines.append((rho, theta, line_cartesian))
99
100
101
102     # Sorting lines by rho and vertical/horizontal to provide
103     # consistent labelling for each eventually found corner
104     vertical_lines = []
105     horizontal_lines = []
106
107     for line in filtered_lines:
108         rho, theta, line_cartesian = line
109         if(abs(theta) < np.deg2rad(50) or abs(theta - np.pi) < np.
110             deg2rad(50)):
111             vertical_lines.append(line)

```

```

98     elif(abs(theta - np.pi/2) < np.deg2rad(50) or abs(theta -
99         3*np.pi/2) < np.deg2rad(50)):
100         horizontal_lines.append(line)
101
102     # Sorting by abs for vertical lines only as they were more
103     # prone to angle switching
104     vertical_lines.sort(key=lambda x: abs(x[0]))
105     horizontal_lines.sort(key=lambda x: x[0])
106
107
108     # Rearranging filtered_lines such that horizontal are always
109     # considered first - helps with consistent labelling
110     filtered_lines = horizontal_lines + vertical_lines
111
112
113     # Finding corners using cross products in HC
114     corners = []
115
116     for i in range(len(filtered_lines)):
117         line_cartesian_test = filtered_lines[i][2]
118
119         endpoint1_hc_test = HC((line_cartesian_test[0],
120         line_cartesian_test[1]))
121         endpoint2_hc_test = HC((line_cartesian_test[2],
122         line_cartesian_test[3]))
123
124         line_hc_test = np.cross(endpoint1_hc_test,
125         endpoint2_hc_test)
126
127         for j in range(len(filtered_lines)):
128             if(i != j):
129                 # If not the same line, extract both line's
130                 # homogeneous coordinates and calculate intersection point
131                 line_cartesian = filtered_lines[j][2]
132
133                 endpoint1_hc = HC((line_cartesian[0],
134                 line_cartesian[1]))
135                 endpoint2_hc = HC((line_cartesian[2],
136                 line_cartesian[3]))
137
138                 line_hc = np.cross(endpoint1_hc, endpoint2_hc)
139
140                 intersection = np.cross(line_hc_test, line_hc)
141
142                 is_valid_corner = True
143                 is_unique_corner = True
144
145                 # Checking if ideal point, if it is, flag as
146                 # invalid as these do not reside in the image
147                 if(intersection[2] != 0):
148                     intersection = intersection / intersection[2]
149                     if(intersection[0] < 0 or intersection[1] < 0
150                     or intersection[1] > image_height or intersection[0] >
151                     image_width):
152                         is_valid_corner = False # Intersection out
153                         of the bounds of the image
154                     else:

```



```

188
189     A = np.array(A)
190
191     _, _, Vt = np.linalg.svd(A) # Using SVD to solve for homography
192     terms
193     H = Vt[-1].reshape((3, 3))
194
195     return H / H[2, 2]
196
197 def find_V(H, i, j):
198     # Creating indexed v-vector for intrinsic calibration in
199     # accordance with Zhang paper
200     return np.array([
201         H[0][i] * H[0][j],
202         H[0][i] * H[1][j] + H[1][i] * H[0][j],
203         H[1][i] * H[1][j],
204         H[2][i] * H[0][j] + H[0][i] * H[2][j],
205         H[2][i] * H[1][j] + H[1][i] * H[2][j],
206         H[2][i] * H[2][j]
207     ])
208
209 def intrinsic_calibration(homographies):
210     # Creating Vb = 0 relationship and solving for b using SVD
211     V = []
212     for H in homographies:
213         V.append(find_V(H, 0, 1))
214         V.append(find_V(H, 0, 0) - find_V(H, 1, 1))
215
216     V = np.array(V)
217     _, _, Vt = np.linalg.svd(V)
218     b = Vt[-1]
219
220     # Extracting omega from b
221     w11, w12, w22, w13, w23, w33 = b
222
223     # Calculating intrinsic parameters from image of absolute conic
224     x0 = (w12 * w13 - w11 * w23) / (w11 * w22 - w12**2)
225     lambda_ = w33 - (w13**2 + x0*(w12 * w13 - w11 * w23)) / w11
226     alpha_x = np.sqrt(lambda_ / w11)
227     alpha_y = np.sqrt((lambda_ * w11) / (w11 * w22 - w12**2))
228     s = -(w12 * alpha_x**2 * alpha_y) / lambda_
229     y0 = s*x0 / alpha_y - (w13 * alpha_x**2) / lambda_
230
231     # Creation of intrinsic calibration matrix from intrinsic
232     # parameters
233     K = np.array([
234         [alpha_x, s, x0],
235         [0, alpha_y, y0],
236         [0, 0, 1]
237     ])
238
239     return K
240
241 def condition_R(R):
242     U, _, Vt = np.linalg.svd(R)
243     R_conditioned = np.dot(U, Vt)

```

```

242     # If determinant is negative, negate last column to create a
243     # proper determinant of 1
244     if np.linalg.det(R_conditioned) < 0:
245         U[:, -1] *= -1
246         R_conditioned = U @ Vt
247
248     # Normalization of the columns of the orthogonal matrix to
249     # produce an orthonormal matrix
250     R_conditioned[:, 0] /= np.linalg.norm(R_conditioned[:, 0])
251     R_conditioned[:, 1] /= np.linalg.norm(R_conditioned[:, 1])
252     R_conditioned[:, 2] /= np.linalg.norm(R_conditioned[:, 2])
253
254     return R_conditioned
255
256 def reproject_points(world_corners, K, Rt):
257     # Applying extrinsic calibration matrix to world points and
258     # extracting image location
259     projected_points = []
260     for x in world_corners:
261         x_hc = np.array([x[0], x[1], x[2], 1])
262         projected_point = K @ Rt @ x_hc
263         projected_point /= projected_point[2]
264         projected_points.append((projected_point[0],
265                                   projected_point[1]))
265
266     return np.array(projected_points)
267
268 def reprojection_error(params, K, world_corners, image_corners):
269     # Helper function that defines reprojection error function used
270     # for LM minimization
271     R = params[:9].reshape(3, 3)
272     t = params[9:].reshape(3, 1)
273
274     projected_points = reproject_points(world_corners, K, np.hstack
275                                         ((R, t)))
276
277     error = (projected_points - image_corners).ravel()
278
279     return error
280
281 def reprojection_visualization(homographies, corners_list,
282                                 extrinsics, optimized_extrinsics, images):
283     # Labels corners, before_lm reprojected corners, and after LM
284     # reprojected corners for visualization on how well the extrinsic
285     # parameters are both before and after LM
286     for j in range(len(homographies)):
287         corners = corners_list[j]
288         extrinsic = extrinsics[j]
289         homography = homographies[j]
290         optimized_extrinsic = optimized_extrinsics[j]
291         detected_corners_image = cv2.cvtColor(images[j], cv2.
292 COLOR_GRAY2BGR)
293         optimized_corners_image = detected_corners_image.copy()
294
295         print("Extrinsic matrix: ", optimized_extrinsic)
296
297         # Using Euclidean distance to calculate reprojection error

```

```

289     projected_points = reproject_points(world_corners, K,
290                                         optimized_extrinsic)
291
292     reprojection_error = []
293
294     for c, corner in enumerate(corners):
295         distance = np.sqrt((corner[0] - projected_points[c][0])
296                             **2 + (corner[1] - projected_points[c][1])**2)
297         reprojection_error.append(distance)
298
299     mean_rproj_error = np.mean(reprojection_error)
300     var_rproj_error = np.var(reprojection_error)
301
302     print("Mean of reprojection error: ", mean_rproj_error)
303     print("Variance of reprojection error: ", var_rproj_error)
304
305     # Reprojecting corners with both non-LM and LM
306     reprojected_corners_initial = reproject_points(
307         world_corners, K, extrinsic)
308     reprojected_corners_optimized = reproject_points(
309         world_corners, K, optimized_extrinsic)
310
311     for i, corner in enumerate(corners):
312         corner_location = (int(corner[0]), int(corner[1]))
313         reprojected_corner_initial = (int(
314             reprojected_corners_initial[i][0]), int(
315             reprojected_corners_initial[i][1]))
316         reprojected_corner_optimized = (int(
317             reprojected_corners_optimized[i][0]), int(
318             reprojected_corners_optimized[i][1]))
319
320         cv2.circle(detected_corners_image, corner_location, 2,
321                    (0, 0, 255), -1)
322         cv2.circle(detected_corners_image,
323                    reprojected_corner_initial, 2, (0, 255, 0), -1)
324
325         cv2.circle(optimized_corners_image, corner_location, 2,
326                    (0, 0, 255), -1)
327         cv2.circle(optimized_corners_image,
328                    reprojected_corner_optimized, 2, (0, 255, 0), -1)
329
330         cv2.putText(detected_corners_image, "{}".format(i),
331                     corner_location, cv2.FONT_HERSHEY_SIMPLEX, 0.5, (255, 0, 0), 1,
332                     cv2.LINE_AA)
333         cv2.putText(optimized_corners_image, "{}".format(i),
334                     corner_location, cv2.FONT_HERSHEY_SIMPLEX, 0.5, (255, 0, 0), 1,
335                     cv2.LINE_AA)
336
337     cv2.imshow("corners", detected_corners_image)
338     cv2.imshow("corners optimized", optimized_corners_image)
339     cv2.waitKey(0)
340     cv2.destroyAllWindows()
341
342
343     world_corners = generate_world_corners()
344     homographies = []
345     corners_list = []

```

```

330 extrinsics = []
331 optimized_extrinsics = []
332 images = []
333
334 # Load in images and estimate homography
335 for i in range(1, 41):
336     image = cv2.imread('HW8-Files/Dataset1/Pic_' + str(i) + '.jpg',
337                         cv2.IMREAD_GRAYSCALE)
338     # image = cv2.resize(cv2.imread('HW8-Files/Dataset2/' + str(i)
339     # + '.jpg', cv2.IMREAD_GRAYSCALE), (640, 480))
340
341     corners = detect_corners(image)
342
343     # Ensuring to discard images that don't have all 80 corners
344     # detected
345     if len(corners) != len(world_corners)):
346         continue
347
348     H = find_homography(world_corners, corners)
349     homographies.append(H)
350     corners_list.append(corners)
351     images.append(image)
352
353     # Calculate intrinsics
354     K = intrinsic_calibration(homographies)
355
356     print("Intrinsic Matrix K: ", K)
357
358     # Calculating extrinsic parameters and thereby extrinsic matrices
359     # for each image
360     for i in range(len(homographies)):
361         H = homographies[i]
362         h1, h2, h3 = H[:, 0], H[:, 1], H[:, 2]
363         zeta = 1 / np.linalg.norm(np.dot(np.linalg.inv(K), h1))
364         r1 = zeta * np.dot(np.linalg.inv(K), h1)
365         r2 = zeta * np.dot(np.linalg.inv(K), h2)
366         r3 = np.cross(r1, r2)
367         t = zeta * np.dot(np.linalg.inv(K), h3)
368
369         R = np.column_stack((r1, r2, r3))
370
371         R = condition_R(R)
372
373         Rt = np.column_stack((R, t))
374         extrinsics.append(Rt)
375
376         # print("Original extrinsic calibration matrix is: ", Rt)
377
378         initial_guess = np.hstack((R.flatten(), t))
379
380         corners_array = np.array(corners_list[i])
381         world_corners_array = np.array(world_corners)
382
383         # Taking the reprojection error function as input to LM and
384         # trying to minimize it using initial extrinsic calibration as
385         # initial guess
386         if (perform_LM):

```

```

381     lm_result = least_squares(reprojection_error, initial_guess
382     , args=(K, world_corners, corners_array), method="lm")
383
383     optimized_params = lm_result.x
384     optimized_R = optimized_params[:9].reshape(3, 3)
385     optimized_t = optimized_params[9:].reshape(3, 1)
386
387     optimized_Rt = np.hstack((optimized_R, optimized_t))
388 else:
389     optimized_Rt = Rt
390
391 optimized_extrinsics.append(optimized_Rt)
392
393 # Visualization of the reprojection error
394 if visualize_reprojection:
395     reprojection_visualization(homographies, corners_list,
396         extrinsics, optimized_extrinsics, images)
396 if visualize_poses:
397     # Parameters for camera pose visualization
398     camera_axes = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]])
399     scaling_factor = 100
400     plane_size = scaling_factor
401     plane_points = np.array([[-plane_size/2, -plane_size/2, 0],
402         [plane_size/2, -plane_size/2, 0], [plane_size/2, plane_size/2,
403         0], [-plane_size/2, plane_size/2, 0]])
402     calibration_board_points = np.array([-[-scaling_factor, -
403         scaling_factor, 0], [scaling_factor, -scaling_factor, 0],
404         [scaling_factor, scaling_factor, 0], [-scaling_factor,
405         scaling_factor, 0]])
403
404 fig = plt.figure()
405 ax = fig.add_subplot(111, projection='3d')
406
407 for i, Rt in enumerate(extrinsics):
408     R = Rt[:, :3]
409     t = Rt[:, 3:]
410
411     # Extracting the camera axes
412     X_cam_x = camera_axes[0].T.reshape((3, 1))
413     X_cam_y = camera_axes[1].T.reshape((3, 1))
414     X_cam_z = camera_axes[2].T.reshape((3, 1))
415
416     # Calculation of camera center
417     camera_center = -R.T @ t
418
419     # Applying extrinsics to find world coordinates of camera
420     # axes
420     X_world_x = R.T @ X_cam_x + camera_center
421     X_world_y = R.T @ X_cam_y + camera_center
422     X_world_z = R.T @ X_cam_z + camera_center
423
424     # Applying extrinsics to plane vertices to find where
425     # camera principal plane lies in world3D
425     world_plane_points = (R.T @ plane_points.T + camera_center)
426     .T
427
427     # Plotting camera center

```

```

428     ax.scatter(camera_center[0], camera_center[1],
429                 camera_center[2], color='r', s=1, label="Camera {}".format(i))
430
431     # Normalization and scaling the camera axes
432     direction_x = X_world_x - camera_center
433     direction_y = X_world_y - camera_center
434     direction_z = X_world_z - camera_center
435
436     direction_x_norm = (direction_x / np.linalg.norm(
437         direction_x)) * scaling_factor
438     direction_y_norm = (direction_y / np.linalg.norm(
439         direction_y)) * scaling_factor
440     direction_z_norm = (direction_z / np.linalg.norm(
441         direction_z)) * scaling_factor
442
443     # Plotting camera axes
444     ax.plot3D(
445         [camera_center[0], camera_center[0] + direction_x_norm
446          [0]],
447         [camera_center[1], camera_center[1] + direction_x_norm
448          [1]],
449         [camera_center[2], camera_center[2] + direction_x_norm
450          [2]], color='r', linewidth=2
451     )
452
453     ax.plot3D(
454         [camera_center[0], camera_center[0] + direction_y_norm
455          [0]],
456         [camera_center[1], camera_center[1] + direction_y_norm
457          [1]],
458         [camera_center[2], camera_center[2] + direction_y_norm
459          [2]], color='g', linewidth=2
460     )
461
462     ax.plot3D(
463         [camera_center[0], camera_center[0] + direction_z_norm
464          [0]],
465         [camera_center[1], camera_center[1] + direction_z_norm
466          [1]],
467         [camera_center[2], camera_center[2] + direction_z_norm
468          [2]], color='b', linewidth=2
469     )
470
471     # Plotting camera principal plane as rectangle
472     plane_color = (random.randint(0, 255) / 255.0, random.
473     randint(0, 255) / 255.0, random.randint(0, 255) / 255.0)
474     vertices = [world_plane_points]
475
476     plane = Poly3DCollection(vertices, facecolors=plane_color,
477                               linewidths=1, edgecolors=plane_color, alpha=0.5)
478     ax.add_collection3d(plane)
479
480     # Plotting calibration board
481     calibration_board_color = (0, 0, 0)
482     calibration_board_vertices = [calibration_board_points]
483     calibration_board = Poly3DCollection(calibration_board_vertices
484                                         , facecolors=calibration_board_color, linewidths=1, edgecolors=

```

```
469     'black', alpha=1)
470     ax.add_collection3d(calibration_board)
471
472     ax.set_xlabel('X')
473     ax.set_ylabel('Y')
474     ax.set_zlabel('Z')
475
476     plt.show()
```

Listing 1: Python Code