

# ECE 661 Homework 9

Michael Goldberg

November 25 2024

## 1 Theoretical Question

### 1.1 Question 1

Let  $\vec{l}$  be the epipolar line in the left camera and  $\vec{l}'$  be the epipolar line in the right camera. Let's explore how an image line backprojects into world coordinates. We know that

$$\vec{x}^T \vec{l} = 0 \quad (1)$$

and

$$\vec{x} = P \vec{X} = \vec{X}^T P^T \quad (2)$$

We can therefore rewrite Equation 1 as:

$$\vec{X}^T P^T \vec{l} = 0 \quad (3)$$

This can be further examined to see that  $P^T \vec{l}$  forms a plane  $\vec{\pi}$  that point  $\vec{x}$  must lie on.

However, we know that the cameras are imaging the same point and their center of projections also exist in world3D. This describes this exact plane  $\pi$  that we just derived, but this plane also pierces camera  $P'$ . Therefore, we know that both cameras must share the same plane, and by extension  $\vec{x}$  must lie on its corresponding line  $\vec{l}'$ , as that line and point must also backproject to the same plane.

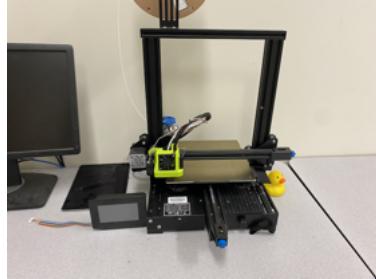
## 2 Programming Tasks

The following task was to perform stereo rectification on a pair of images to result in an eventual pointcloud of the object being imaged. This was broken down into the following steps:

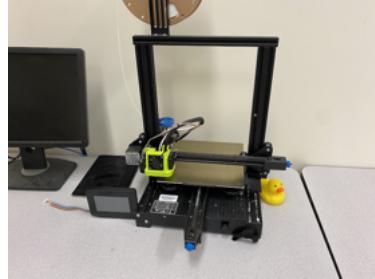
1. 8-Point Correspondence Estimation of F
2. Epipole Estimation
3. Camera Matrix Estimation

4. Refinement of  $P$  and  $P'$
5. Refinement of  $F$
6. Estimation of  $H$  and  $H'$
7. Interest Point Detection
8. Projective Reconstruction

The following shows the two images used:



(a) Left Stereo Image

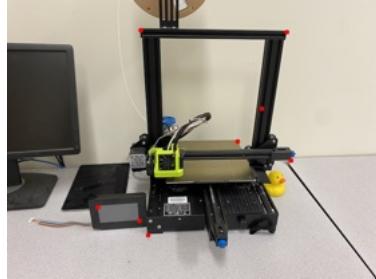


(b) Right Stereo Image

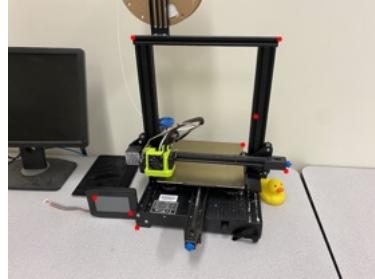
I will now go over how I addressed each task.

## 2.1 8-Point Correspondence Estimation of $F$

I first established 8 point-to-point correspondences between the left and right images, which can be seen here:



(a) Left Stereo Image Points



(b) Right Stereo Image Points

The first step is to normalize the image coordinates. From here on, I will refer to the left image points as  $x$  and the right image points as  $x'$ . This normalization matrix is:

$$T = \begin{bmatrix} s & 0 & -s\bar{x} \\ 0 & s & -s\bar{y} \\ 0 & 0 & 1 \end{bmatrix} \quad (4)$$

and

$$T' = \begin{bmatrix} s' & 0 & -s'\bar{x}' \\ 0 & s' & -s'\bar{y}' \\ 0 & 0 & 1 \end{bmatrix} \quad (5)$$

, where  $\bar{x}$  and  $\bar{y}$  are the means of the x and y coordinates of the left points and  $\bar{x}'$  and  $\bar{y}'$  are the means of the x' and y' coordinates of the right points, and  $s$  and  $s'$  are  $\sqrt{2}/\sigma$  and  $\sqrt{2}/\sigma'$  respectively.

These normalization matrices can be applied to each set of image points to form  $x_n$  and  $x'_n$ . We can now construct the equation to solve for F. It must be that  $AF = 0$ , where A is constructed in the following manner. For each point correspondence  $x_{n,i} <-> x'_{n,i}$ :

$$\begin{bmatrix} x'_{n,i,x}x_{n,i,x} & x'_{n,i,x}x_{n,i,y} & x'_{n,i,x} & x'_{n,i,y}x_{n,i,x} & x'_{n,i,y}x_{n,i,y} & x'_{n,i,y} & x_{n,i,x} & x_{n,i,y} & 1 \end{bmatrix} \begin{bmatrix} F_{11} \\ F_{12} \\ F_{13} \\ F_{21} \\ F_{22} \\ F_{23} \\ F_{31} \\ F_{32} \\ F_{33} \end{bmatrix} = 0 \quad (6)$$

Since we have 8 unknowns (due to ratios) we must have 8 correspondences added to A. The solution to this is therefore the SVD decomposition of A to get the elements of F. However, this needs to be conditioned to be of rank 2. This can be done by performing SVD on F and setting its smallest singular value to 0.

## 2.2 Epipole Estimation

Once F is estimated, we can calculate the epipoles. We know that these must satisfy the following conditions:

$$Fe = 0 \quad (7)$$

and

$$e'F = 0 \quad (8)$$

This can again be done through SVD with  $F$  for  $e$  and  $F^T$  for  $e'$ .

## 2.3 Camera Matrix Estimation

Once the fundamental matrix and the epipoles have been estimated, the camera matrices  $P$  and  $P'$  can be estimated. We assume canonical form, thereby setting

$$P = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \quad (9)$$

$P'$  can be found through the construction of a skew-symmetric matrix of rank 3 derived from  $e'$ :

$$s = \begin{bmatrix} 0 & -e'_2 & e'_1 \\ e'_2 & 0 & -e'_0 \\ -e'_1 & e'_0 & 0 \end{bmatrix} \quad (10)$$

$P'$  can then be found through:  $P' = (sF|e')$

## 2.4 Refinement of P and P'

We now want to refine our estimation for  $P'$ . This can be done using Levenberg-Marqardt (LM) with the cost function being defined as the squared euclidean distance between the reprojected point and the original point correspondence. This is achieved through triangulation in the following manner: Suppose

$$A = \begin{bmatrix} x[0]P[2] - P[0] \\ x[1]P[2] - P[1] \\ x'[0]P'[2] - P'[0] \\ x'[1]P'[2] - P'[1] \end{bmatrix} \quad (11)$$

The SVD of A results in the triangulated world3D coordinate  $X$  of  $x$  and  $x'$ .

Then,

$$x_{reprojected} = PX \quad (12)$$

and

$$x'_{reprojected} = P'X \quad (13)$$

LM optimizes  $P'$  to minimize the reprojection error.

## 2.5 Refinement of F

Once  $P'$  is refined, the refined version of  $F$  can be calculated again through the help of the skew-symmetric matrix  $s$ .  $F$  is simply defined as:

$$F = sP'_{refined}P^+ \quad (14)$$

This  $F$  needs to again be conditioned to be rank 2, and the refined epipoles can be calculated through finding the left and right null vectors of the refined  $F$  as seen from before.

## 2.6 Estimation of H and H'

$H$  and  $H'$  can now be found. The first step is to transform the image center to the origin through the following homography:

$$T_1 = \begin{bmatrix} 1 & 0 & -w/2 \\ 0 & 1 & -h/2 \\ 0 & 0 & 1 \end{bmatrix} \quad (15)$$

We then want to send the right epipole to  $[f \ 0 \ 1]$ . This can be done through the following rotation matrix:

$$R = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (16)$$

where  $\theta = -\arctan e'_y/e'_x$

Upon performing  $Re'$ , we can extract  $f$ .

Then, we can send this epipole to infinity along the x-axis through the homography:

$$G = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -1/f & 0 & 1 \end{bmatrix} \quad (17)$$

Finally, we want to transform the image center back using

$$T_2 = \begin{bmatrix} 1 & 0 & w/2 \\ 0 & 1 & h/2 \\ 0 & 0 & 1 \end{bmatrix} \quad (18)$$

Finally,  $H' = T_2 G R T_1$ .

Computation of  $H$  follows a similar process, however now we perform the above steps using the left epipole and set its result to  $\hat{H}$ . We then construct A and b matrices in the following manner:

$$A_i = [\hat{H}x_i[0] \ \hat{H}x_i[1] \ 1] \quad (19)$$

and

$$b_i = [H'x'_i[0]] \quad (20)$$

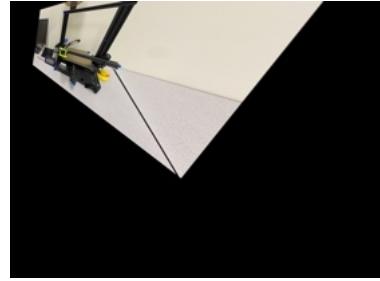
Solving for the minimization, we get:

$$m = A^+ b \quad (21)$$

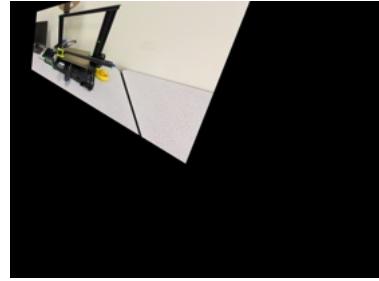
This allows us to set up:

$$H_a = \begin{bmatrix} m[0] & m[1] & m[2] \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (22)$$

Finally,  $H = H_a \hat{H}$  Applying these homographies to rectify the images, we are left with the following rectified images:



(a) Left Rectified Image



(b) Right Rectified Image

Additionally, here are the correspondence lines drawn between the two:

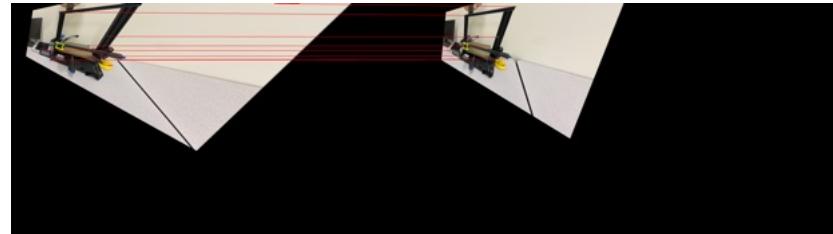


Figure 4: Rectified Images with Correspondence Lines

## 2.7 Interest Point Detection

Now that the images are rectified, we can continue with interest point detection. First, a canny edge detector was applied with threshold values of 600 and 800. The results of this can be seen below (please zoom in):



(a) Left Canny Detection



(b) Right Canny Detection

The pixel locations of true values could then be extracted and used as point-to-point correspondences to further refine  $F$ . This process entailed sweeping a window across each pixel in the left canny image and sweeping a corresponding window across a small region around the same row in the right canny image. The sum of squared differences (SSD) across both windows were computed, and the corresponding center pixel in the right image with the lowest SSD score was deemed to be the left image pixel's corresponding match. Lines were drawn between each correspondence, and can be seen in the dense set below:

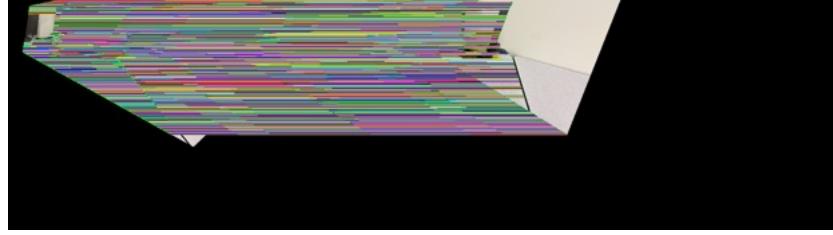


Figure 6: Left Rectified Image

## 2.8 Projective Reconstruction

Finally, these dense correspondences were then used to refine  $P'$  once more using the process from above. With the final refined  $P'$ , the world coordinates of the object could be extracted using the triangulation technique from above as well, resulting in the following 3D pointcloud with projective distortion:

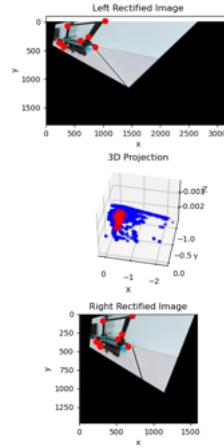
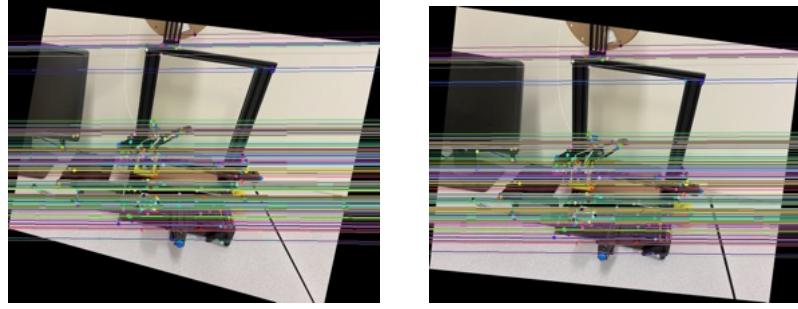


Figure 7: Rectification result, with original 8-point correspondences in Red

## 3 Loop and Zhang

Loop and Zhang is an image rectification algorithm that removes projective and affine distortion from images. It first uses the vanishing points of orthogonal lines to estimate the projective transformation, and aligns them to be parallel. Then, it enforces perpendicularity between the orthogonal lines, resulting in a rectified image true to the original proportions. Here is the result of Loop and Zhang on my stereo image pair:



(a) Left Image Rectification Loop and (b) Right Image Rectification Loop  
Zhang and Zhang

Loop and Zhang has a lot less distortion than my implementation, because that is what it is designed to circumvent. However, this algorithm produces much less correspondences than my implementation does. This may be a good thing, as I did not perform any outlier rejection like RANSAC on my dense correspondences which may lead to inaccuracies. I do like how my implementation combines the two rectified images into a single image so you can really see which points correspond to which other points, something Loop and Zhang does not do.

## 4 Disparity Map Using Census Transform

This task focused on calculation of disparity between two rectified images. Disparity essentially describes how much any given shared point between two images is shifted along a row in one image compared to another. Large disparity means large shifts and indicates an object in the foreground, while small disparity indicates objects in the background. The census transform is an algorithm to calculate this disparity, and works in the following manner:

- A parent pixel is determined through raster order of the left image and its bitvector is calculated
- In the right image, a window is swept along the image in raster order and its bitvector is calculated (The bitvector is calculated by looking inside of the window and counting how many pixels are strictly larger in intensity than the center pixel)
- The two bitvectors are then XORed together, and the cost is the number of 1s are in the resultant word
- The window is then shifted in the x-direction from 0 to  $d_{max}$  and the bitvector and cost are then recalculated. The disparity at the parent point is the pixel shift that results in the least cost between the two bitvectors

I wrote my own implementation of this census transform algorithm and compared it with the ground truth disparity maps given to us. Disparity was deemed accurate if the difference between calculated and ground truth disparity was less than  $\delta = 2$ . I tested this with a sliding window size of 5x5 and 11x11. The 5x5 sliding window resulted in 40 accuracy, while the 11x11 window resulted in 51 accuracy.

Below are my calculated disparity maps along with the ground truth disparity maps:



(a) 5x5 Left Disparity Map



(b) 11x11 Left Disparity Map

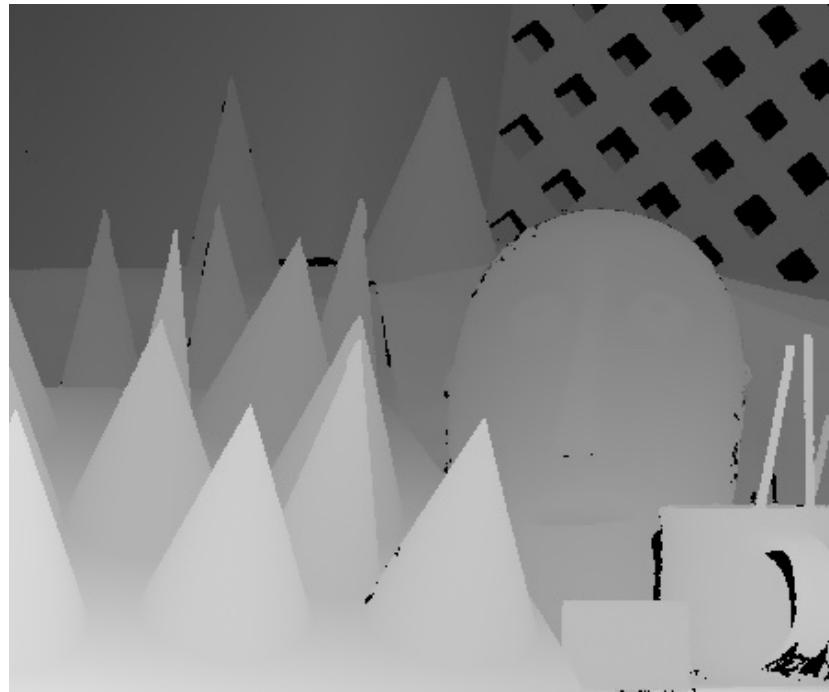
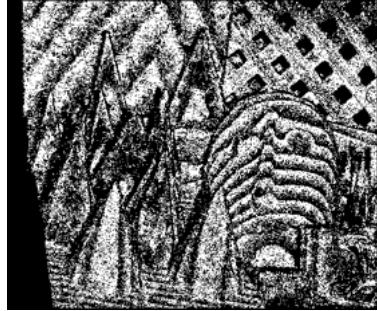


Figure 10: Ground Truth Left Disparity Map

Below are the error masks for these two tests:



(a) 5x5 Left Error Mask



(b) 11x11 Left Error Mask

It seems that increasing the window size makes the disparity map denser, while also increasing the accuracy. However, this increases the processing time as this increases the size of the bitvectors that are calculated.

## 5 Automatic Extraction of Dense Correspondences Using Depth Maps

This refers to finding corresponding points between two images using just the shared depth data between their depth maps. Essentially, if you image the same object, you can determine the camera's pose from the 3D geometry and determine the corresponding homographies between each image using the 3D geometry as a middle ground to project to. This allows for much more accurate correspondences, as it uses the known 3D geometry to register points to and backproject back to their own camera coordinates. In addition, depth maps are very commonly used in robotics, so this has many useful applications for figuring out which 2 points describe the same object in a stereo pair. This is particularly useful for identifying similar points across varying contexts, for example with images taken at different times or perhaps under inclement weather.

This process works in the following manner:

- Find the inverse of the left camera intrinsic matrix, and apply to each point you want to find a correspondence to  $X_{cam,left} = K^{-1}x_{left}$  - this is simply the projection of the pixel along a ray stemming from the camera
- Multiply  $X_{cam,left}$  by its corresponding depth value, resulting in  $X_{cam,left}$  being transformed to be a 3D point in the camera's coordinate frame
- Convert  $X_{cam,left}$  into the world3D coordinate frame using extrinsics
- Using the right camera's extrinsics, we can estimate this world coordinate's location  $X_{cam,right}$ , which includes its estimated depth  $Z_{cam,right}$

- Backproject  $X_{cam,right}$  to  $x_{right}$  and find the true depth using the right image's depth map
- Compare estimated depth and true depth, see if lie below threshold. If yes,  $x_{left}$  and  $x_{right}$  are correspondences.

The following shows the successful execution of this process for a variety of images with their corresponding depth maps. Below are the image and depth pairs:

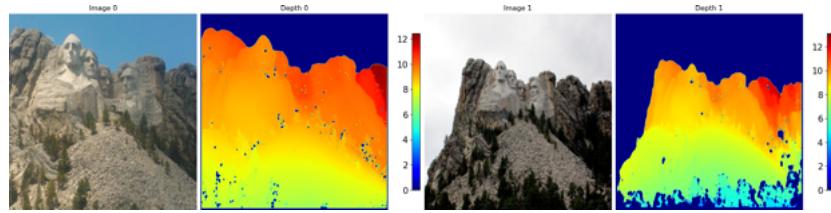


Figure 12: Image and Depth Pair 0

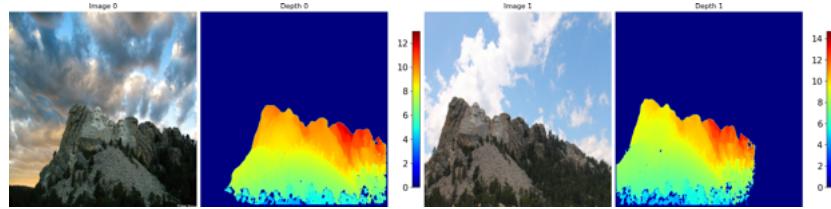


Figure 13: Image and Depth Pair 1

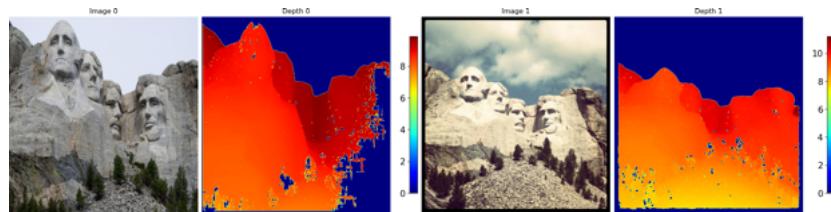


Figure 14: Image and Depth Pair 2

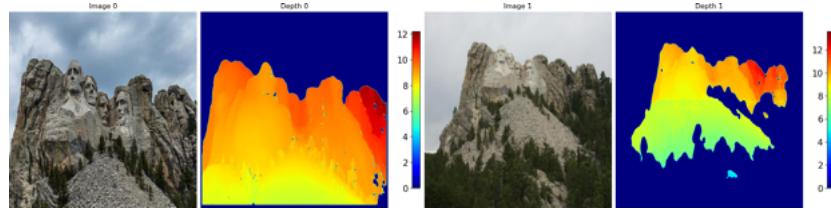


Figure 15: Image and Depth Pair 3

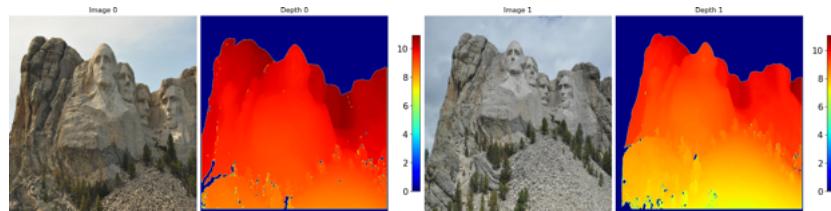


Figure 16: Image and Depth Pair 4

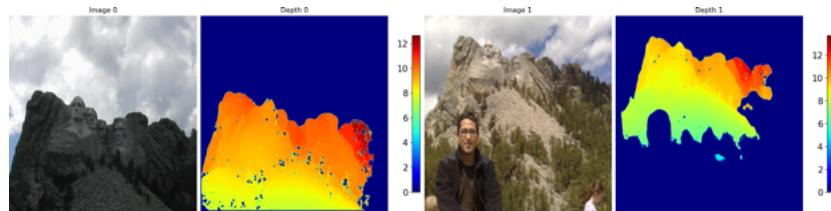


Figure 17: Image and Depth Pair 5

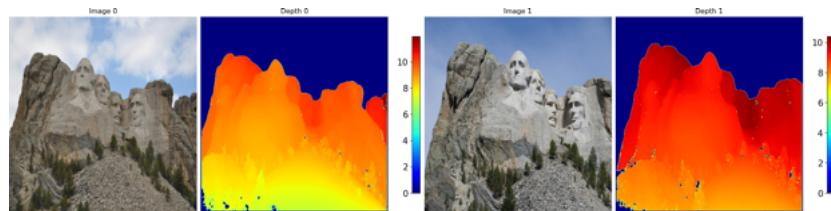


Figure 18: Image and Depth Pair 6

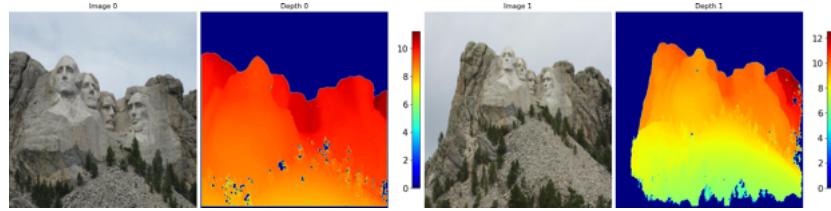


Figure 19: Image and Depth Pair 7

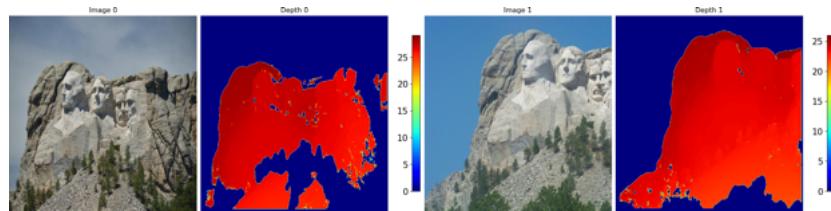


Figure 20: Image and Depth Pair 8

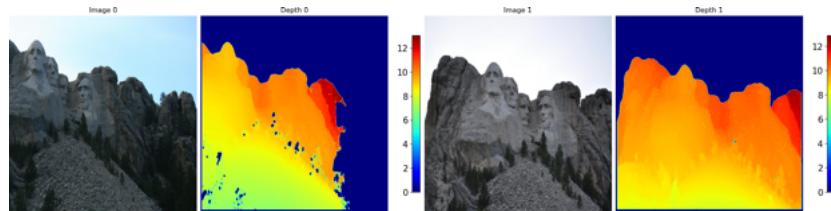


Figure 21: Image and Depth Pair 9

Below are the successfully found correspondences:

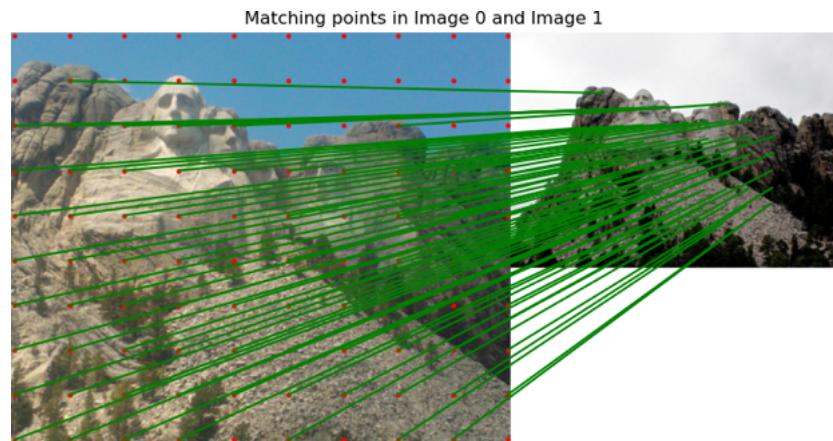


Figure 22: Depth Check Pair 0

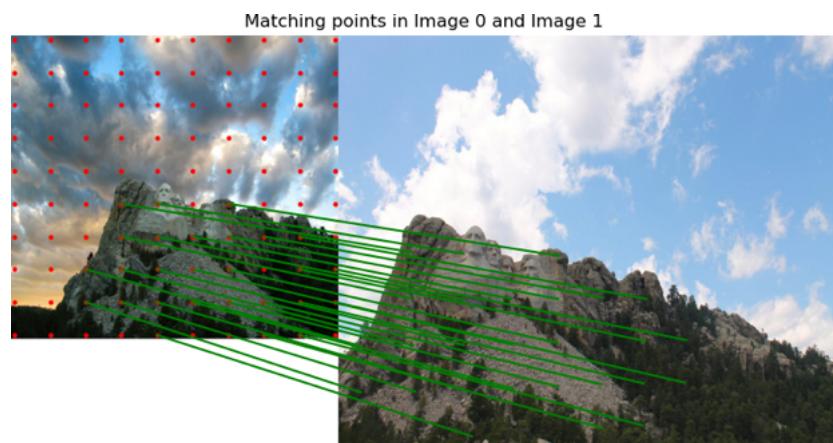


Figure 23: Depth Check Pair 1

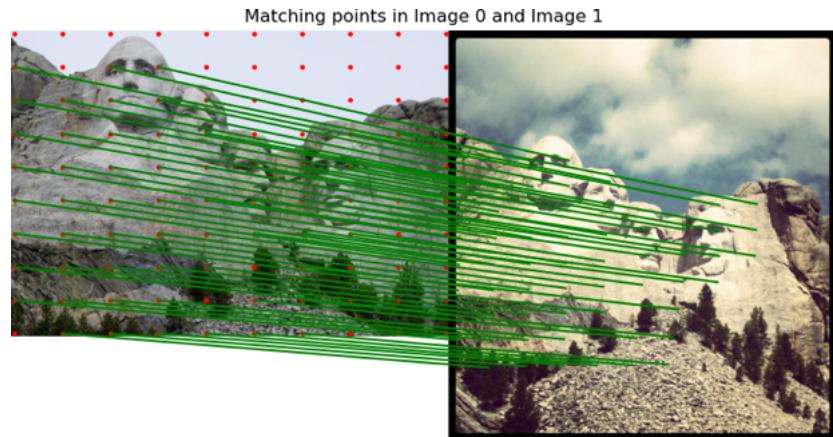


Figure 24: Depth Check Pair 2

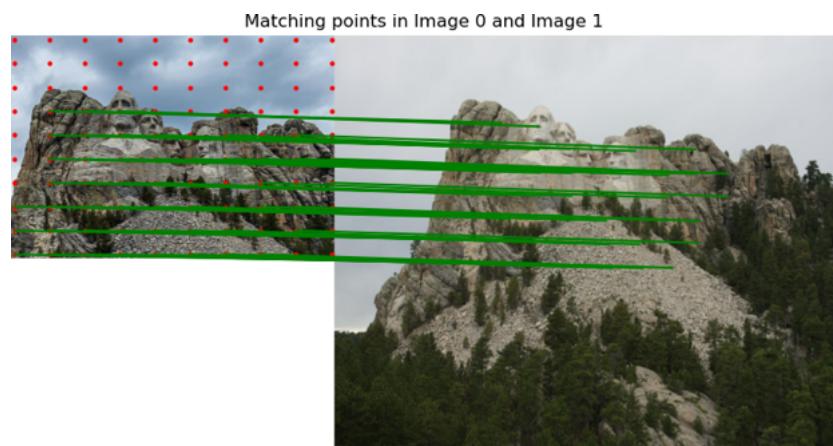


Figure 25: Depth Check Pair 3

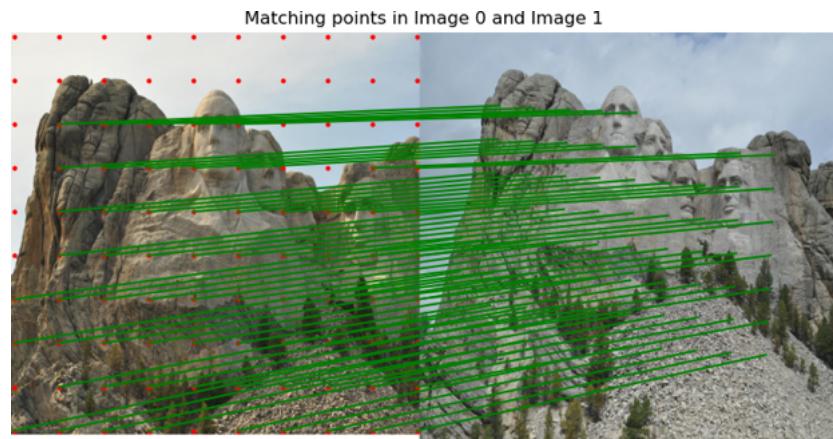


Figure 26: Depth Check Pair 4

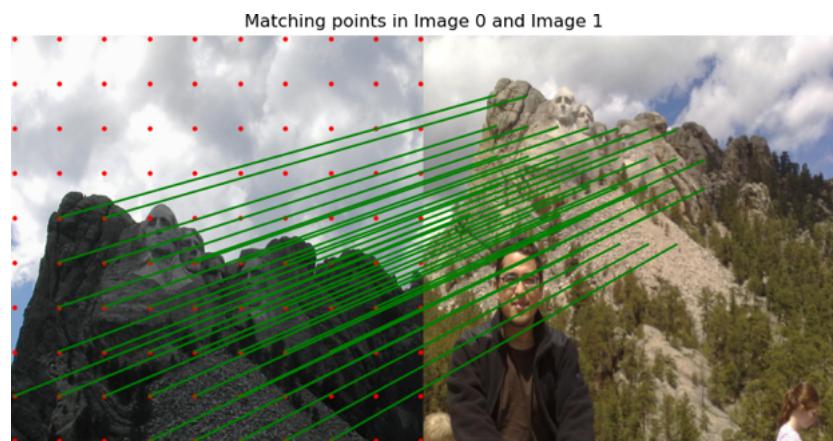


Figure 27: Depth Check Pair 5

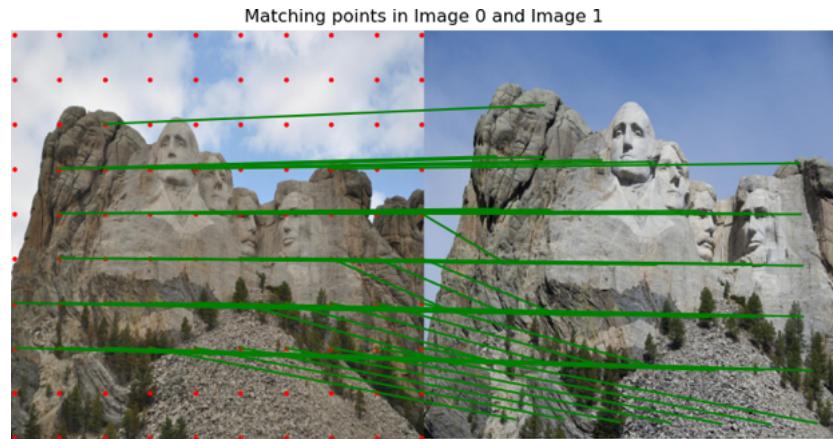


Figure 28: Depth Check Pair 6

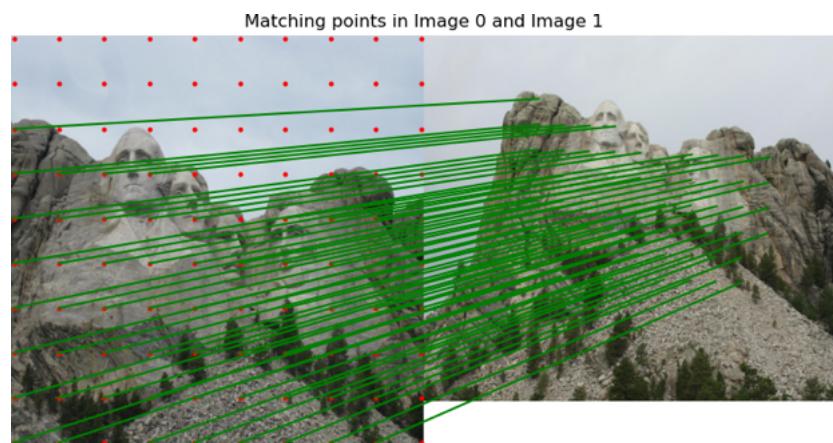


Figure 29: Depth Check Pair 7

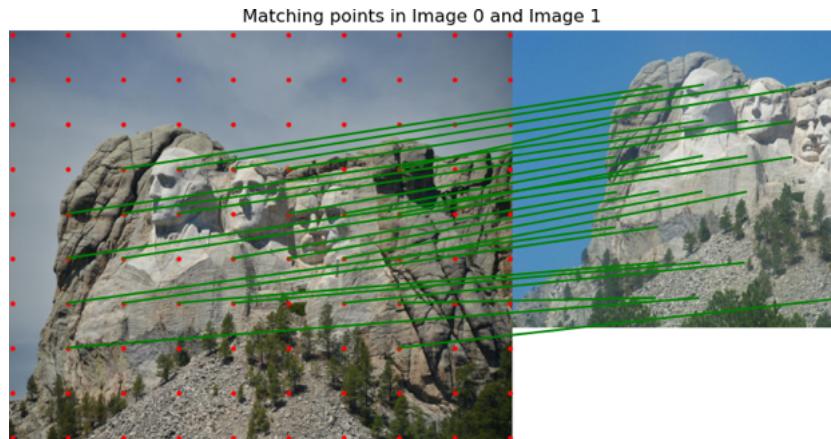


Figure 30: Depth Check Pair 8

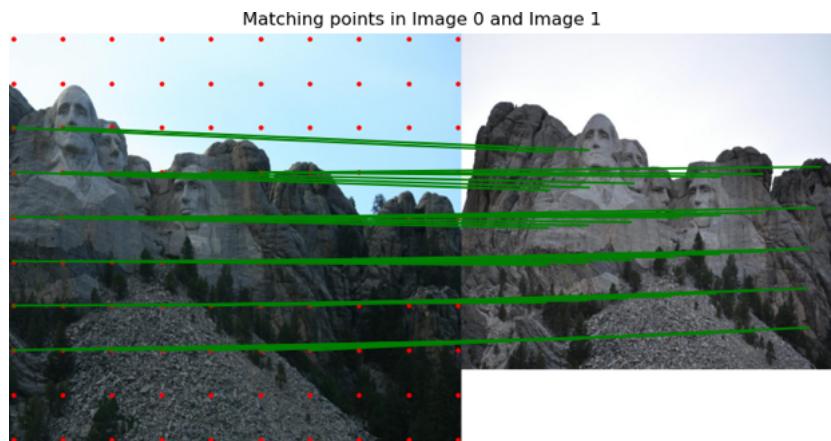


Figure 31: Depth Check Pair 9

Finally, we can generate an even denser set of correspondences (Here, I generated 250,000 correspondences) to form a pointcloud of the world 3D coordinates of the object in all 10 scenarios:

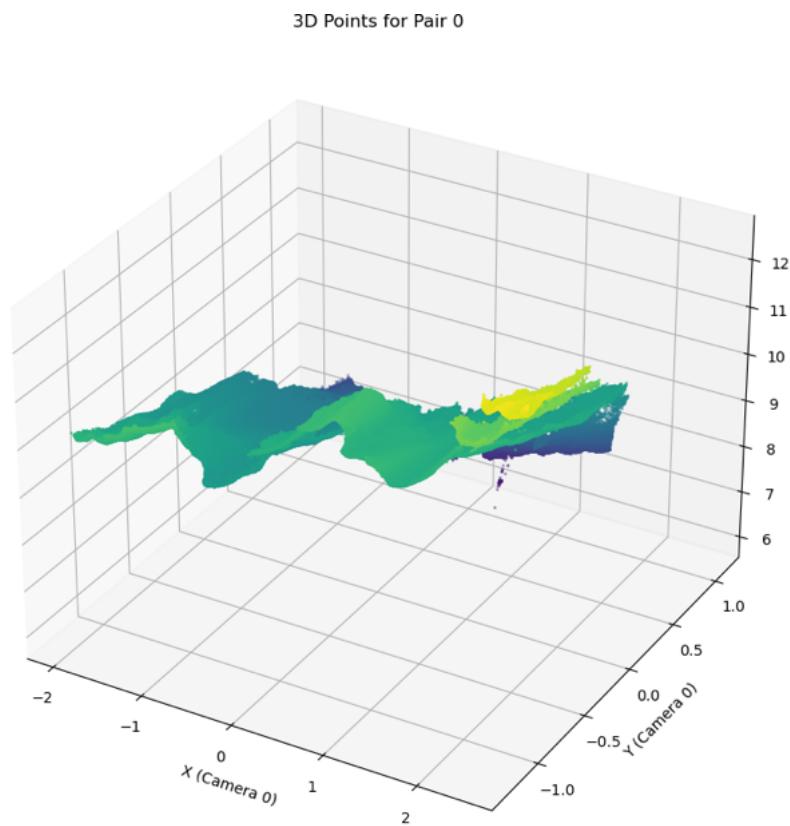


Figure 32: 3D Scatter Plot for Pair 0

3D Points for Pair 1

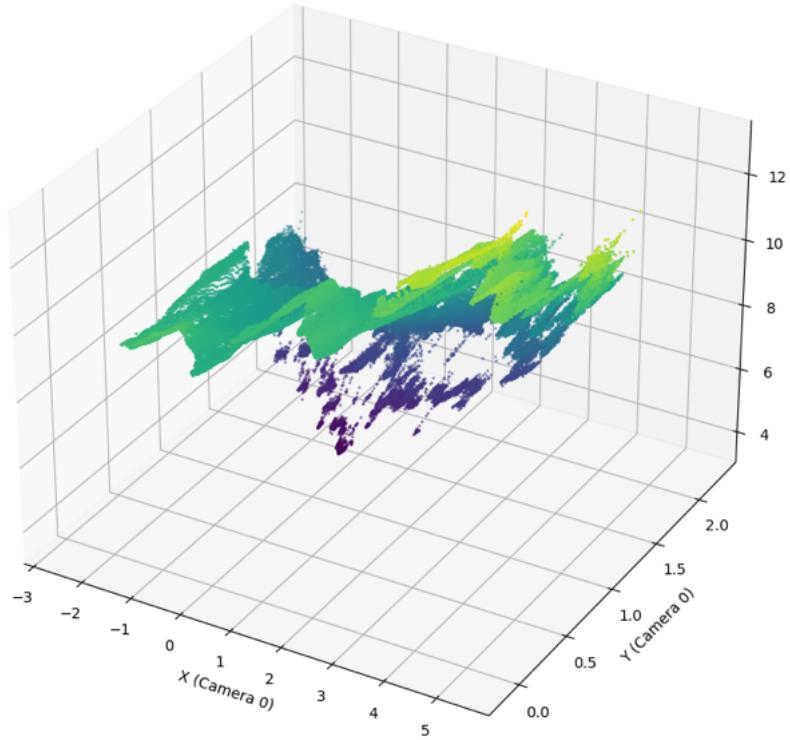


Figure 33: 3D Scatter Plot for Pair 1

3D Points for Pair 2

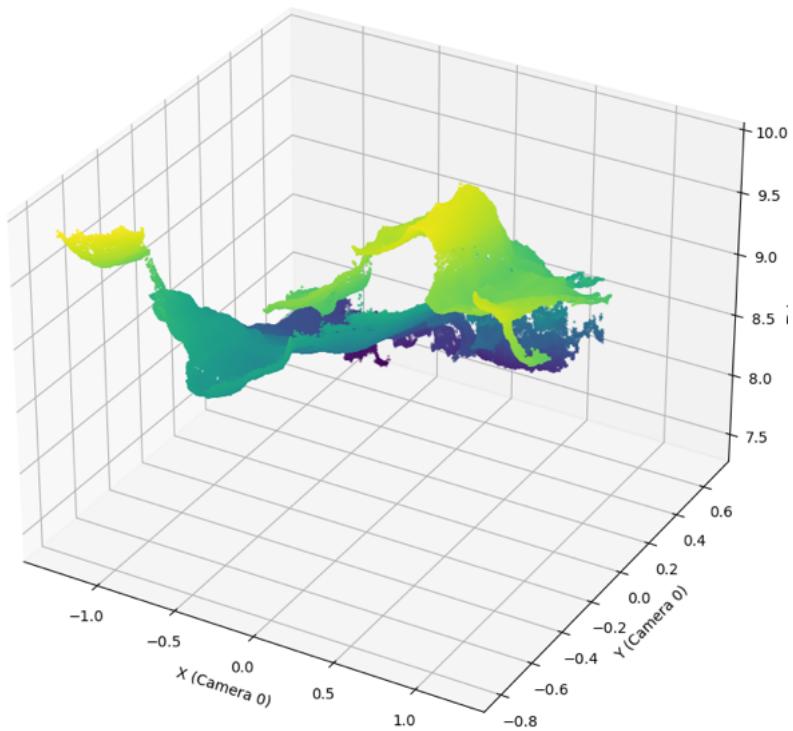


Figure 34: 3D Scatter Plot for Pair 2

3D Points for Pair 3

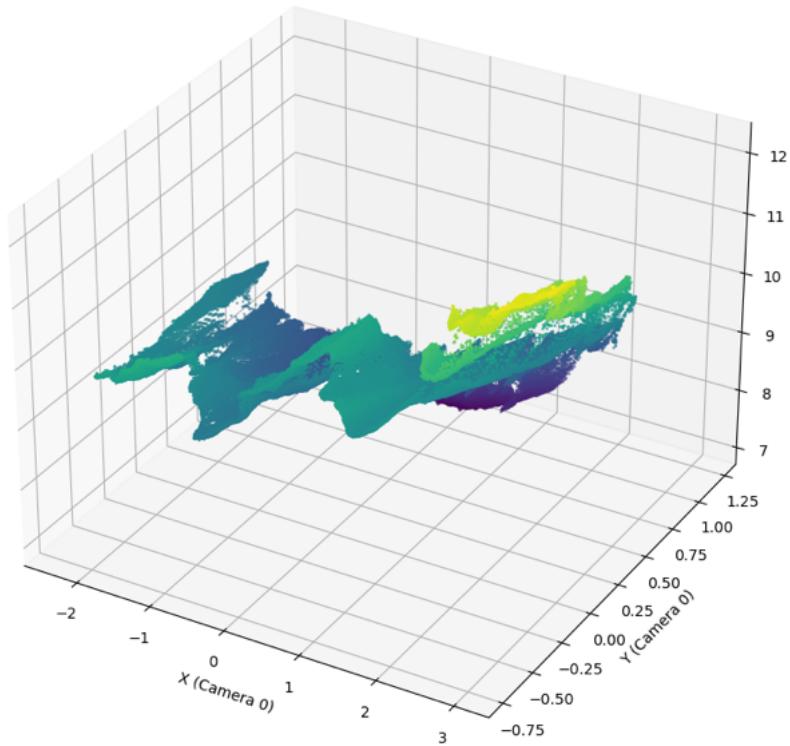


Figure 35: 3D Scatter Plot for Pair 3

3D Points for Pair 4

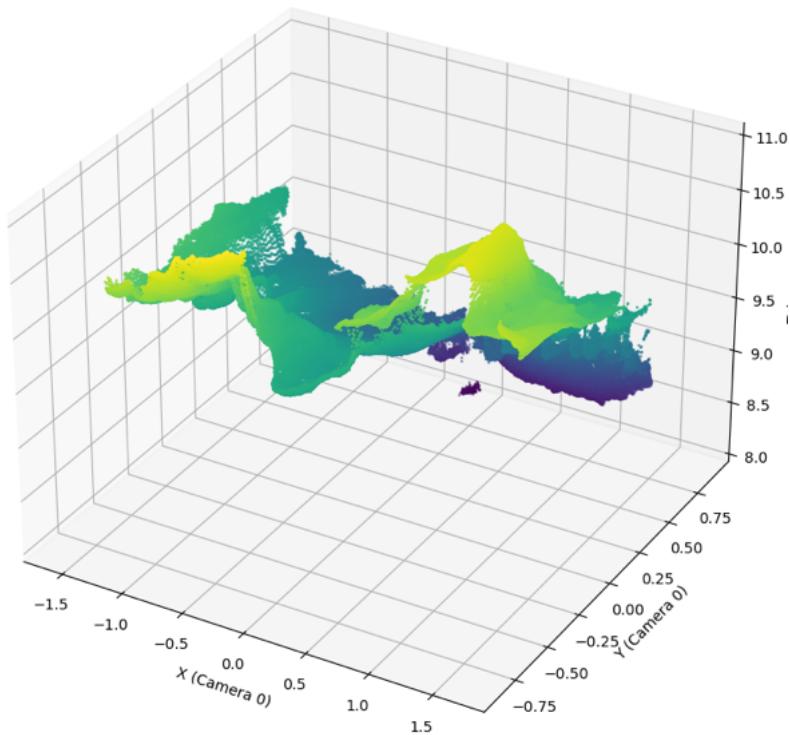


Figure 36: 3D Scatter Plot for Pair 4

3D Points for Pair 5

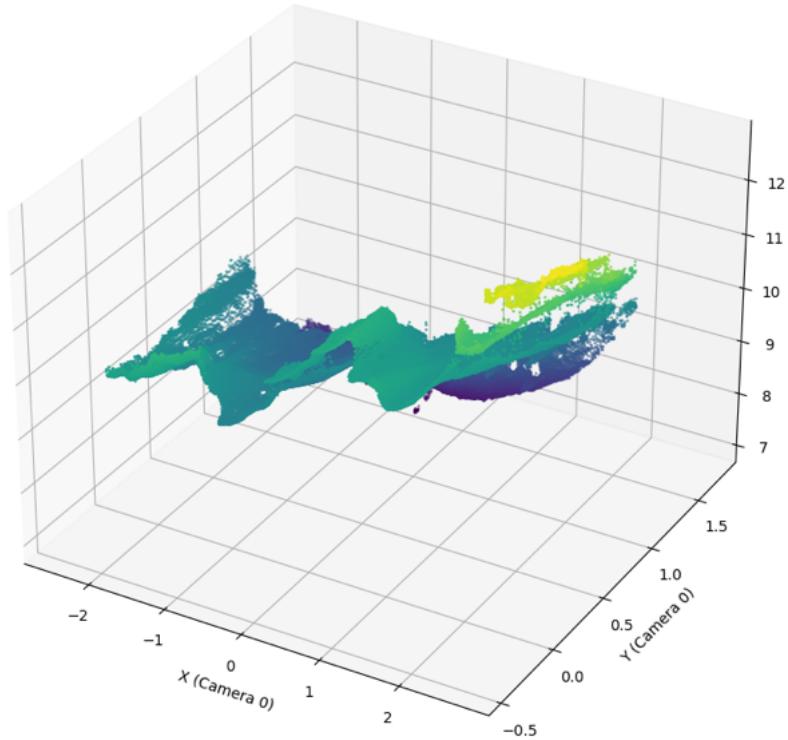


Figure 37: 3D Scatter Plot for Pair 5

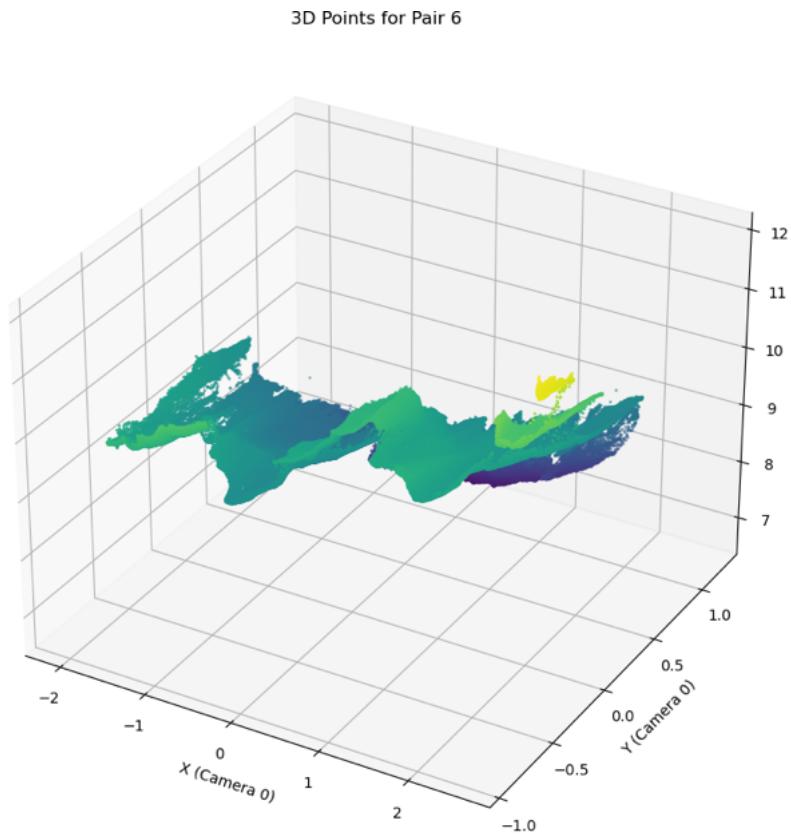


Figure 38: 3D Scatter Plot for Pair 6

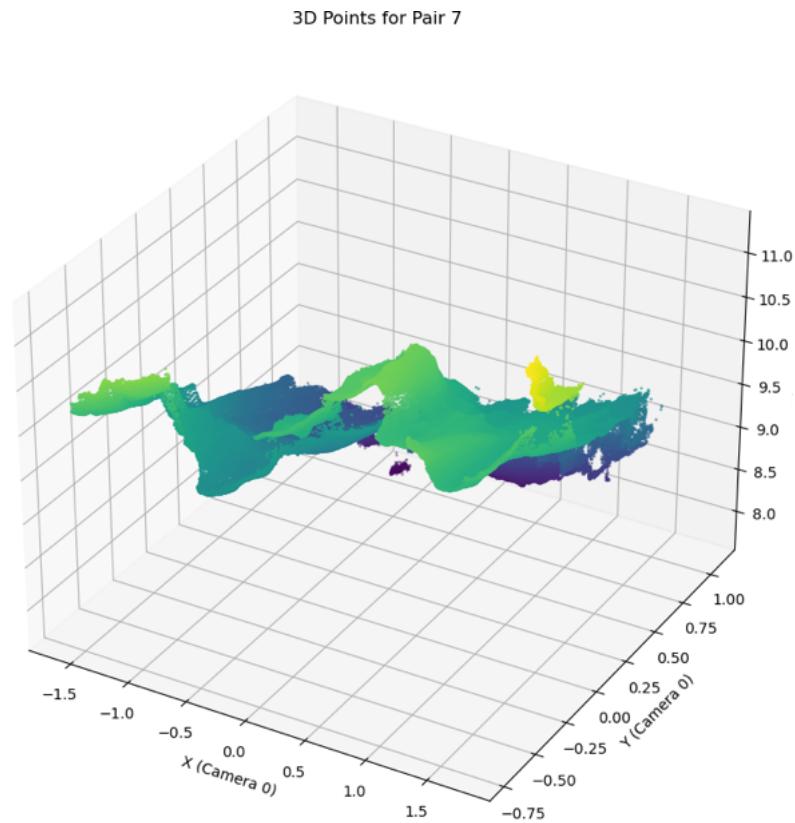


Figure 39: 3D Scatter Plot for Pair 7

3D Points for Pair 8

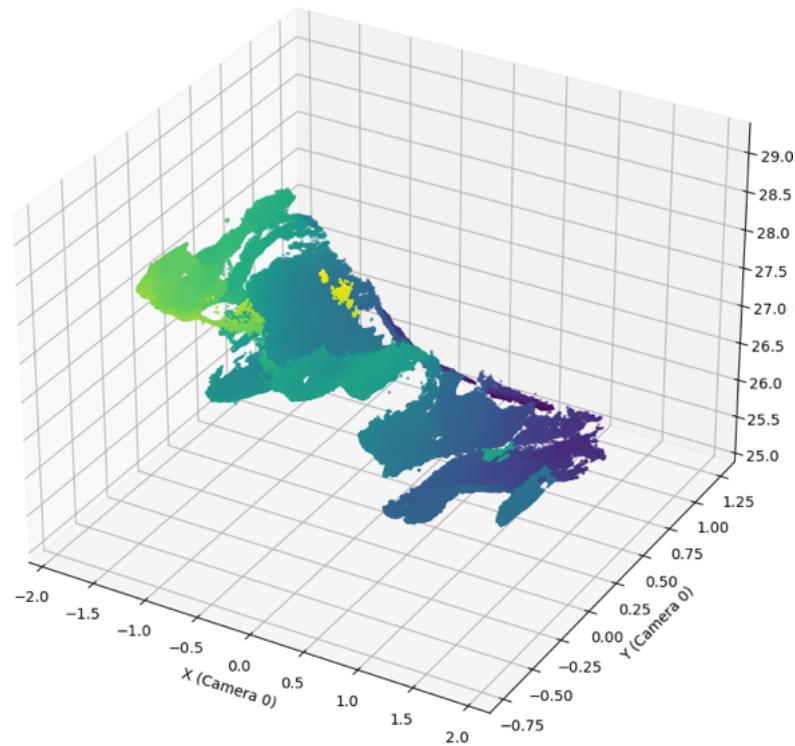


Figure 40: 3D Scatter Plot for Pair 8

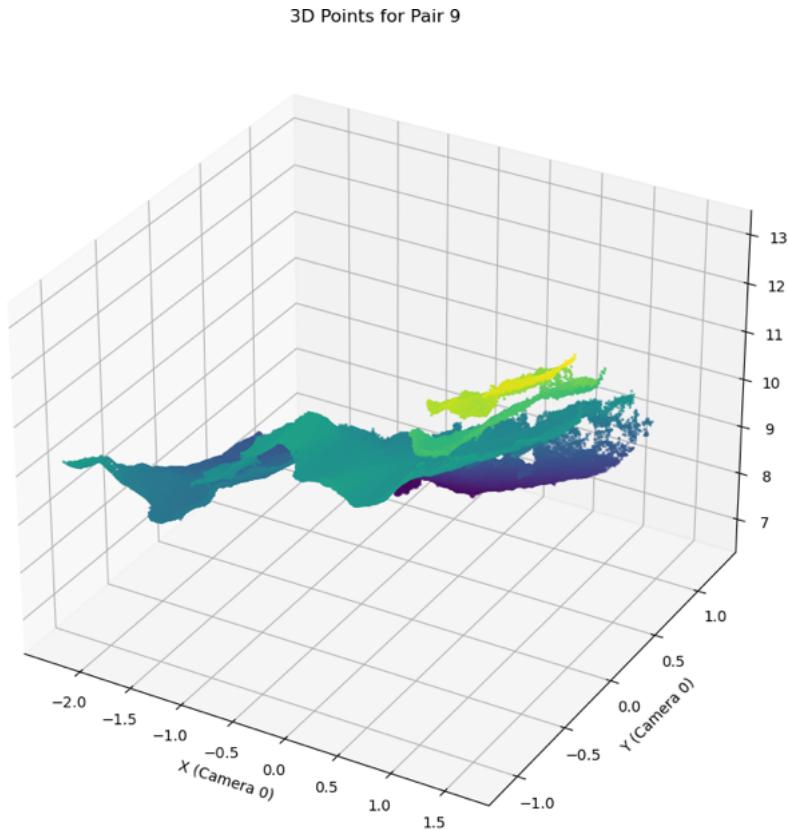


Figure 41: 3D Scatter Plot for Pair 9

## 6 Code

```

1 import cv2
2 import numpy as np
3 import random
4 import time
5 import matplotlib.pyplot as plt
6
7 from matplotlib.patches import ConnectionPatch
8 from scipy.optimize import least_squares
9 from scipy.linalg import null_space
10 from mpl_toolkits.mplot3d import Axes3D
11
12
13 image1 = cv2.imread("img1_alt.png")
14 image2 = cv2.imread("img2_alt.png")
15
16 image1_height, image1_width, _ = image1.shape

```

```

17     image2_height, image2_width, _ = image2.shape
18
19 # cv2.imshow("image1", image1)
20 # cv2.imshow("image2", image2)
21 # cv2.waitKey(0)
22 # cv2.destroyAllWindows()
23
24 visualize = True
25
26 # Since we are using canonical form this is constant
27 P = np.array([[1, 0, 0, 0], [0, 1, 0, 0], [0, 0, 1, 0]])
28
29
30 def HC(point): # Helper function to turn pixels into HC
31     return np.array([point[0], point[1], 1]).reshape((3, 1))
32
33 def dist(p1, p2):
34     return np.sqrt((p1[0] - p2[0])**2 + (p1[1] - p2[1])**2)
35
36 def condition_F(F):
37     # Conditioning F by performing SVD and setting its smallest
38     # singular value in D to 0
39     U, D, Vt = np.linalg.svd(F)
40     D[-1] = 0
41
42     D = np.diag(D)
43
44     F_conditioned = np.dot(U, np.dot(D, Vt))
45
46     return F_conditioned
47
48 def calculate_raster(image, H):
49     height, width, _ = image.shape
50
51     # Transforming corners through homography and choosing max
52     # bounds
53     corners = np.array([
54         [0, 0, 1],
55         [width, 0, 1],
56         [width, height, 1],
57         [0, height, 1]
58     ]).T
59
60     transformed_corners = H @ corners
61     transformed_corners /= (transformed_corners[-1, :] + 1e-6)
62
63     min_x = min(transformed_corners[0, :])
64     max_x = max(transformed_corners[0, :])
65     min_y = min(transformed_corners[1, :])
66     max_y = max(transformed_corners[1, :])
67
68     w_out = int(max_x - min_x)
69     h_out = int(max_y - min_y)
70
71     return w_out, h_out

```

```

72 def visualize_correspondences():
73
74     for (x1, y1), (x2, y2) in zip(image1_points, image2_points):
75         cv2.circle(combined_image, (x1, y1), radius=5, color=(0, 0,
76         255), thickness=-1)
77         cv2.circle(combined_image, (x2 + image1_width, y2), radius
78 =5, color=(0, 0, 255), thickness=-1)
79
80         cv2.line(combined_image, (x1, y1), (x2 + image1_width, y2),
81         color=(0, 255, 0), thickness=2)
82
83 plt.imshow(combined_image)
84 plt.show()
85
86 def normalize_points(points):
87     # Creation of normalization homography for 8-point method
88     mean = np.mean(points, axis=0)
89     std_dev = np.std(points, axis=0)
90     scale = np.sqrt(2) / std_dev
91
92     T = np.array([
93         [scale[0], 0, -scale[0] * mean[0]],
94         [0, scale[1], -scale[1] * mean[1]],
95         [0, 0, 1]
96     ])
97
98     normalized_points = np.dot(T, np.column_stack((points, np.ones(
99         len(points))))).T.T
100
101     return normalized_points[:, :2], T
102
103 def estimate_F(image1_points, image2_points):
104     # Normalize points
105     points1_normalized, T1 = normalize_points(image1_points)
106     points2_normalized, T2 = normalize_points(image2_points)
107
108     # Construct A matrix
109     A = np.array([
110         [x2 * x1, x2 * y1, x2, y2 * x1, y2 * y1, y2, x1, y1, 1]
111         for (x1, y1), (x2, y2) in zip(points1_normalized,
112         points2_normalized)
113     ])
114
115     # SVD to solve for fundamental matrix parameters
116     _, _, Vt = np.linalg.svd(A) # Using SVD to solve for homography
117     terms
118     F_normalized = Vt[-1].reshape((3, 3))
119
120     # Condition F to be rank 2
121     F_conditioned = condition_F(F_normalized)
122
123     F = T2.T @ F_conditioned @ T1
124
125     F /= F[2, 2]
126
127     return F

```

```

123
124 def estimate_epipoles(F):
125     # Epipoles are left and right null vectors of F
126     e = null_space(F)
127     e_p = null_space(F.T)
128
129     e /= e[-1]
130     e_p /= e_p[-1]
131
132     return e, e_p
133
134
135 def calculate_camera_matrices(F, e_p):
136     e_p = e_p.flatten()
137
138     # Skew symmetric matrix for construction of P'
139     s = np.array([[0, -e_p[2], e_p[1]], [e_p[2], 0, -e_p[0]], [-e_p[1], e_p[0], 0]])
140
141     P_p = np.hstack((s @ F, e_p.reshape(-1, 1)))
142
143     return P_p
144
145 def compute_rectification_homography_right(right_epipole):
146     # Translational homographies to bring epipole to image origin
147     # and back
148     T1 = np.array([[1, 0, -image2_width / 2], [0, 1, -image2_height / 2], [0, 0, 1]])
149     T2 = np.array([[1, 0, image2_width / 2], [0, 1, image2_height / 2], [0, 0, 1]])
150
151     right_epipole = (T1 @ right_epipole).ravel()
152
153     right_epipole = right_epipole / right_epipole[2]
154
155     # Calculate angle to bring epipole to infinity
156     theta = -np.arctan2(right_epipole[1], right_epipole[0])
157
158     R = np.array([[np.cos(theta), -np.sin(theta), 0], [np.sin(theta), np.cos(theta), 0], [0, 0, 1]])
159
160     transformed_epipole = (R @ right_epipole.reshape((3, 1))).ravel()
161
162     transformed_epipole = transformed_epipole / transformed_epipole[2]
163
164     f = transformed_epipole[0]
165
166     G = np.array([[1, 0, 0], [0, 1, 0], [-1 / f, 0, 1]])
167
168     H_p = T2 @ G @ R @ T1
169
170     H_p /= H_p[2, 2]
171
172     return H_p

```

```

173 def compute_rectification_homography_left(H_p, left_epipole,
174     image1_points, image2_points):
175     # Initial estimate for left homography
176     H_hat = compute_rectification_homography_right(left_epipole)
177
178     A = []
179     b = []
180
181     # Use minimization technique to find left homography
182     for point_left, point_right in zip(image1_points, image2_points):
183         point_left_hc = HC(point_left)
184         point_right_hc = HC(point_right)
185
186         transformed_point_left = (H_hat @ point_left_hc).flatten()
187         transformed_point_right = (H_p @ point_right_hc).flatten()
188
189         A.append([transformed_point_left[0], transformed_point_left[1], 1])
190         b.append(transformed_point_right[0])
191
192     A = np.array(A)
193     b = np.array(b)
194
195     min_solution = np.linalg.pinv(A) @ b
196
197     Ha = np.array([[min_solution[0], min_solution[1], min_solution[2]], [0, 1, 0], [0, 0, 1]])
198
199     H = Ha @ H_hat
200
201     H /= H[2, 2]
202
203     return H
204
205 # def verify_F(F):
206 #     for point1, point2 in zip(image1_points, image2_points):
207 #         point1_hc = HC(point1)
208 #         point2_hc = HC(point2)
209 #
210 #         print("x'^T * F * x: ", point2_hc.T @ F @ point1_hc)
211
212 def triangulate_point(P_p, point1, point2):
213     # Standard triangulation using P' and P
214     A = [
215         point1[0] * P_p[2] - P_p[0],
216         point1[1] * P_p[2] - P_p[1],
217         point2[0] * P_p[2] - P_p[0],
218         point2[1] * P_p[2] - P_p[1]
219     ]
220     A = np.array(A)
221     _, _, Vt = np.linalg.svd(A.T @ A)
222     X = Vt[-1]
223
224     return X / X[3]
225

```

```

226 def backprojection(P_p, image1_points, image2_points):
227     errors = []
228     world_points = []
229
230     for point1, point2 in zip(image1_points, image2_points):
231         X = triangulate_point(P_p, HC(point1), HC(point2)) #  
Triangulate the 3D point
232         world_points.append(X.flatten()[:3])
233
234     # Apply corresponding camera matrix to world point
235     reprojected_point1 = P @ X
236     reprojected_point2 = P_p @ X
237
238     # Normalize the projected points
239     reprojected_point1 /= reprojected_point1[2]
240     reprojected_point2 /= reprojected_point2[2]
241
242     # Calculate the Euclidean distance (geometric error)
243     error1 = np.linalg.norm(reprojected_point1[:2] - point1)**2
244     error2 = np.linalg.norm(reprojected_point2[:2] - point2)**2
245
246     errors.extend([error1, error2])
247
248     return np.array(errors), np.array(world_points)
249
250 def geometric_error(P_p_vector, image1_points, image2_points):
251
252     P_p = P_p_vector.reshape(3, 4) # Reshape the flattened P'  
vector
253
254     errors, _ = backprojection(P_p, image1_points, image2_points)
255
256     return errors
257
258 def refine_P_p(P_p, image1_points, image2_points):
259
260     P_p_vector = P_p.flatten() # Flatten P' for optimization
261
262     # Perform LM on P_p to reduce euclidean distance error
263     result = least_squares(
264         geometric_error,
265         P_p_vector,
266         method='lm',
267         args=(image1_points, image2_points)
268     )
269
270     refined_P_p = result.x.reshape(3, 4) # Reshape optimized  
vector back to matrix
271
272     return refined_P_p
273
274 def rectify_images(image1, image2, image1_points, image2_points,  
collect_world_points):
275
276     # Estimate F from point correspondences
277     F = estimate_F(image1_points, image2_points)

```

```

279     # Calculate initial estimates for epipoles
280     initial_e, initial_e_p = estimate_epipoles(F)
281     initial_e = initial_e.reshape((3, 1))
282     initial_e_p = initial_e_p.reshape((3, 1))
283
284     # Use F and epipole estimates to construct P'
285     initial_P_p = calculate_camera_matrices(F, initial_e_p)
286
287     # Use LM to refine P' estimate
288     refined_P_p = refine_P_p(initial_P_p, image1_points,
289     image2_points)
290
291     # Extract new epipole from P', construct skew-symmetric matrix
292     # for refined F calculation
293     refined_e_p = refined_P_p[:, 3]
294     s = np.array([[0, -refined_e_p[2], refined_e_p[1]], [
295     refined_e_p[2], 0, -refined_e_p[0]], [-refined_e_p[1],
296     refined_e_p[0], 0]])
297
298     # Calculate refined F and condition to enforce rank 2
299     refined_F = s @ refined_P_p @ np.linalg.pinv(P)
300
301     refined_F /= refined_F[2, 2]
302     refined_F = condition_F(refined_F)
303
304     # Extract refined epipoles from refined F
305     refined_e, refined_e_p = estimate_epipoles(refined_F)
306
307     # Construct rectification homographies from
308     H_p = compute_rectification_homography_right(refined_e_p)
309     H = compute_rectification_homography_left(H_p, refined_e,
310     image1_points, image2_points)
311
312     if collect_world_points:
313         _, world_points = backprojection(refined_P_p, image1_points,
314         image2_points)
315         return H, H_p, refined_P_p, world_points
316
317     return H, H_p, refined_P_p
318
319 def create_visualization(image1_rectified, image2_rectified, H, H_p):
320     combined_rectified_image = np.hstack((image1_rectified,
321     image2_rectified))
322
323     for (x1, y1), (x2, y2) in zip(image1_points, image2_points):
324         point_left_hc = HC((x1, y1))
325         point_right_hc = HC((x2, y2))
326
327         transformed_point_left = (H @ point_left_hc).flatten()
328         transformed_point_right = (H_p @ point_right_hc).flatten()
329
330         pixel_left = (int(transformed_point_left[0] /
331         transformed_point_left[2]), int(transformed_point_left[1] /
332         transformed_point_left[2]))
333         pixel_right = (int(transformed_point_right[0] /
334         transformed_point_right[2]), int(transformed_point_right[1] /

```

```

    transformed_point_right[2])))

325
326
327     cv2.circle(image1_rectified, pixel_left, radius=5, color
328             =(255, 0, 0), thickness=-1)
329     cv2.circle(image2_rectified, pixel_right, radius=5, color
330             =(0, 255, 0), thickness=-1)
331
332     cv2.circle(combined_rectified_image, pixel_left, radius=5,
333             color=(255, 0, 0), thickness=-1)
334     cv2.circle(combined_rectified_image, (pixel_right[0] +
335             new_width, pixel_right[1]), radius=5, color=(0, 255, 0),
336             thickness=-1)
337
338     cv2.line(combined_rectified_image, pixel_left, (pixel_right
339             [0] + new_width, pixel_right[1]), color=(0, 0, 255), thickness
340             =2)
341
342
343     return combined_rectified_image
344
345 def extract_interest_points(rectified_image, canny_threshold1,
346     canny_threshold2):
347     # Extract edges using canny detector, collect points where mask
348     # is true
349     edges = cv2.Canny(rectified_image, canny_threshold1,
350             canny_threshold2)
351     points = np.column_stack(np.where(edges > 0))
352
353     # Sort by row
354     points = points[np.lexsort((points[:, 1], points[:, 0]))]
355
356     if visualize:
357         cv2.imshow("Interest Points", edges)
358         cv2.waitKey(0)
359         cv2.destroyAllWindows()
360
361     return points
362
363 def get_candidate_points(point, interest_points2, row_window=3):
364     row_range = range(point[0] - row_window, point[0] + row_window
365             + 1)
366     return interest_points2[np.isin(interest_points2[:, 0],
367             row_range)]
368
369 def extract_patch(image, point, patch_size=5):
370     half_size = patch_size // 2
371     x, y = point
372     h, w = image.shape[:2]
373
374     # Out of bounds checking
375     if(x - half_size < 0 or x + half_size >= w or y - half_size < 0
376             or y > half_size >= h):
377         return None
378
379     # Collect all points within patch window
380     return image[x - half_size:x + half_size + 1, y - half_size:y +
381             half_size + 1]

```

```

367
368 def SSD(patch1, patch2):
369     # Sum of squared differences if patches are both all full
370     if patch1 is None or patch2 is None:
371         return float("inf")
372     return np.sum((patch1 - patch2)**2)
373
374 def compute_dense_correspondences(image1_rectified,
375                                     image2_rectified, interest_points1, interest_points2):
376     point1_correspondences = []
377     point2_correspondences = []
378
379     for point1 in interest_points1:
380         # Find potential candidates through buffered row
381         candidates = get_candidate_points(point1, interest_points2)
382         if len(candidates) == 0:
383             continue
384
385         # Go through each candidate and find lowest SSD between
386         # parent patch and potential patch
387         patch1 = extract_patch(image1_rectified, point1)
388         best_candidate, best_score = None, float("inf")
389
390         for candidate in candidates:
391             patch2 = extract_patch(image2_rectified, candidate)
392             score = SSD(patch1, patch2)
393
394             if score < best_score:
395                 best_score = score
396                 best_candidate = candidate
397
398         # Add best candidate to correspondence list
399         if best_candidate is not None:
400             point1_correspondences.append([point1[1], point1[0]])
401             point2_correspondences.append([best_candidate[1],
402                                            best_candidate[0]])
403
404     return np.array(point1_correspondences), np.array(
405                    point2_correspondences)
406
407 if __name__ == "__main__":
408     image1_points = np.array([[602, 68], [287, 64], [496, 304],
409                             [303, 506], [611, 345], [195, 446], [284, 475], [548, 233]])
410     image2_points = np.array([[584, 83], [268, 80], [505, 311],
411                             [277, 490], [598, 368], [185, 425], [263, 459], [533, 250]])
412
413     # Calculate rectification homographies
414     H, H_p, _ = rectify_images(image1, image2, image1_points,
415                                image2_points, False)
416
417     # Calculate updated buffers from homographies
418     width_new1, height_new1 = calculate_raster(image1, H)
419     width_new2, height_new2 = calculate_raster(image2, H_p)
420
421     if visualize:
422         new_width = max(width_new1, width_new2)

```

```

417     new_height = max(height_new1, height_new2)
418
419     image1_rectified = cv2.warpPerspective(image1, H, (
420         new_width, new_height))
421     image2_rectified = cv2.warpPerspective(image2, H_p, (
422         new_width, new_height))
423 else:
424     image1_rectified = cv2.warpPerspective(image1, H, (
425         width_new1, height_new1))
426     image2_rectified = cv2.warpPerspective(image2, H_p, (
427         width_new2, height_new2))
428
429 # if visualize:
430 #     combined_image = create_visualization(image1_rectified,
431 #                                             image2_rectified, H, H_p)
432
433
434 # Canny edge detection parameters
435 canny_threshold1 = 600
436 canny_threshold2 = 800
437
438 # Extract interest points
439 interest_points1 = extract_interest_points(image1_rectified,
440                                              canny_threshold1, canny_threshold2)
441 interest_points2 = extract_interest_points(image2_rectified,
442                                              canny_threshold1, canny_threshold2)
443
444 # Find correspondences using SSD patch scoring
445 new_image1_points, new_image2_points =
446     compute_dense_correspondences(image1_rectified,
447                                     image2_rectified, interest_points1, interest_points2)
448
449 # Refine P' from dense correspondences
450 _, _, refined_P_p, world_points = rectify_images(
451     image1_rectified, image2_rectified, new_image1_points,
452     new_image2_points, True)
453
454 if visualize:
455     combined_rectified_image = np.hstack((image1_rectified,
456                                           image2_rectified))
457
458     for point1, point2 in zip(new_image1_points,
459                               new_image2_points):
460         color = (random.randint(0, 255), random.randint(0, 255),
461                  random.randint(0, 255))
462
463         offset = image1_rectified.shape[1]
464         point2 = (point2[0] + offset, point2[1])
465
466         cv2.circle(combined_rectified_image, point1, radius=2,
467                    color=(0, 255, 0), thickness=-1)

```

```

459         cv2.circle(combined_rectified_image, point2, radius=2,
460                     color=(0, 255, 0), thickness=-1)
461
462         cv2.line(combined_rectified_image, point1, point2,
463                     thickness=2, color=color)
464
465         cv2.imshow("combined_image", combined_rectified_image)
466         cv2.waitKey(0)
467         cv2.destroyAllWindows()
468
469     # Populate 2D and 3D points to form pointcloud
470     visualization_points_3d = []
471     visualization_points_2d_left = []
472     visualization_points_2d_right = []
473
474     for (x1, y1), (x2, y2) in zip(image1_points, image2_points):
475         point_left_hc = HC((x1, y1))
476         point_right_hc = HC((x2, y2))
477
478         transformed_point_left = (H @ point_left_hc).flatten()
479         transformed_point_right = (H_p @ point_right_hc).flatten()
480
481         visualization_points_2d_left.append([transformed_point_left[0] / transformed_point_left[2], transformed_point_left[1] / transformed_point_left[2]])
482         visualization_points_2d_right.append([
483             transformed_point_right[0] / transformed_point_right[2],
484             transformed_point_right[1] / transformed_point_right[2]])
485
486         world_visualization_point = triangulate_point(refined_P_p,
487             transformed_point_left / transformed_point_left[2],
488             transformed_point_right / transformed_point_right[2])
489         visualization_points_3d.append(world_visualization_point)
490
491
492     visualization_points_3d = np.array(visualization_points_3d)
493
494     print("world_points shape: ", world_points.shape)
495     print("visualization_points shape: ", visualization_points_3d.shape)
496
497     # Plotting
498     X = world_points[:, 0]
499     Y = world_points[:, 1]
500     Z = world_points[:, 2]
501
502     X_v = visualization_points_3d[:, 0]
503     Y_v = visualization_points_3d[:, 1]
504     Z_v = visualization_points_3d[:, 2]
505
506     fig = plt.figure(figsize=(6, 9))
507
508     ax1 = fig.add_subplot(312, projection='3d')
509     ax2 = fig.add_subplot(311)
510     ax3 = fig.add_subplot(313)
511
512     ax1.scatter(X, Y, Z, c='b', s=10)

```

```

507     ax1.scatter(X_v, Y_v, Z_v, c='r', s=100)
508     ax1.set_xlabel("X")
509     ax1.set_ylabel("Y")
510     ax1.set_zlabel("Z")
511     ax1.set_title("3D Projection")
512
513
514     ax2.imshow(image1_rectified)
515     ax3.imshow(image2_rectified)
516
517     for i in range(len(visualization_points_3d)):
518         left_point = visualization_points_2d_left[i]
519         right_point = visualization_points_2d_right[i]
520
521         ax2.scatter(left_point[0], left_point[1], color='r', s=50)
522         ax3.scatter(right_point[0], right_point[1], color='r', s
523 =50)
524
525     ax2.set_xlabel('x')
526     ax2.set_ylabel('y')
527     ax2.set_title("Left Rectified Image")
528
529     ax3.set_xlabel('x')
530     ax3.set_ylabel('y')
531     ax3.set_title("Right Rectified Image")
532
533
534     plt.tight_layout()
535     plt.show()

```

Listing 1: Stereo Rectification

```

1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5
6 def compute_census_transform(image, window_size):
7     rows, cols = image.shape
8     half_window = window_size // 2
9     census = np.zeros((rows, cols, window_size**2), dtype=np.uint8)
10
11    for i in range(half_window, rows - half_window):
12        for j in range(half_window, cols - half_window):
13            center_pixel = image[i, j]
14            bit_vector = []
15
16            for m in range(-half_window, half_window + 1):
17                for n in range(-half_window, half_window + 1):
18                    neighbor_pixel = image[i + m, j + n]
19                    if neighbor_pixel > center_pixel:
20                        bit_vector.append(1)
21                    else:
22                        bit_vector.append(0)
23
24    census[i, j] = bit_vector

```

```

25     return census
26
27
28 def census_disparity_map(image1, image2, window_size, dmax):
29     rows, cols = image1.shape
30
31     disparity_map = np.zeros((rows, cols), dtype=np.uint8)
32     half_window = window_size // 2
33
34     census1 = compute_census_transform(image1, window_size)
35     census2 = compute_census_transform(image2, window_size)
36
37     for i in range(half_window, rows - half_window):
38         for j in range(half_window, cols - half_window):
39             min_cost = float("inf")
40             best_disparity = 0
41
42             for d in range(dmax + 1):
43                 if j - d < half_window:
44                     # Skipping disparities that go out of bounds
45                     continue
46
47                 bitvector1 = census1[i, j]
48                 bitvector2 = census2[i, j - d]
49
50                 xor = np.bitwise_xor(bitvector1, bitvector2)
51                 data_cost = np.sum(xor)
52
53                 if data_cost < min_cost:
54                     min_cost = data_cost
55                     best_disparity = d
56
57             disparity_map[i, j] = best_disparity
58
59     return disparity_map
60
61
62 if __name__ == "__main__":
63     image1 = cv2.imread("Task3Images/im2.png", cv2.IMREAD_GRAYSCALE)
64     image2 = cv2.imread("Task3Images/im6.png", cv2.IMREAD_GRAYSCALE)
65
66     ground_truth = cv2.imread("Task3Images/disp2.png", cv2.
67     IMREAD_GRAYSCALE)
68     ground_truth = np.array(np.array(ground_truth, dtype=np.float32)
69     / 4, dtype=np.uint8)
70
71     window_size = 11
72     dmax = 64
73     disparity_threshold = 2
74     num_valid_points = 0
75     num_accurate_points = 0
76
77     disparity_map = census_disparity_map(image1, image2,
78     window_size, dmax)
79     error_mask = np.zeros_like(disparity_map)

```

```

77
78     for i in range(disparity_map.shape[0]):
79         for j in range(disparity_map.shape[1]):
80             ground_truth_value = ground_truth[i, j]
81
82             if ground_truth_value > 0:
83                 num_valid_points += 1
84
85             if(ground_truth_value == 0):
86                 continue
87
88             if(abs(ground_truth_value - disparity_map[i, j]) <=
disparity_threshold):
89                 error_mask[i, j] = 255
90                 num_accurate_points += 1
91
92
93     print("accuracy: ", num_accurate_points / num_valid_points)
94
95     cv2.imshow("Disparity Map", disparity_map)
96     cv2.imshow("Ground Truth", ground_truth)
97     cv2.imshow("Error Mask", error_mask)
98     cv2.waitKey(0)
99     cv2.destroyAllWindows()
100
101 # 5x5 40% accuracy
102 # 11x11

```

Listing 2: Disparity Census Transform

```

1 import numpy as np
2 import pickle as pkl
3 import matplotlib.pyplot as plt
4 import h5py # for reading depth maps
5 import time
6
7 """
8 A few notes on the scene_info dictionary:
9 - depth maps are stored as h5 files. Depth is the distance of the
object from the camera (ie Z coordinate in camera coordinates).
The depth map can contain invalid points (depth = 0) which
correspond to points where the depth could not be estimated.
10 - The intrinsics are stored as a 3x3 matrix.
11 - The poses [R,t] are stored as a 4x4 matrix to allow for easy
transformation of points from one camera to the other. The
resulting transformation matrix is a 4x4 matrix is of the form:
T = [[R, t]
      [0, 1]] where R is a 3x3 rotation matrix and t is a 3x1
translation vector.
12 """
13
14
15 DEPTH_THR = 0.1
16 dense = False
17 if dense:
18     num_points_in_mesh = 500
19 else:
20     num_points_in_mesh = 10

```

```

22
23
24 def plot_image_and_depth(img0, depth0, img1, depth1, plot_name):
25     # Enable constrained layout for uniform subplot sizes
26     fig, ax = plt.subplots(1, 4, figsize=(20, 5),
27                           constrained_layout=True)
28
29     # Image 0
30     ax[0].imshow(img0, aspect='auto')
31     ax[0].set_title('Image 0')
32     ax[0].axis('off')
33
34     # Depth 0
35     im1 = ax[1].imshow(depth0, cmap='jet', aspect='auto')
36     ax[1].set_title('Depth 0')
37     ax[1].axis('off')
38     cbar1 = fig.colorbar(im1, ax=ax[1], shrink=0.8, aspect=20)
39     cbar1.ax.yaxis.set_ticks_position('left')
40     cbar1.ax.yaxis.set_label_position('left')
41     cbar1.ax.tick_params(labelsize=15)
42
43     # Image 1
44     ax[2].imshow(img1, aspect='auto')
45     ax[2].set_title('Image 1')
46     ax[2].axis('off')
47
48     # Depth 1
49     im2 = ax[3].imshow(depth1, cmap='jet', aspect='auto')
50     ax[3].set_title('Depth 1')
51     ax[3].axis('off')
52     cbar2 = fig.colorbar(im2, ax=ax[3], shrink=0.8, aspect=20)
53     cbar2.ax.yaxis.set_ticks_position('left')
54     cbar2.ax.yaxis.set_label_position('left')
55     cbar2.ax.tick_params(labelsize=15)
56
57     plt.savefig(plot_name, bbox_inches='tight', pad_inches=0)
58     plt.close()
59
60 if __name__ == "__main__":
61     scene_info = pkl.load(open('./data/scene_info/1589_subset.pkl',
62                            'rb'))
63
64     for i_pair in range(len(scene_info)):
65         # print(scene_info[i_pair].keys())
66         # ['image0', 'image1', 'depth0', 'depth1', 'K0', 'K1', 'T0',
67         # 'T1', 'overlap_score']
68         # print(scene_info[i_pair]['image0']) # path to image0
69         # print(scene_info[i_pair]['image1']) # path to image1
70         # print(scene_info[i_pair]['depth0']) # path to depth0
71         # print(scene_info[i_pair]['depth1']) # path to depth1
72         # print(scene_info[i_pair]['K0']) # intrinsic matrix of
73         # camera 0 [3,3]
74         # print(scene_info[i_pair]['K1']) # intrinsic matrix of
75         # camera 1 [3,3]
76         # print(scene_info[i_pair]['T0']) # pose matrix of camera
77         # [4,4]
78         # print(scene_info[i_pair]['T1']) # pose matrix of camera

```

```

1 [4,4]
# print('-----')
73
74
75     # read images
76     img0 = plt.imread(scene_info[i_pair]['image0'])
77     img1 = plt.imread(scene_info[i_pair]['image1'])
78
79     # read depth
80     with h5py.File(scene_info[i_pair]['depth0'], 'r') as f:
81         depth0 = f['depth'][:]
82     with h5py.File(scene_info[i_pair]['depth1'], 'r') as f:
83         depth1 = f['depth'][:]
84
85     # check shapes
86     h0, w0 = img0.shape[:-1]
87     h1, w1 = img1.shape[:-1]
88     assert img0.shape[:-1] == depth0.shape, f"depth and image shapes do not match: {img0}, {depth0}"
89     assert img1.shape[:-1] == depth1.shape, f"depth and image shapes do not match: {img1}, {depth1}"
90
91
92     # plot image and depth
93     plot_name = f'./pics/image_and_depth_pair_{i_pair}.png'
94     plot_image_and_depth(img0, depth0, img1, depth1, plot_name)
95
96     #(1) make meshgrid of points in image 0
97     x = np.linspace(10, img0.shape[1] - 10, num_points_in_mesh)
98     # ignore a border of 10 pxls
99     y = np.linspace(10, img0.shape[0] - 10, num_points_in_mesh)
100
101    # meshgrid of x and y coordinates
102    xx, yy = np.meshgrid(x, y)
103    xx, yy = xx.flatten(), yy.flatten()
104
105    # Make homogeneous coordinates for points0 [3, N]
106    points0 = np.vstack((xx, yy, np.ones_like(xx)))
107
108    # Get depth values at points0
109    depth_values0 = depth0[yy.astype(int), xx.astype(int)]
110
111    # Remove points with invalid depth (depth == 0)
112    valid_points = depth_values0 > 0
113    points0 = points0[:, valid_points]
114    depth_values0 = depth_values0[valid_points]
115
116    # (3) Find the 3D coordinates of these points in camera 0
117    frame
118        K0 = scene_info[i_pair]['K0'] # [3,3]
119        T0 = scene_info[i_pair]['T0'] # [4,4]
120
121        # inverse of K0
122        K0_inv = np.linalg.inv(K0)
123        # convert points0 to camera coordinates
124        xyz_cam0 = K0_inv @ points0
125        # normalize xyz_cam0 to set z = 1 (sanity check)

```

```

125     xyz_cam0 /= xyz_cam0[2, :]
126     # get the point at depth
127     xyz_cam0 *= depth_values0
128     # make homogeneous coordinates [4,N]
129     xyz_cam0_hc = np.vstack((xyz_cam0, np.ones((1, xyz_cam0.
shape[1]))))
130     # convert to world frame [4,N]
131     xyz_world_hc = np.linalg.inv(T0) @ xyz_cam0_hc
132
133     # (4) Transform these points to camera 1 frame
134     K1 = scene_info[i_pair]['K1']
135     T1 = scene_info[i_pair]['T1']
136
137     # transform points to camera 1 frame
138     xyz_cam1_hc = T1 @ xyz_world_hc
139     # convert to camera 1 coordinates
140     xyz_cam1 = xyz_cam1_hc[:3, :]
141     # get z coordinates for depth check
142     estimated_depth_values1 = xyz_cam1[2, :]
143
144
145     # project to image 1
146     points1 = K1 @ xyz_cam1 # [3, N]
147     # normalize by dividing by last row
148     points1 = points1 / points1[2, :] # [3, N]
149     # check if points1 are within image bounds
150     # Create mask
151     valid_projections = (
152         (points1[0, :] >= 0) & (points1[0, :] < img1.shape[1])
153     & (points1[1, :] >= 0) & (points1[1, :] < img1.shape[0])
154         )
155     # Apply mask
156     points1 = points1[:, valid_projections]
157     estimated_depth_values1 = estimated_depth_values1[
158     valid_projections]
159     # get the depth values at these points using the depth map
160     true_depth_values1 = depth1[points1[1, :].astype(int),
161     points1[0, :].astype(int)]
162
163     if(not dense):
164         # (5) plot matching points in image 0 and image 1 with
165         # depth check such that the depth values match
166         fig, ax = plt.subplots(1, 1, figsize=(10, 5))
167
168         # Horizontally stack the images
169         combined_img = np.ones((max(img0.shape[0], img1.shape
170             [0]), img0.shape[1] + img1.shape[1], 3), dtype=np.uint8) * 255
171         combined_img[:img0.shape[0], :img0.shape[1]] = img0
172         combined_img[:img1.shape[0], img0.shape[1]:] = img1
173
174         ax.imshow(combined_img, aspect='auto')
175         ax.scatter(xx, yy, c='r', s=5)
176         ax.set_title('Matching points in Image 0 and Image 1')
177         ax.axis('off')
178
179         # draw lines between matching points

```

```

176     for i in range(points1.shape[1]):
177         # if depth values match
178         if np.abs(estimated_depth_values1[i] -
179 true_depth_values1[i]) < DEPTH THR and true_depth_values1[i] != 0:
180             ax.plot([points0[0,i], points1[0, i] + img0.
181 shape[1]], [points0[1,i], points1[1, i]], 'g')
182
183         plt.savefig(f'./pics/depth_check_pair_{i_pair}.png',
184 bbox_inches='tight', pad_inches=0)
185         plt.close()
186         print(f"Done with pair {i_pair}")
187
188     if dense:
189         # (6) Plot all 3D points for the pair
190
191         fig = plt.figure(figsize=(10, 10))
192         ax = fig.add_subplot(111, projection='3d')
193
194         X, Y, Z = xyz_cam0[0, :], xyz_cam0[1, :], xyz_cam0[2,
195 :]
196
197         ax.scatter(X, Y, Z, c=Z, cmap='viridis', s=1)
198         ax.set_xlabel('X (Camera 0)')
199         ax.set_ylabel('Y (Camera 0)')
200         ax.set_zlabel('Z (Depth)')
201         plt.title(f'3D Points for Pair {i_pair}')
202         plt.savefig(f'./pics/3d_points_pair_{i_pair}.png',
203 bbox_inches='tight', pad_inches=0)
204         plt.close()
205         print(f"Done with pair {i_pair}")

```

Listing 3: Depth Check