

ECE 661 Homework 10

Michael Goldberg

December 6 2024

1 Theoretical Question

1.1 Overfitting

The ultimate goal of a machine learning algorithm is to find the solution that reduces the error the most across the training set yet generalizes well to the testing set. Overfitting refers to a found solution that cannot generalize well to the test set, as it has become overly reliant on the structure of the training set. For example, given an upwardly-trending set of points on a 2D scatter plot, one can draw a best-fit line through the data - this would be the desired solution, such that any test point can be predicted with the best-fit line. However, the overfit solution would be a high-degree polynomial that intersects all of the scatter points - while this provides us with very low training error, it cannot generalize to test data that lies outside of the polynomial.

1.2 Reparametrization Trick

I understand this in the following manner. The encoder maps the input data to a latent space defined by a mean and standard deviation that describes a probability distribution for the latent variable. However, this latent variable is sampled from the output of the encoder. This sampling makes backpropagation difficult, as a process with randomness like this is nondifferentiable.

Therefore, the trick is to introduce the latent variable as:

$$x = \mu + \epsilon\sigma \tag{1}$$

where μ is the mean, ϵ is a randomly sampled value between 0 and 1, and σ is the standard deviation. This allows us to still have a randomly sampled process, while remaining differentiable with respect to μ and σ for backpropagation.

2 Programming Tasks

The following task was to perform face recognition by using the PCA and LDA algorithms on a labelled dataset.

PCA I designed my own implementation of PCA and LDA. Let's start with PCA. The first step was to load in each image and vectorize them by flattening them into a single dimension. The mean of these vectors were found and subtracted from each value, transforming the vector to be zero-mean. Then, the data was normalized.

Once this was done, the goal of PCA is to reduce the dimensionality of the feature vector to a lower-dimensional space. This allows for rich encoding of the features for use within a classifier, ensuring computational inexpensiveness. Therefore, PCA suggests to encode the image vectors by using the eigenvectors that correspond to the p-largest eigenvalues in the eigendecomposition of the covariance matrix of the input data. However, the covariance matrix of the input data is incredibly large, yet a computational trick can be performed for less expensive eigendecomposition. Let's say the collection of all of the input images is in a data matrix called X , which is of size $N \times C$, where N is the number of pixels in the input image (in my case $64 \times 64 = 16384$), and C is the number of images (in my case 630). The following submatrix can be computed:

$$s = X^T * X \quad (2)$$

, which is of size $C \times C$, which is much smaller than the covariance matrix of size $N \times N$. We can then perform eigendecomposition on s , and find the p-largest eigenvectors, called v_s . However, these are within the submatrix space, and are not true eigenvectors of the covariance matrix. Therefore, we must map them back using:

$$v = X * v_s \quad (3)$$

Then, I normalized these true eigenvectors to form the feature vectors for the image.

2.1 LDA

LDA was more intensive, as this requires a process that uses the between- and within-class scatter matrices. First, the dimension of the input data is first reduced by running PCA to get the PCA eigenvectors, and mapping the input data to the PCA space through:

$$X_r = v * X \quad (4)$$

where v is the PCA eigenvectors, X is the input data, and X_r is the dimensionally reduced data. This then allows us to compute within-class and between-class variance using the following formulas:

Between-class variance:

$$S_B = \frac{1}{|M|} \sum_{i=1}^{|M|} (m_i - m)(m_i - m)^T \quad (5)$$

where M is the total number of unique faces within the training data, m_i is the single class mean and m is the global mean of the reduced data.

Within-class variance:

$$S_W = \frac{1}{|M|} \sum_{i=1}^{|M|} \frac{1}{|M_i|} \sum_{k=1}^{|M_i|} (x_k^i - m_i)(x_k^i - m_i)^T \quad (6)$$

where x_k^i is the kth image vector in the ith class.

We now want to solve the generalized eigenvalue problem, producing eigenvectors that satisfy the equation

$$S_B * v = \lambda * S_w * v \quad (7)$$

With these eigenvectors determined, we can then again collect the p-largest eigenvectors and treat that as our feature vector. Below is a comparison between the PCA and LDA methods for facial detection on the test set. First, we want to see how well these perform given different choices of p. The following plots show this relationship:

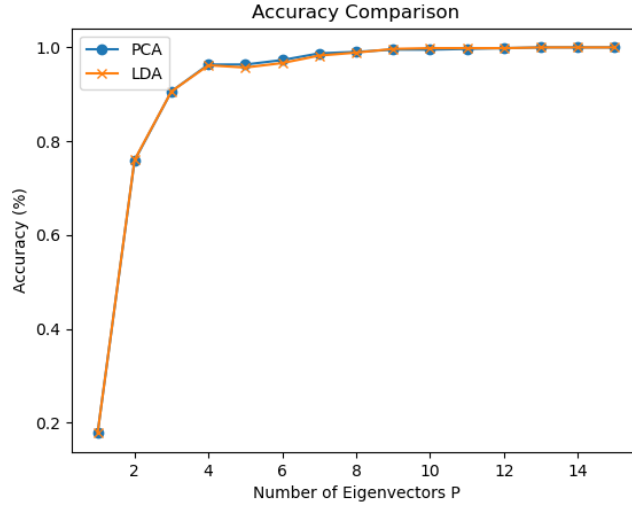




Figure 1: PCA and LDA Accuracy as a Function of p

It can be seen that both perform very well, and reach 100 accuracy when p becomes greater than 10. However, there are small mismatches between the two, with PCA generally performing better for this dataset.

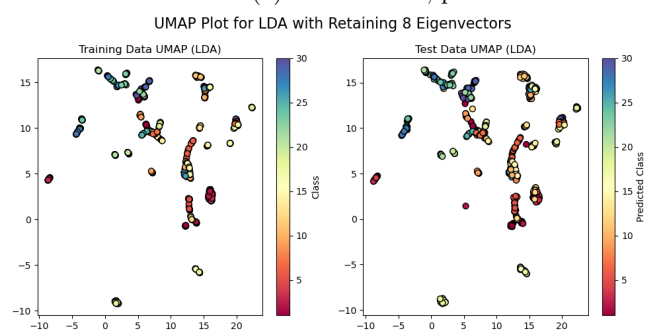
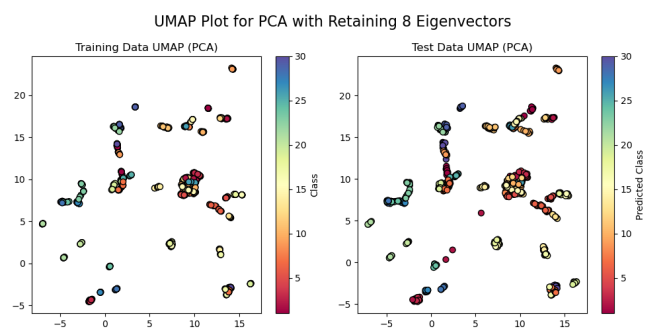
Next, let's look at how the different encodings look when projected to a 2D space for visualization:

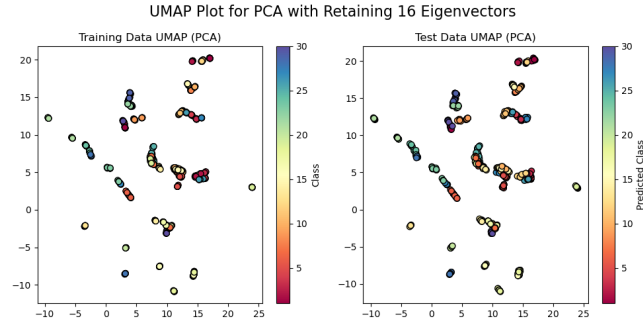


(a) PCA UMAP, $p=3$




(b) LDA UMAP, $p=3$





(a) PCA UMAP, $p=16$



(b) LDA UMAP, $p=16$

We can clearly see that the different strategies encode the rich features differently for the same value of p and same input data. However, when p increases to larger values, the 2D representations of the manifolds tend to converge to look more and more similar between the two approaches.

2.2 Autoencoder

Here, I used a variational autoencoder (AE) to encode the input data. I wanted to compare performance with PCA and LDA to see if this machine-learning-based approach had richer embeddings.


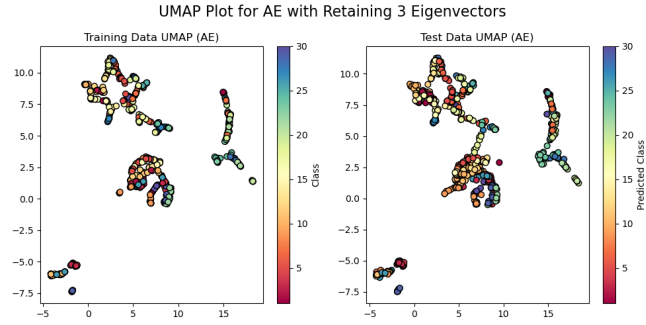



Figure 5: PCA and LDA and AE Accuracy as a Function of p


It can be seen that all three approaches match very closely to one another in terms of classification accuracy, with all tending towards 100 accuracy by the point that $p = 10$. Below are the UMAP projections for the AE approach at the same values of p (3, 8, 16):



(a) AE UMAP, $p=3$



(b) AE UMAP, $p=8$



(c) AE UMAP, $p=16$

It can again be seen that the autoencoder has a vastly different projection than the other two methods, yet as p tends to larger values, it converges to look more similar to PCA and LDA UMAP representations.

2.3 AdaBoost Cascade Classifier

We now are looking at a task that correctly identifies if an image is of a car or not. We used a cascade classifier built from weak adaboost classifiers to form an overall strong classifier. My approach was as follows:

First, I transformed the input images into integral images, and passed them through a vertical and horizontal haar feature detector to extract small vertical and horizontal features, reducing the dimensionality of the dataset for training. These features were then fed into a cascade classifier of multiple stages. Each stage, an adaboost classifier was trained on the data. Any particular feature vector was classified as true positive, true negative, false positive, and false negative determined from learned weights and learned threshold values. These threshold values were updated on each iteration of the adaboost training loop to discriminate more between positive and negative classification.

The cascade then determined how many samples from the input data pass the current stage, and only kept those for the next iteration. This repeats for a specified number of stages, until a strong classifier is built. The test data could then be passed through this final classifier and the false positive and false negative detection rates could be determined. Here are the results for a 2-stage classifier:

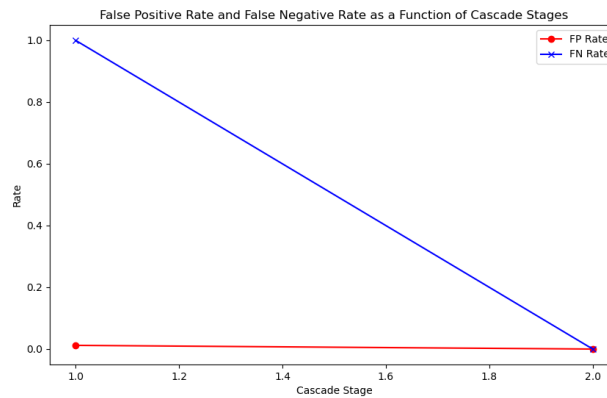


Figure 7: Cascaded Adaboost Classifier False Positive and False Negative Rate

3 Code

```

1 import os
2 import numpy as np
3 import cv2
4 import matplotlib.pyplot as plt
5 import time
6 import scipy
7 import umap
8
9 from sklearn.metrics import accuracy_score
10 from autoencoder import get_data
11
12
13 base_directory = 'FaceRecognition/'

```

```

14 train_directory = base_directory + 'train'
15 test_directory = base_directory + 'test'
16
17
18 def load_and_vectorize(directory):
19     image_vectors = []
20     normalized_vectors = []
21     labels = []
22
23     for filename in os.listdir(directory):
24         if filename.endswith(".png"):
25             # Read in grayscale image
26             image_path = os.path.join(directory, filename)
27             image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
28
29             # Vectorizing image and collecting label
30             image_vector = image.flatten()
31             image_vectors.append(image_vector)
32
33             labels.append(int(filename.split("_")[0]))
34
35     # Normalization and centering to be zero-mean
36     image_vectors = np.array(image_vectors)
37     mean = np.mean(image_vectors, axis=0)
38     centered_data = image_vectors - mean
39
40     normalized_data = centered_data / np.linalg.norm(centered_data,
41         axis=1, keepdims=True)
42
43     return normalized_data.T, labels
44
45 def pca(data, p):
46     # Compute submatrix, computational trick
47     submatrix = data.T @ data
48
49     # Eignedecomposition of the submatrix
50     sub_eigenvalues, sub_eigenvectors = np.linalg.eigh(submatrix)
51
52     # Sort in descending order, choose eigenvectors that correspond
53     # to the p-largest eigenvalues
54     descending_sorted_indices = np.argsort(sub_eigenvalues)[::-1]
55     sub_eigenvectors = sub_eigenvectors[:,
56         descending_sorted_indices][:, :p]
57
58     # Mapping back to be true eigenvectors of the covariance matrix
59     # + nnormalization
60     eigenvectors = data @ sub_eigenvectors
61     eigenvectors = eigenvectors / np.linalg.norm(eigenvectors, axis
62         =0)
63
64     return eigenvectors
65
66 def lda(data, labels, p):
67     # Perform PCA for dimensionality reduction
68     pca_eigenvectors = pca(data, p)
69     reduced_data = pca_eigenvectors.T @ data

```

```

66 # Compute within-class and between-class scatter matrices in
    PCA space
67
68 # Calculating within-class and between-class scatter matrices
69 S_W_reduced = np.zeros((reduced_data.shape[0], reduced_data.
    shape[0]))
70 S_B_reduced = np.zeros((reduced_data.shape[0], reduced_data.
    shape[0]))
71
72
73 unique_classes = np.unique(labels)
74 overall_mean = np.mean(reduced_data, axis=1)
75
76 # Calculation of within-class and between-class scatter
77 for class_ in unique_classes:
78     class_data = reduced_data[:, np.array(labels) == class_]
79     class_mean = np.mean(class_data, axis=1)
80     class_scatter = (class_data - class_mean[:, np.newaxis]) @
    (class_data - class_mean[:, np.newaxis]).T
81     S_W_reduced += class_scatter
82
83     mean_difference = (class_mean - overall_mean).reshape(-1,
    1)
84     S_B_reduced += class_data.shape[1] * (mean_difference @
    mean_difference.T)
85
86
87 # Eigendecomposition and keeping the eigenvectors corresponding
    to the p-largest eigenvalues
88 lda_eigenvalues, lda_eigenvectors = scipy.linalg.eigh(
    S_B_reduced, S_W_reduced + np.eye(S_W_reduced.shape[0]) * 1e-6)
89 descending_sorted_indices = np.argsort(lda_eigenvalues)[::-1]
90 lda_eigenvectors = lda_eigenvectors[:,
    descending_sorted_indices][:, :p]
91
92 # Step 4: Map LDA eigenvectors back to original space +
    normalization
93 lda_eigenvectors = pca_eigenvectors @ lda_eigenvectors
94 lda_eigenvectors = lda_eigenvectors / np.linalg.norm(
    lda_eigenvectors, axis=0)
95
96
97 return lda_eigenvectors
98
99 def project_to_subspace(data, pca_feature_set):
100     return pca_feature_set.T @ data
101
102 def nearest_neighbor(train_projected, train_labels, test_projected)
    :
103     predictions = []
104
105     for test_sample in test_projected.T:
106         # Calculate L2 norm to all training samples
107         distances = np.linalg.norm(train_projected.T - test_sample,
            axis=1)
108
109         # Find the index of the closest training sample

```

```

110         nearest_index = np.argmin(distances)
111
112         # Assign label
113         predictions.append(train_labels[nearest_index])
114
115     return predictions
116
117
118 def plot_umap(train_data, train_labels, test_data, test_labels,
119               predicted_labels, method, p):
120     # Apply UMAP to reduce the data to 2D
121     reducer = umap.UMAP(n_components=2, random_state=42)
122
123     # Apply UMAP on training data
124     train_umap = reducer.fit_transform(train_data.T)
125     # Apply UMAP on test data
126     test_umap = reducer.transform(test_data.T)
127
128     plt.figure(figsize=(10, 5))
129
130     # Plot training data
131     plt.subplot(1, 2, 1)
132     plt.scatter(train_umap[:, 0], train_umap[:, 1], c=train_labels,
133                cmap='Spectral', edgecolors='k', s=40)
134     plt.title(f"Training Data UMAP ({method})")
135     plt.colorbar(label="Class")
136
137     # Plot test data
138     plt.subplot(1, 2, 2)
139     plt.scatter(test_umap[:, 0], test_umap[:, 1], c=
140                predicted_labels, cmap='Spectral', edgecolors='k', s=40)
141     plt.title(f"Test Data UMAP ({method})")
142     plt.colorbar(label="Predicted Class")
143
144     plt.suptitle(f"UMAP Plot for {method} with Retaining {p}
145                 Eigenvectors", fontsize=16)
146     plt.tight_layout()
147     plt.show()
148
149
150 if __name__ == "__main__":
151     train_data, train_labels = load_and_vectorize(train_directory)
152     test_data, test_labels = load_and_vectorize(test_directory)
153
154     print("train data shape: ", train_data.shape)
155
156     time.sleep(100)
157
158     p_set = range(1, 17)
159     ae_p_set = [3, 8, 16]
160     pca_accuracies = []
161     lda_accuracies = []
162     ae_accuracies = []
163
164     plot_umaps = True
165
166     for p in p_set:

```

```

163         # Perform PCA
164
165         # Collect feature set
166         pca_feature_set = pca(train_data, p)
167         # Project to PCA subspace
168         pca_train_feature_vector = project_to_subspace(train_data,
169         pca_feature_set)
170         pca_test_feature_vector = project_to_subspace(test_data,
171         pca_feature_set)
172         # print("lda_train_feature_vector shape: ",
173         pca_train_feature_vector.shape)
174         # Predict labels with nearest neighbor algorithm
175         pca_predicted_labels = nearest_neighbor(
176         pca_train_feature_vector, train_labels, pca_test_feature_vector
177         )
178         # Calculate accuracy through #correct_predictions / #
179         total_images
180         pca_accuracy = accuracy_score(test_labels,
181         pca_predicted_labels)
182
183         # Perform LDA
184         # Collect feature set
185         lda_feature_set = lda(train_data, train_labels, p)
186         # Project to LDA subspace
187         lda_train_feature_vector = project_to_subspace(train_data,
188         lda_feature_set)
189         lda_test_feature_vector = project_to_subspace(test_data,
190         lda_feature_set)
191         # Predict labels with nearest neighbor algorithm
192         lda_predicted_labels = nearest_neighbor(
193         lda_train_feature_vector, train_labels, lda_test_feature_vector
194         )
195         # Calculate accuracy through #correct_predictions / #
196         total_images
197         lda_accuracy = accuracy_score(test_labels,
198         lda_predicted_labels)
199
200         pca accuracies.append(pca_accuracy)
201         lda accuracies.append(lda_accuracy)
202
203         print(f"p = {p}: PCA Accuracy = {pca_accuracy:.2%}, LDA
204         Accuracy = {lda_accuracy:.2%}")
205
206         if (p == 3 or p == 8 or p == 16):
207             # Load autoencoder vectors and labels
208             ae_train_feature_vector, ae_train_labels,
209             ae_test_feature_vector, ae_test_labels = get_data(training=
210             False, p=p)
211             # Converting to be of shape (p, C) rather than (C, p)
212             ae_train_feature_vector = ae_train_feature_vector.T
213             ae_test_feature_vector = ae_test_feature_vector.T
214
215             # Use nearest neighbors to predict test labels with
216             autoencoder embeddings

```

```

202         ae_predicted_labels = nearest_neighbor(
203             ae_train_feature_vector, ae_train_labels,
204             ae_test_feature_vector)
205         # Calculate autoencoder accuracy
206         ae_accuracy = accuracy_score(test_labels,
207             ae_predicted_labels)
208         ae accuracies.append(ae_accuracy)
209
210         print(f"p = {p}: AE Accuracy = {ae_accuracy:.2%}")
211
212         if(plot_umaps):
213             plot_umap(pca_train_feature_vector, train_labels,
214                 pca_test_feature_vector, test_labels, pca_predicted_labels, "
215                 PCA", p)
216             plot_umap(lda_train_feature_vector, train_labels,
217                 lda_test_feature_vector, test_labels, lda_predicted_labels, "
218                 LDA", p)
219             plot_umap(ae_train_feature_vector, ae_train_labels,
220                 ae_test_feature_vector, ae_test_labels, ae_predicted_labels, "
221                 AE", p)
222
223     plt.plot(p_set, pca_accuracies, "-o", label="PCA")
224     plt.plot(p_set, lda_accuracies, "-x", label="LDA")
225     plt.plot(ae_p_set, ae_accuracies, "-*", label="AE")
226     plt.title("Accuracy Comparison")
227     plt.xlabel("Number of Eigenvectors P")
228     plt.ylabel("Accuracy (%)")
229     plt.legend()
230     plt.show()

```

Listing 1: PCA, LDA, and Autoencoder Classification

```

1  import os
2
3  import numpy as np
4  import torch
5  from torch import nn, optim
6  from PIL import Image
7  from torch.autograd import Variable
8  from torch.utils.data import Dataset, DataLoader
9  from torchvision import transforms
10
11
12  class DataBuilder(Dataset):
13      def __init__(self, path):
14          self.path = path
15          self.image_list = [f for f in os.listdir(path) if f.
16              endswith('.png')]
17          self.label_list = [int(f.split('_')[0]) for f in self.
18              image_list]
19          self.len = len(self.image_list)
20          self.aug = transforms.Compose([
21              transforms.Resize((64, 64)),
22              transforms.ToTensor(),
23          ])
24
25      def __getitem__(self, index):

```

```

24     fn = os.path.join(self.path, self.image_list[index])
25     x = Image.open(fn).convert('RGB')
26     x = self.aug(x)
27     return {'x': x, 'y': self.label_list[index]}
28
29     def __len__(self):
30         return self.len
31
32
33 class Autoencoder(nn.Module):
34
35     def __init__(self, encoded_space_dim):
36         super().__init__()
37         self.encoded_space_dim = encoded_space_dim
38         ### Convolutional section
39         self.encoder_cnn = nn.Sequential(
40             nn.Conv2d(3, 8, 3, stride=2, padding=1),
41             nn.LeakyReLU(True),
42             nn.Conv2d(8, 16, 3, stride=2, padding=1),
43             nn.LeakyReLU(True),
44             nn.Conv2d(16, 32, 3, stride=2, padding=1),
45             nn.LeakyReLU(True),
46             nn.Conv2d(32, 64, 3, stride=2, padding=1),
47             nn.LeakyReLU(True)
48         )
49         ### Flatten layer
50         self.flatten = nn.Flatten(start_dim=1)
51         ### Linear section
52         self.encoder_lin = nn.Sequential(
53             nn.Linear(4 * 4 * 64, 128),
54             nn.LeakyReLU(True),
55             nn.Linear(128, encoded_space_dim * 2)
56         )
57         self.decoder_lin = nn.Sequential(
58             nn.Linear(encoded_space_dim, 128),
59             nn.LeakyReLU(True),
60             nn.Linear(128, 4 * 4 * 64),
61             nn.LeakyReLU(True)
62         )
63         self.unflatten = nn.Unflatten(dim=1,
64                                       unflattened_size=(64, 4, 4))
65         self.decoder_conv = nn.Sequential(
66             nn.ConvTranspose2d(64, 32, 3, stride=2,
67                               padding=1, output_padding=1),
68             nn.BatchNorm2d(32),
69             nn.LeakyReLU(True),
70             nn.ConvTranspose2d(32, 16, 3, stride=2,
71                               padding=1, output_padding=1),
72             nn.BatchNorm2d(16),
73             nn.LeakyReLU(True),
74             nn.ConvTranspose2d(16, 8, 3, stride=2,
75                               padding=1, output_padding=1),
76             nn.BatchNorm2d(8),
77             nn.LeakyReLU(True),
78             nn.ConvTranspose2d(8, 3, 3, stride=2,
79                               padding=1, output_padding=1)
80         )

```

```

81
82     def encode(self, x):
83         x = self.encoder_cnn(x)
84         x = self.flatten(x)
85         x = self.encoder_lin(x)
86         mu, logvar = x[:, :self.encoded_space_dim], x[:, self.
encoded_space_dim:]
87         return mu, logvar
88
89     def decode(self, z):
90         x = self.decoder_lin(z)
91         x = self.unflatten(x)
92         x = self.decoder_conv(x)
93         x = torch.sigmoid(x)
94         return x
95
96     @staticmethod
97     def reparameterize(mu, logvar):
98         std = logvar.mul(0.5).exp_()
99         eps = Variable(std.data.new(std.size()).normal_())
100         return eps.mul(std).add_(mu)
101
102
103 class VaeLoss(nn.Module):
104     def __init__(self):
105         super(VaeLoss, self).__init__()
106         self.mse_loss = nn.MSELoss(reduction="sum")
107
108     def forward(self, xhat, x, mu, logvar):
109         loss_MSE = self.mse_loss(xhat, x)
110         loss_KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar
.exp())
111         return loss_MSE + loss_KLD
112
113
114 def train(epoch):
115     model.train()
116     train_loss = 0
117
118     for batch_idx, data in enumerate(trainloader):
119         optimizer.zero_grad()
120         mu, logvar = model.encode(data['x'])
121         z = model.reparameterize(mu, logvar)
122         xhat = model.decode(z)
123         loss = vae_loss(xhat, data['x'], mu, logvar)
124         loss.backward()
125         train_loss += loss.item()
126         optimizer.step()
127
128     print('====> Epoch: {} Average loss: {:.4f}'.format(
129         epoch, train_loss / len(trainloader.dataset)))
130
131
132 def get_data(training=False, p=3):
133     #####
134     TRAIN_DATA_PATH = 'FaceRecognition/train'
135     EVAL_DATA_PATH = 'FaceRecognition/test'

```



```

136 LOAD_PATH = f'weights/model_{p}.pt'
137 OUT_PATH = LOAD_PATH
138 #####
139
140 model = Autoencoder(p)
141
142 if training:
143     epochs = 100
144     log_interval = 1
145     trainloader = DataLoader(
146         dataset=DataBuilder(TRAIN_DATA_PATH),
147         batch_size=12,
148         shuffle=True,
149     )
150     optimizer = optim.Adam(model.parameters(), lr=1e-3)
151     vae_loss = VaeLoss()
152     for epoch in range(1, epochs + 1):
153         train(epoch)
154         torch.save(model.state_dict(), os.path.join(OUT_PATH, f'
model_{p}.pt'))
155     else:
156         trainloader = DataLoader(
157             dataset=DataBuilder(TRAIN_DATA_PATH),
158             batch_size=1,
159         )
160         model.load_state_dict(torch.load(LOAD_PATH, weights_only=
True))
161         model.eval()
162
163         X_train, y_train = [], []
164         for batch_idx, data in enumerate(trainloader):
165             mu, logvar = model.encode(data['x'])
166             z = mu.detach().cpu().numpy().flatten()
167             X_train.append(z)
168             y_train.append(data['y'].item())
169         X_train = np.stack(X_train)
170         y_train = np.array(y_train)
171
172         testloader = DataLoader(
173             dataset=DataBuilder(EVAL_DATA_PATH),
174             batch_size=1,
175         )
176         X_test, y_test = [], []
177         for batch_idx, data in enumerate(testloader):
178             mu, logvar = model.encode(data['x'])
179             z = mu.detach().cpu().numpy().flatten()
180             X_test.append(z)
181             y_test.append(data['y'].item())
182         X_test = np.stack(X_test)
183         y_test = np.array(y_test)
184
185         return X_train, y_train, X_test, y_test

```

Listing 2: Autoencoder Helper Script

```

1 import os
2 import cv2

```

```

3 import numpy as np
4 import matplotlib.pyplot as plt
5 from tqdm import tqdm
6
7 num_iterations = 20
8
9 class WeakClassifier:
10     def __init__(self, feature, threshold, polarity):
11         self.feature = feature
12         self.threshold = threshold
13         self.polarity = polarity
14
15     def predict(self, features):
16         feature_value = features[self.feature]
17         return 1 if (self.polarity == 1 and feature_value >= self.
18             threshold) or (self.polarity == -1 and feature_value < self.
19             threshold) else -1
20
21 class AdaBoost:
22     def __init__(self, T):
23         self.T = T
24         self.alphas = []
25         self.classifiers = []
26
27     def fit(self, X, y):
28         w = np.ones(len(X)) / len(X) # Initial weight for each
29         sample
30
31         for t in tqdm(range(self.T), desc="Training AdaBoost",
32             ncols=100):
33             best_classifier = None
34             min_error = float('inf')
35
36             for feature_index in range(len(X[0])):
37                 # Sorting features
38                 feature_values = [features[feature_index] for
39                     features in X]
40                 sorted_indices = np.argsort(feature_values)
41                 sorted_features = np.array(feature_values)[
42                     sorted_indices]
43                 sorted_weights = w[sorted_indices]
44                 sorted_labels = np.array(y)[sorted_indices]
45
46                 # Calculating next threshold value
47                 T_plus = np.sum(w * (y == 1))
48                 T_minus = np.sum(w * (y == -1))
49
50                 for i in range(1, len(X)):
51                     threshold = (sorted_features[i - 1] +
52                         sorted_features[i]) / 2
53                     S_plus = np.sum(sorted_weights[:i] * (
54                         sorted_labels[:i] == 1))
55                     S_minus = np.sum(sorted_weights[:i] * (
56                         sorted_labels[:i] == -1))
57
58                     error_pos_1 = S_plus + (T_minus - S_minus)

```

```

51         error_neg_1 = S_minus + (T_plus - S_plus)
52
53         # Choosing minimum between two error metrics
54         error = min(error_pos_1, error_neg_1)
55
56         if error < min_error:
57             min_error = error
58             best_classifier = WeakClassifier(
feature_index, threshold, 1)
59
60         # Find new parameters to compute next weak classifier
61         alpha = 0.5 * np.log((1 - min_error) / (min_error + 1e
-10))
62         self.alphas.append(alpha)
63         self.classifiers.append(best_classifier)
64
65         predictions = np.array([best_classifier.predict(
features) for features in X])
66         w = w * np.exp(-alpha * y * predictions)
67         w = w / np.sum(w) # Normalize weights
68
69
70     def predict(self, X):
71         strong_preds = np.zeros(len(X))
72         for alpha, classifier in zip(self.alphas, self.classifiers)
:
73             predictions = np.array([classifier.predict(features)
for features in X])
74             strong_preds += alpha * predictions
75         return np.sign(strong_preds)
76
77
78     class CascadeClassifier:
79         def __init__(self, false_positive_target, true_detection_target
, num_stages):
80             self.false_positive_target = false_positive_target
81             self.true_detection_target = true_detection_target
82             self.num_stages = num_stages
83             self.stages = []
84             self.fp_rates = []
85             self.fn_rates = []
86
87         def train(self, X, y):
88
89             for stage_index in range(self.num_stages):
90                 print(f"Training stage {stage_index + 1}/{self.
num_stages}...")
91                 # Perform adaboost to fit weak classifier to data
92                 adaboost = AdaBoost(T=num_iterations)
93                 adaboost.fit(X, y)
94                 self.stages.append(adaboost)
95
96                 # Get predictions on the current stage's data
97                 predictions = adaboost.predict(X)
98
99                 # Compute False Positive and False Negative rates
100                 fp = np.sum((predictions == 1) & (y == 0)) # Positive

```

```

classified as negative
101         fn = np.sum((predictions == -1) & (y == 1)) # Negative
classified as positive
102
103         fp_rate = fp / np.sum(y == 0) if np.sum(y == 0) > 0
else 0
104         fn_rate = fn / np.sum(y == 1) if np.sum(y == 1) > 0
else 0
105
106         self.fp_rates.append(fp_rate)
107         self.fn_rates.append(fn_rate)
108
109         print(f"Stage {stage_index + 1}: FP rate = {fp_rate},
FN rate = {fn_rate}")
110
111         # If the stage does not meet target FP and FN rates,
stop training
112         if fp_rate > self.false_positive_target or fn_rate <
self.true_detection_target:
113             print(f"Stage {stage_index + 1} did not meet target
rates. Stopping early.")
114             break
115
116         # Keep only the samples that pass the current stage
passed_indices = (predictions == 1)
117         X = np.array(X)[passed_indices.astype(int)]
118         y = np.array(y)[passed_indices.astype(int)]
119
120         if len(X) == 0:
121             break
122
123         self.plot_performance()
124
125     def predict(self, X):
126         for stage in self.stages:
127             predictions = stage.predict(X)
128             if np.any(predictions == -1): # Reject if any stage
129                 fails
130                 return -1
131                 return 1
132
133     def plot_performance(self):
134         # Plotting FP and FN rates
135         stages = np.arange(1, len(self.fp_rates) + 1)
136
137         plt.figure(figsize=(10, 6))
138
139         # Plot FP rate
140         plt.plot(stages, self.fp_rates, label="FP Rate", color="red
", marker='o')
141         # Plot FN rate
142         plt.plot(stages, self.fn_rates, label="FN Rate", color="
blue", marker='x')
143
144         # Labels and title
145         plt.xlabel("Cascade Stage")
146         plt.ylabel("Rate")

```

```

147     plt.title("False Positive Rate and False Negative Rate as a
148             Function of Cascade Stages")
149
150     # Show a legend
151     plt.legend()
152
153     # Show the plot
154     plt.show()
155
156 def extract_haar_features(integral_image):
157     img_height, img_width = integral_image.shape
158     features = []
159
160     # Horizontal 1x2 feature
161     for y in range(img_height):
162         for x in range(img_width - 1):
163             left = integral_image[y, x]
164             right = integral_image[y, x + 1]
165             horizontal_feature = right - left
166             features.append(horizontal_feature)
167
168     # Vertical 2x1 feature
169     for y in range(img_height - 1):
170         for x in range(img_width):
171             top = integral_image[y, x]
172             bottom = integral_image[y + 1, x]
173             vertical_feature = bottom - top
174             features.append(vertical_feature)
175
176     return np.array(features)
177
178
179 def compute_integral_image(image):
180     return image.cumsum(axis=0).cumsum(axis=1)
181
182
183 def load_data(positive_dir, negative_dir):
184     images = []
185     labels = []
186
187     # Loading data and labels
188     for filename in os.listdir(positive_dir):
189         img = cv2.imread(os.path.join(positive_dir, filename), cv2.
190             IMREAD_GRAYSCALE)
191         images.append(img)
192         labels.append(1)
193
194     for filename in os.listdir(negative_dir):
195         img = cv2.imread(os.path.join(negative_dir, filename), cv2.
196             IMREAD_GRAYSCALE)
197         images.append(img)
198         labels.append(0)
199
200     images = np.array(images)
201     labels = np.array(labels)

```

```

201     return images, labels
202
203
204 # Load training data
205 train_dir = "CarDetection/train"
206 positive_train_dir = os.path.join(train_dir, "positive")
207 negative_train_dir = os.path.join(train_dir, "negative")
208
209 positive_train_images, positive_train_labels = load_data(
210     positive_train_dir, negative_train_dir)
211
212 train_integral_images = [compute_integral_image(img) for img in
213     positive_train_images]
214 train_features = [extract_haar_features(integral_image) for
215     integral_image in train_integral_images]
216
217 # Train cascade classifier
218 cascade_classifier = CascadeClassifier(false_positive_target=0.1,
219     true_detection_target=0.9, num_stages=5)
220 cascade_classifier.train(train_features, positive_train_labels)
221
222 # Test the classifier
223 test_dir = "CarDetection/test"
224 positive_test_dir = os.path.join(test_dir, "positive")
225 negative_test_dir = os.path.join(test_dir, "negative")
226
227 positive_test_images, positive_test_labels = load_data(
228     positive_test_dir, negative_test_dir)
229
230 test_integral_images = [compute_integral_image(img) for img in
231     positive_test_images]
232 test_features = [extract_haar_features(integral_image) for
233     integral_image in test_integral_images]
234
235 test_predictions = cascade_classifier.predict(test_features)
236 # test_predictions = [cascade_classifier.predict(features) for
237     features in test_features]
238 accuracy = np.mean(np.array(test_predictions) == np.array(
239     positive_test_labels))
240 print(f"Cascade Classifier Test Accuracy: {accuracy * 100:.2f}%")

```

Listing 3: Cascaded Adaboost Classifier